

Modeling with k-means, Part 3

In this notebook I investigate how well the k-means algorithm distinguishes between the INLAND California housing districts and the districts closer to the ocean. The INLAND districts are already identified as such. I start with 11 predictors. The longitude and latitude variables are removed in order to make the challenge meaningful. Variables remaining include, among others, the median house value of each district, the median housing age of a district, the population per household in the district, and so forth. For more information on this dataset, see my CA housing analysis project. (This project involves a set of Jupyter notebooks. See the AnalyzeCAhousingData repository.)

As with the downer cow dataset, I want to see whether we can construct a better classification model using k-means in combination with one or more supervised learning algorithms. Applying weights to the variables in this dataset will be more difficult because there are more variables to work with and quite a bit more data, so finding the optimal weights will take some computing power.

* * * * *

Preliminaries

```
In [ ]: require(car)
require(repr)
require(ggplot2)
require(stringr)
require(parallel)
require(faraway)
require(randomForest)
require(gbm)
require(plyr)
require(e1071)
```

```
In [2]: options(digits= 5, show.signif.stars= FALSE)
```

Construct training and test sets

```
In [3]: traindat <- read.csv("/home/greg/Documents/stat/Geron_ML/datasets/housing/R_train_Part02-2_
header = T, row.names= 1, colClasses= c("character", rep("numeric", 9), r
rep("numeric", 4), "character", "numeric", "character"))
dim(traindat)
```

```
16308 · 18
```

```
In [4]: colnames(traindat)
```

```
'longitude' · 'latitude' · 'housing_median_age' · 'total_rooms' · 'total_bedrooms' · 'population' · 'households' ·
'median_income' · 'median_house_value' · 'ocean_proximity' · 'HHdensity' · 'HHdens_ln' · 'rooms_per_hh' ·
'bdrms_per_room' · 'pop_per_hh' · 'op_transf' · 'long_transf' · 'income_cat'
```

```
In [22]: table(as.factor(traindat$op_transf))
```

```
INLAND NEAR BAY      OCEAN
5172    1816    9320
```

```
In [23]: testdat <- read.csv("/home/greg/Documents/stat/Geron_ML/datasets/housing/R_test_Part02-2_St
header = T, row.names= 1,
```

```
colClasses= c("character", rep("numeric", 9),
              rep("character", 2), rep("numeric", 4),
              "character", "numeric", "character"))
dim(testdat)
3888 · 18
```

In [24]: colnames(testdat)

```
'longitude' · 'latitude' · 'housing_median_age' · 'total_rooms' · 'total_bedrooms' · 'population' · 'households' ·
'median_income' · 'median_house_value' · 'ocean_proximity' · 'HHdensity' · 'HHdens_In' · 'rooms_per_hh' ·
'bdrms_per_room' · 'pop_per_hh' · 'op_transf' · 'long_transf' · 'income_cat'
```

In [25]: dat <- rbind(traindat, testdat)
dim(dat)
table(as.factor(dat\$op_transf))

```
20196 · 18
```

```
INLAND NEAR BAY OCEAN
6467 2235 11494
```

In [26]: dat\$Inland <- 0
dat[which(dat\$op_transf== "INLAND"),]\$Inland <- 1
table(dat\$Inland)

```
0 1
13729 6467
```

In [27]: newcols <- colnames(dat)[which(!(colnames(dat) %in% c("HHdensity", "income_cat",
"ocean_proximity", "op_transf")))]
newcols

```
'longitude' · 'latitude' · 'housing_median_age' · 'total_rooms' · 'total_bedrooms' · 'population' · 'households' ·
'median_income' · 'median_house_value' · 'HHdens_In' · 'rooms_per_hh' · 'bdrms_per_room' · 'pop_per_hh' ·
'long_transf' · 'Inland'
```

In [28]: dat <- dat[, newcols]
dim(dat)
colnames(dat)

```
20196 · 15
```

```
'longitude' · 'latitude' · 'housing_median_age' · 'total_rooms' · 'total_bedrooms' · 'population' · 'households' ·
'median_income' · 'median_house_value' · 'HHdens_In' · 'rooms_per_hh' · 'bdrms_per_room' · 'pop_per_hh' ·
'long_transf' · 'Inland'
```

In [29]: *# Create training and test sets for use in testing*
new k-means approach.

```
set.seed(8763)
smp <- sample(rownames(dat), 10196, replace=FALSE)
traindat <- dat[smp,]
testdat <- dat[which(!(rownames(dat) %in% smp)),]
dim(testdat)
```

```
10000 · 15
```

In [30]: rm(dat)

In [31]: *# Remove latitude, longitude, and long_transf from*
both traindat and testdat. If these variables
remained in the dataset, it would be far too easy

```
# to distinguish between the inland and oceanside
# districts.

newcols <- colnames(traindat)[which(!(colnames(traindat) %in%
                                     c("latitude", "longitude", "long_transf")))]
newcols

'housing_median_age' · 'total_rooms' · 'total_bedrooms' · 'population' · 'households' · 'median_income' ·
'median_house_value' · 'HHdens_In' · 'rooms_per_hh' · 'bdrms_per_room' · 'pop_per_hh' · 'Inland'
```

```
In [32]: traindat <- traindat[, newcols]
testdat <- testdat[, newcols]
dim(traindat)

10196 · 12
```

```
In [33]: # Move the response variable, Inland, to the first
# column.

newcols <- colnames(traindat)
newcols <- newcols[which(!(newcols %in% c("Inland")))]
newcols <- c("Inland", newcols)
newcols

'Inland' · 'housing_median_age' · 'total_rooms' · 'total_bedrooms' · 'population' · 'households' · 'median_income' ·
'median_house_value' · 'HHdens_In' · 'rooms_per_hh' · 'bdrms_per_room' · 'pop_per_hh'
```

```
In [34]: traindat <- traindat[, newcols]
testdat <- testdat[, newcols]
```

```
In [35]: # Save out the files.

write.csv(traindat, file="/home/greg/Documents/stat/github_repos/cows/CAhousing_train_10196
          row.names= TRUE)

write.csv(testdat, file="/home/greg/Documents/stat/github_repos/cows/CAhousing_test_10K.csv
          row.names= TRUE)
```

```
In [5]: traindat <- read.csv("/home/greg/Documents/stat/github_repos/cows/CAhousing_train_10196.csv
                             header=TRUE, row.names=1, colClasses= c("character", rep("numeric",12))
dim(traindat)

10196 · 12
```

```
In [6]: testdat <- read.csv("/home/greg/Documents/stat/github_repos/cows/CAhousing_test_10K.csv",
                             header=TRUE, row.names=1, colClasses= c("character", rep("numeric",12))
dim(testdat)

10000 · 12
```

```
In [5]: # Get ratio of Inland to non-Inland in traindat.

round(mean(traindat$Inland), 4)

0.3187
```

```
In [6]: # Get ratio of Inland to non-Inland in testdat.

round(mean(testdat$Inland), 4)

0.3218
```

Basic functions

```
In [7]: # Function for computing "proportion of deviance explained", an
# R-sqrd statistic for link= logit. See p.41 of Julian Faraway's
# "Extending the Linear Model with R" (2006: Chapman & Hall).
# This statistic is due to N. Nagelkerke (1991; see Faraway's
# bibliography).

get_RsqrdDev <- function(modl) {

  n_obs <- length(modl$fitted.values)
  ans <- (1-exp((modl$deviance - modl$null.deviance)/n_obs))/(1-exp(-modl$null.deviance/n_obs))
  return(round(ans, 4))
}
```

```
In [8]: # Function to compute f-score for a 2x2 confusion matrix.

get_fscores <- function(mat) {

  FN <- as.numeric(mat[2,1])
  TP <- as.numeric(mat[2,2])
  FP <- as.numeric(mat[1,2])
  recall <- TP/(TP + FN)
  precision <- TP/(TP + FP)
  f_score <- 2* (recall*precision)/(recall + precision)
  return(round(f_score, 4))
}
```

```
In [9]: # Function to output a confusion matrix and the f-score
# for that matrix (if it is 2x2).

get_confusion <- function(preds, df_actual) {

  # df_actual is a one-column dataframe;
  # preds is a named vector of predictions;
  # preds is of type factor; it is assumed there are
  # at least 2 factor levels

  levs <- levels(preds)
  n_levs <- length(levs)
  if(n_levs== 1) { levs <- c('0', '1') }
  n_levs <- max(n_levs, 2)
  actual <- as.vector(df_actual[, 1])
  names(actual) <- rownames(df_actual)

  datout <- rep(0, n_levs * (n_levs + 1))
  dim(datout) <- c(n_levs, n_levs + 1)
  datout <- as.data.frame(datout)
  colnames(datout) <- c(levs, "class.error")
  rownames(datout) <- levs

  result <- vector("list", length= 2)
  names(result) <- c("matrix", "f_score")

  # for each factor level, identify the rcd names
  # which should be classed as such
  for(rowlev in levs) {
    actlev_names <- names(actual[actual == rowlev])
    # columns are for the predicted values:
    for(collev in levs) {
      predlev_names <- names(preds[preds == collev])
      if(length(predlev_names > 0)) {
        datout[rowlev, collev] <- sum(predlev_names %in% actlev_names)
      }
    }
  }
  nonrow_cols <- levs[!(levs %in% rowlev)]
}
```

```

        datout[rowlev, "class.error"] <- round(sum(as.vector(datout[rowlev, nonrow_cols]))/
                                              sum(as.vector(datout[rowlev, levs])), 4)
    }

    result$matrix <- datout
    if(n_levs == 2) {
        result[[2]] <- get_fscore(as.matrix(datout))
    } else {
        result$f_score <- NA
    }

    return(result)
}

```

In [10]: *# The following function is from Robert Kabacoff's "R in Action", pp.379-380.*

```

wssplot <- function(data, title="", nc=15, seed=1233) {

    # wss[1] is the total sum of squares when there is only
    # one cluster. In R's kmeans help this is called 'totss'.
    # Here is another way to compute totss:
    # ss <- function(x) sum(scale(x, scale = FALSE)^2)
    wss <- (nrow(data) - 1)*sum(apply(data, 2, var))
    for(i in 2:nc) {
        set.seed(seed)
        km_model <- suppressWarnings(kmeans(data, centers=i, iter.max=50,
                                              nstart=5))

        wss[i] <- sum(km_model$withinss)
    }
    plot(1:nc, wss, type='b', xlab="Number of clusters",
         ylab="Within groups sums of squares",
         main= title)
}

```

In [11]: *# Function for identifying which cluster each record
belongs to.*

```

getCluster <- function(x, centers) {

    # x is a row of a dataframe; its columns need
    # to be in the same order as centers (a matrix)
    # constructed from kmeans)

    cl_dist <- apply(centers, 1, function(y) sqrt(sum((x-y)^2)))
    return(which.min(cl_dist)[1])
}

```

In [12]: *# Function to generate combination of parameters for gridSearch;
each combination must add to a number ~1. Returns a dataframe,
each row of which is a valid combination.*

```

# I re-factored this ftn using R's expand.grid ftn. expand.grid
# actually takes more time to run. This is probably due to
# type-checking. It appears that we also run out of memory more
# quickly when using expand.grid. So at the moment I am
# reverting to the deprecated section.

generate_combs <- function(arglist, tol=0.0001) {
    # arglist is a named list; each name is a column
    # name of the dataframe which goes to k-means

    #####
    # this next section is an alternative to expand.grid
    #####
    # if(FALSE) {
    n_args <- length(arglist)

    param_vlens <- rep(NA, n_args)

```

```

    for(i in 1:n_args) {
      param_vlens[i] <- length(arglist[[i]])
    }
    n_rows <- prod(param_vlens)

    datout <- rep(NA, n_args*n_rows)
    dim(datout) <- c(n_rows, n_args)
    datout <- as.data.frame(datout)
    colnames(datout) <- names(arglist)

    cprod <- 1
    for(j in 1:n_args) {
      vect <- arglist[[j]]
      val <- rep(vect, rep(cprod, length(vect)))
      datout[, j] <- rep(val, n_rows/length(val))
      cprod <- cprod*length(vect)
    }
    # } ## end of 'if(FALSE)'
    #####

    # datout <- expand.grid(arglist, KEEP.OUT.ATTRS= FALSE)
    # colnames(datout) <- names(arglist)

    row_sums <- round(rowSums(datout), 4)
    names(row_sums) <- rownames(datout)
    tol <- tol
    row_sums <- row_sums[which((as.numeric(row_sums) <= (1 + tol)) & (as.numeric(row_sums) > 0))]
    datout <- datout[names(row_sums),]
    return(datout)
  }

```

In [13]: # Function to constrain range of data between 0 and 1.

```
range01 <- function(x) {(x - min(x))/(max(x) - min(x))}
```

In [11]: # Function for obtaining average of confusion matrix
f-score and percent correctly answered. This function
is called from gridSearch06.

```

get_cvScore_kmp1 <- function(traindat, valdat, wghts) {

  # Apply pca to traindat (4 components).
  pca <- prcomp(traindat[, km_predictors], center=TRUE, scale.=TRUE,
               rank.=4, retx=TRUE)

  # Apply min-max scaling.
  traindat_scaled <- apply(pca$x, MARGIN=2, range01)
  traindat_scaled <- as.data.frame(traindat_scaled)
  colnames(traindat_scaled) <- paste0("pc", 1:4)

  # Construct random forest model.
  set.seed(1493)
  rfclf <- randomForest(I(as.factor(Inland)) ~ .,
                       data= traindat[, rfclf_columns],
                       ntree=900, mtry= 3, nodesize= 1,
                       importance= TRUE)

  # Add prob01 column.
  preds_val <- predict(rfclf, newdata= traindat[, rfclf_columns], type="prob")
  traindat_scaled$prob01 <- preds_val[, 2]

  # Apply weights. The sqrt should have been taken in the
  # calling function.
  cols <- names(wghts)
  traindat_wghts <- t(t(traindat_scaled[, cols]) * as.numeric(wghts[cols]))

  #####
  # Prepare valdat.

```

```

# Apply pca.
tmpdat <- predict(pca, valdat[, km_predictors])

# Apply min-max scaling.
valdat_scaled <- apply(tmpdat, MARGIN=2, range01)
valdat_scaled <- as.data.frame(valdat_scaled)
colnames(valdat_scaled) <- paste0("pc", 1:4)

# Add prob01 column.
preds_val <- predict(rfclf, newdata= valdat[, rfclf_columns], type="prob")
valdat_scaled$prob01 <- preds_val[, 2]

# Apply weights to valdat.
valdat_wghts <- t(t(valdat_scaled[, cols]) * as.numeric(wghts[cols]))

#####
# Construct k-means model on traindat.

kmod <- suppressWarnings(kmeans(traindat_wghts, 2, iter.max = 50, nstart=5))

# See how the clusters are associated with Inland.
kmtrain_Inland_percent <- mean(traindat$Inland)
# Note that traindat, and not traindat_wghts, is being used
# in the following statement. The order of the rows needs to
# remain the same among the different versions of traindat.
dfout <- as.data.frame(cbind(traindat$Inland, kmod$cluster))
colnames(dfout) <- c("Inland", "cluster")
rownames(dfout) <- rownames(traindat)
dat_c1 <- dfout[which(dfout$cluster== 1),]
datc1_Inland_percent <- mean(dat_c1$Inland)
c1_to_InlandYES <- FALSE
if(datc1_Inland_percent >= kmtrain_Inland_percent) { c1_to_InlandYES <- TRUE }

#####
# Apply the k-means model to valdat_wghts.

# Each element of the following list is a row of valdat_wghts.
valdat_asList <- split(valdat_wghts[, colnames(kmod$centers)],
                      seq(nrow(valdat_wghts)))

ctr_list <- vector("list", length= nrow(valdat))
for(i in 1:nrow(valdat)) {
  ctr_list[[i]] <- kmod$centers
}
names(ctr_list) <- rownames(valdat)

# Get the predictions for the validation set.
preds <- mcmapply(getCluster, valdat_asList, ctr_list,
                  SIMPLIFY=TRUE, mc.cores=6)
valdat$cluster <- as.numeric(preds)

valdat$pred_Inland <- NA
if(c1_to_InlandYES) {
  valdat[which(valdat$cluster==1),]$pred_Inland <- 1
  valdat[which(valdat$cluster==2),]$pred_Inland <- 0
} else {
  valdat[which(valdat$cluster==1),]$pred_Inland <- 0
  valdat[which(valdat$cluster==2),]$pred_Inland <- 1
}

# Generate confusion matrix for the k-means clusters and
# the corresponding f-score.
preds <- as.factor(valdat$pred_Inland)
names(preds) <- rownames(valdat)
ans <- get_confusion(preds, valdat[, "Inland", drop=FALSE])

# The result returned is a Type2 score (which is a mixture
# of accuracy and f-score).

```

```

mat <- as.matrix(ans[[1]])
percent_correct <- sum(diag(mat))/floor(sum(mat))
result <- round((0.4 * percent_correct + 0.6 * ans[[2]]), 6)
return(result)
}

```

In [12]: *# This grid search searches for the best set of weights to use
in our k-means clustering model. The best weights are those
which generalize best to the validation set.*

```

gridSearch06 <- function(seed_vector, dat, df_params, folds=5) {

  datout <- rep(NA, 2*nrow(df_params))
  dim(datout) <- c(nrow(df_params), 2)
  datout <- as.data.frame(datout)
  colnames(datout) <- c("row", "Type2")
  datout$row <- rownames(df_params)

  # We want the sqrt of the weights.
  df_params <- df_params^0.5

  params_rows <- rownames(df_params)

  #####
  # Partition the data into folds.

  # divide dat by the number of folds
  segment_size <- round(nrow(dat)/folds)
  diff <- nrow(dat) - folds * segment_size
  last_seg_size <- segment_size + diff
  segmentsv <- c(rep(segment_size, (folds - 1)), last_seg_size)
  stopifnot(sum(segmentsv) == nrow(dat))

  # Create a dataframe, each row for a distinct seed.
  # Each column of the dataframe is for a distinct set
  # of weights. The entries in the cells are Type2
  # scores.
  seedv_len <- length(seed_vector)
  df_scores <- rep(NA, seedv_len*nrow(df_params))
  dim(df_scores) <- c(seedv_len, nrow(df_params))
  df_scores <- as.data.frame(df_scores)
  colnames(df_scores) <- rownames(df_params)
  rownames(df_scores) <- as.character(seed_vector)

  for(h in 1:seedv_len) {
    # shuffle dat
    cur_seed <- seed_vector[h]
    set.seed(cur_seed)
    smp <- sample(rownames(dat), dim(dat)[1], replace=FALSE)
    dat <- dat[smp,]

    # Each element of row_list will be the rows we pick
    # out for one of the folds. E.g., the first element
    # of row_list will contain the rows we want for the
    # first fold, the second element of row_list will
    # contain the rows we want for the second fold, and
    # so forth.
    row_list <- vector("list", length=folds)
    names(row_list) <- as.character(1:folds)
    startpt <- 1
    for(k in 1:folds) {
      endpt <- startpt + segmentsv[k] - 1
      stopifnot(endpt <= nrow(dat))
      row_list[[k]] <- rownames(dat)[startpt:endpt]
      startpt <- endpt + 1
    }

    for(i in 1:nrow(df_params)) {
      cur_row <- params_rows[i]

```



```

      wghts <- as.numeric(df_params[i,])
      names(wghts) <- colnames(df_params)

      train_list <- test_list <- vector("list", length= folds)
      for(j in 1:folds) {
        testdat <- dat[row_list[[j]],]
        traindat <- dat[which(!(rownames(dat) %in% rownames(testdat))),]
        stopifnot((length(rownames(traindat)) + length(rownames(testdat))) == nrow(dat))
        test_list[[j]] <- testdat
        train_list[[j]] <- traindat
      }
      # When there are only 5 folds, only 5 cores get used.
      ### NOTE: I change the following function call depending on
      ### the model I am scoring.
      scores <- mcmapapply(get_cvScore_kmpl, train_list, test_list,
                          MoreArgs= list(wghts=wghts),
                          SIMPLIFY= TRUE, mc.cores=5)

      # For the current seed, store the average of the Type2
      # scores, the average taken over the folds.
      df_scores[as.character(cur_seed), cur_row] <- round(mean(scores), 5)

    } # end of for-loop, index i
  } ## end of for-loop, index h

  # Compute the average over the seeds of the Type2 scores
  # obtained for each set of parameters in df_params.
  datout$Type2 <- round(apply(df_scores, MARGIN=2, mean), 5)
  return(datout)
}

```

Section 1: Get best models for traindat

When looking for good models, I am interested in also reducing the number of predictors, assuming it makes sense to do so. The more predictors used for the hybrid k-means model, the more difficult it will be to find a good set of weights.

Find best random forest model

Variable selection for the random forest model

```

In [50]: set.seed(123)
         rfclf <- randomForest(I(as.factor(Inland)) ~ ., data= traindat,
                               ntree= 900, mtry= 3, nodesize= 2,
                               importance=TRUE)

         print(rfclf)
         print(get_fscore(as.matrix(rfclf$confusion)))
         # [1] 0.8316

```

Call:

```
randomForest(formula = I(as.factor(Inland)) ~ ., data = traindat, ntree = 900, mtry = 3, nodesize = 2, importance = TRUE)
```

Type of random forest: classification

Number of trees: 900

No. of variables tried at each split: 3

OOB estimate of error rate: 10.38%

Confusion matrix:

```

      0      1 class.error
0 6526 421      0.060602
1 637 2612      0.196060

```

"

```
[1] 0.8316
```

```
In [51]: print(round(rfclf$importance, 3))
```

	0	1	MeanDecreaseAccuracy	MeanDecreaseGini
housing_median_age	0.012	0.035	0.019	255.78
total_rooms	0.022	0.023	0.022	152.88
total_bedrooms	0.021	0.016	0.019	138.42
population	0.016	0.016	0.016	146.76
households	0.022	0.019	0.021	140.07
median_income	0.042	0.030	0.038	344.15
median_house_value	0.136	0.246	0.171	1499.71
HHdens_ln	0.032	0.127	0.063	795.74
rooms_per_hh	0.031	0.051	0.037	353.05
bdrms_per_room	0.031	0.039	0.033	292.52
pop_per_hh	0.013	0.017	0.014	233.32

```
In [52]: # We can model nearly as well without households.
```

```
newcols <- colnames(traindat)[which(!(colnames(traindat) %in%
                                     c("households")))]
df <- traindat[, newcols]
```

```
In [53]: set.seed(123)
rfclf <- randomForest(I(as.factor(Inland)) ~ ., data= df,
                      ntree= 900, mtry= 3, nodesize= 2,
                      importance=TRUE)
print(rfclf)
''
print(get_fscore(as.matrix(rfclf$confusion)))
# [1] 0.8323
```

Call:

```
randomForest(formula = I(as.factor(Inland)) ~ ., data = df, ntree = 900, mtry = 3, n
odesize = 2, importance = TRUE)
```

Type of random forest: classification

Number of trees: 900

No. of variables tried at each split: 3

OOB estimate of error rate: 10.33%

Confusion matrix:

	0	1	class.error
0	6530	417	0.060026
1	636	2613	0.195753

"

```
[1] 0.8323
```

```
In [54]: print(round(rfclf$importance, 3))
```

	0	1	MeanDecreaseAccuracy	MeanDecreaseGini
housing_median_age	0.012	0.035	0.019	261.66
total_rooms	0.020	0.018	0.019	167.49
total_bedrooms	0.016	0.015	0.016	154.71
population	0.015	0.013	0.014	159.34
median_income	0.043	0.031	0.039	358.66
median_house_value	0.140	0.248	0.175	1516.44
HHdens_ln	0.032	0.131	0.064	822.02
rooms_per_hh	0.030	0.052	0.037	369.25
bdrms_per_room	0.031	0.038	0.033	299.12
pop_per_hh	0.014	0.017	0.015	245.84

```
In [55]: # We can also remove total_bedrooms.
```

```
newcols <- colnames(traindat)[which(!(colnames(traindat) %in%
```

```

                                c("households", "total_bedrooms")))]
df <- traindat[, newcols]

set.seed(123)
rfclf <- randomForest(I(as.factor(Inland)) ~ ., data= df,
                      ntree= 900, mtry= 3, nodesize= 2,
                      importance=TRUE)

print(rfclf)
''
print(get_fscore(as.matrix(rfclf$confusion)))
# [1] 0.8313

```

Call:

```
randomForest(formula = I(as.factor(Inland)) ~ ., data = df, ntree = 900, mtry = 3, n
odesize = 2, importance = TRUE)
Type of random forest: classification
Number of trees: 900
No. of variables tried at each split: 3

OOB estimate of error rate: 10.41%
Confusion matrix:
      0      1 class.error
0 6520  427   0.061465
1  634 2615   0.195137
"
```

[1] 0.8313

In [56]: `print(round(rfclf$importance, 3))`

	0	1	MeanDecreaseAccuracy	MeanDecreaseGini
housing_median_age	0.012	0.036	0.020	271.47
total_rooms	0.015	0.014	0.015	193.91
population	0.011	0.012	0.012	179.60
median_income	0.041	0.030	0.037	347.53
median_house_value	0.141	0.254	0.177	1566.67
HHdens_ln	0.032	0.134	0.065	856.16
rooms_per_hh	0.029	0.053	0.037	379.26
bdrms_per_room	0.030	0.038	0.033	300.42
pop_per_hh	0.014	0.018	0.015	261.03

In [57]: `# Let's also remove population.`

```

newcols <- colnames(traindat)[which(!(colnames(traindat) %in%
                                c("households", "total_bedrooms",
                                "population")))]

df <- traindat[, newcols]

set.seed(123)
rfclf <- randomForest(I(as.factor(Inland)) ~ ., data= df,
                      ntree= 900, mtry= 3, nodesize= 2,
                      importance=TRUE)

print(rfclf)
''
print(get_fscore(as.matrix(rfclf$confusion)))
# [1] 0.8332

```

Call:

```
"
```

```
[1] 0.8332
```

In [58]: `print(round(rfclf$importance, 3))`

	0	1	MeanDecreaseAccuracy	MeanDecreaseGini
housing_median_age	0.013	0.038	0.021	283.54
total_rooms	0.009	0.007	0.008	232.39
median_income	0.043	0.027	0.038	344.08
median_house_value	0.144	0.258	0.180	1607.21
HHdens_ln	0.033	0.138	0.067	886.04
rooms_per_hh	0.030	0.053	0.037	399.14
bdrms_per_room	0.031	0.039	0.033	318.50
pop_per_hh	0.015	0.017	0.016	289.04

In [59]: `# Let's also remove total_rooms.`

```
newcols <- colnames(traindat)[which(!(colnames(traindat) %in%
                                     c("households", "total_bedrooms",
                                       "population", "total_rooms")))]

df <- traindat[, newcols]

set.seed(123)
rfclf <- randomForest(I(as.factor(Inland)) ~ ., data= df,
                      ntree= 900, mtry= 3, nodesize= 2,
                      importance=TRUE)

print(rfclf)
''
print(get_fscore(as.matrix(rfclf$confusion)))
# [1] 0.8310
```

Call:

```
randomForest(formula = I(as.factor(Inland)) ~ ., data = df, ntree = 900, mtry = 3, n
odesize = 2, importance = TRUE)
```

Type of random forest: classification

Number of trees: 900

No. of variables tried at each split: 3

OOB estimate of error rate: 10.43%

Confusion matrix:

	0	1	class.error
0	6519	428	0.061609
1	635	2614	0.195445

"

```
[1] 0.831
```

In [60]: `# We are now left with 7 predictors.`

```
print(round(rfclf$importance, 3))
```

	0	1	MeanDecreaseAccuracy	MeanDecreaseGini
housing_median_age	0.015	0.037	0.022	305.28
median_income	0.043	0.026	0.038	345.36
median_house_value	0.144	0.262	0.182	1674.29
HHdens_ln	0.034	0.139	0.067	936.38
rooms_per_hh	0.034	0.054	0.040	436.17
bdrms_per_room	0.031	0.041	0.034	336.19
pop_per_hh	0.017	0.019	0.017	326.43

In []: `#&* Bookmark`

```
In [14]: # Identify the columns we will use for random forest modeling.

rfclf_columns <- c('Inland', 'housing_median_age', 'median_income',
                  'median_house_value', 'HHdens_ln', 'rooms_per_hh',
                  'bdrms_per_room', 'pop_per_hh')
```

```
In [83]: # Restrict variables in traindat and testdat for random forest
# modeling.

rftrain <- traindat[, rfclf_columns]
rftest <- testdat[, rfclf_columns]
```

Tune the random forest model

```
In [63]: # This function is called from get_cvScore_rf.

get_Type2_rf <- function(traindat, valdat, ntrees, mtry, nodesize) {

  rfmod <- randomForest(I(as.factor(Inland)) ~ .,
                        data= traindat, ntree= ntrees,
                        mtry= mtry, nodesize= nodesize)

  preds <- predict(rfmod, newdata= valdat, type="response")
  ans <- get_confusion(preds, valdat[, "Inland", drop=FALSE])

  # Type2 score is average of accuracy and f-score.
  mat <- as.matrix(ans[[1]])
  percent_correct <- sum(diag(mat))/floor(sum(mat))
  result <- round((0.4 * percent_correct + 0.6 * ans[[2]]), 4)
  return(result)
}
```

```
In [64]: # Function to obtain a cross-validation score, averaging the
# Type2 scores of the folds. This function is called from
# avg_seedScores_rf.

get_cvScore_rf <- function(seed, dat, ntrees, mtry,
                           nodesize, folds= 5) {

  # divide dat by the number of folds
  segment_size <- round(dim(dat)[1]/folds)
  diff <- nrow(dat) - folds * segment_size
  last_seg_size <- segment_size + diff
  segmentsv <- c(rep(segment_size, (folds - 1)), last_seg_size)
  # print(segmentsv)
  stopifnot(sum(segmentsv) == nrow(dat))

  # shuffle dat
  set.seed(seed)
  smp <- sample(rownames(dat), nrow(dat), replace= FALSE)
  dat <- dat[smp,]

  # split the data into the folds
  row_list <- vector("list", length= folds)
  names(row_list) <- as.character(1:folds)
  startpt <- 1
  for(i in 1:folds) {
    endpt <- startpt + segmentsv[i] - 1
    stopifnot(endpt <= dim(dat)[1])
    row_list[[i]] <- rownames(dat)[startpt:endpt]
    startpt <- endpt + 1
  }

  train_list <- test_list <- vector("list", length= folds)
  for(j in 1:folds) {
    testdat <- dat[row_list[[j]],]
```

```

    traindat <- dat[which(!(rownames(dat) %in% rownames(testdat))),]
    stopifnot((length(rownames(traindat)) + length(rownames(testdat))) == dim(dat)[1])
    test_list[[j]] <- testdat
    train_list[[j]] <- traindat
  }

  scores <- mcmapply(get_Type2_rf, train_list, test_list,
                    MoreArgs= list(ntrees= ntrees, mtry=mtry,
                                   nodesize=nodesize),
                    SIMPLIFY= TRUE, mc.cores=5)

  # The average is of Type2 scores.
  return(round(mean(scores), 5))
}

```

In [65]: *# Since the seed value can have a big effect on the results,
I take the average over a number of seeds. This ftn is
called from gridSearch02.*

```

avg_seedScores_rf <- function(seed_vector, traindat, n_trees,
                              mtry, nodesize, folds= 5) {

  seed_len <- length(seed_vector)
  outv <- rep(NA, seed_len)
  for(i in 1:seed_len) {
    seed <- seed_vector[i]
    outv[i] <- get_cvScore_rf(seed, traindat, n_trees,
                              mtry, nodesize, folds=folds)
  }
  return(round(mean(outv), 5))
}

```

In [66]: *# This grid search is specific to finding the best random forest
classifier for traindat.*

```

gridSearch02 <- function(seed_vector, traindat, ntree_vector,
                          mtry_vector, nodesizes, folds=5) {

  tree_len <- length(ntree_vector)
  mtry_len <- length(mtry_vector)
  node_len <- length(nodesizes)
  # We need to capture the gridSearch parameters as well as
  # the cross-val scores.
  datout <- rep(NA, 2 * tree_len * mtry_len * node_len)
  dim(datout) <- c((tree_len * mtry_len * node_len), 2)
  datout <- as.data.frame(datout)
  colnames(datout) <- c("params", "Type2")
  datout$params <- ""

  index <- 0
  for(i in 1:tree_len) {
    n_trees <- ntree_vector[i]
    for(j in 1:mtry_len) {
      mtry <- mtry_vector[j]
      for(k in 1:node_len) {
        index <- index + 1
        nodesize <- nodesizes[k]
        param_string <- paste(as.character(n_trees),
                              as.character(mtry),
                              as.character(nodesize), sep= "--")
        datout$params[index] <- param_string
        datout$Type2[index] <- avg_seedScores_rf(seed_vector, traindat, n_trees,
                                                  folds=folds, mtry=mtry,
                                                  nodesize=nodesize)
      }
    }
  }
  return(datout)
}

```

```

In [67]: # Run grid search to get better parameters for the
# random forest model. Test with 10 seeds. For each
# seed, an average is taken over 5 folds.

set.seed(7543)
seed_smp <- sample(1:9999, 10, replace=FALSE)
tree_vector <- c(500, 900, 1200)
mtry_vector <- c(2, 3)
node_vector <- c(1, 2, 3)

start <- Sys.time()
paste("Start time: ", start, sep="")
ans <- gridSearch02(seed_smp, rftrain, tree_vector, mtry_vector, node_vector)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 20 mins

(best_params <- ans[which(ans$Type2 == max(ans$Type2)),]$params)
# '900--3--1'

(best_rf_Type2 <- ans[which(ans$Type2 == max(ans$Type2)),]$Type2)
# 0.85547

'Start time: 2021-04-22 15:17:00'

Time difference of 19.61 mins

'900--3--1'

0.85547

```

```

In [68]: ans

```

A data.frame: 18 × 2

params	Type2
<chr>	<dbl>
500--2--1	0.85411
500--2--2	0.85408
500--2--3	0.85375
500--3--1	0.85527
500--3--2	0.85531
500--3--3	0.85483
900--2--1	0.85448
900--2--2	0.85403
900--2--3	0.85396
900--3--1	0.85547
900--3--2	0.85518
900--3--3	0.85463
1200--2--1	0.85443
1200--2--2	0.85420
1200--2--3	0.85407
1200--3--1	0.85537
1200--3--2	0.85536
1200--3--3	0.85484

Get scores for rfc1f best on traindat

```

In [16]: # Get stable scores for the best random forest model. I will
# refer to this model as rfc1f_best. Note that 1000 seeds
# are being used. [* 150 seeds would have been plenty *]

set.seed(1433)
seed_smp <- sample(1:9999, 1000, replace=FALSE)

datout <- rep(NA, 6 * length(seed_smp))
dim(datout) <- c(length(seed_smp), 6)
datout <- as.data.frame(datout)
colnames(datout) <- c("seed", "fscore", "Acc", "Type2", "FN", "FP")
datout$seed <- seed_smp

start <- Sys.time()
for(i in 1:length(seed_smp)) {


  set.seed(seed_smp[i])
  rfmod <- randomForest(I(as.factor(Inland)) ~ .,
                        data= rftrain, ntree=900,
                        mtry= 3, nodesize= 1)

  # preds <- predict(rfmod, newdata= dat, type="response")
  # ans <- get_confusion(preds, dat[, "Outcome", drop=FALSE])
  # mat <- as.matrix(ans[[1]])
  mat <- rfmod$confusion
  # percent_correct <- sum(diag(mat))/floor(sum(mat))
  # datout[i, c("Acc")] <- round(percent_correct, 4)
  datout[i, c("Acc")] <- acc <- round(1-median(rfmod$err.rate[,1]), 4)
  # datout[i, c("fscore")] <- round(ans[[2]], 4)
  datout[i, c("fscore")] <- fscore <- round(get_fscore(mat), 4)
  datout[i, c("Type2")] <- round(0.4*acc + 0.6*fscore, 4)
  datout[i, c("FN")] <- as.numeric(mat[2,1])
  datout[i, c("FP")] <- as.numeric(mat[1,2])
}
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 1.76 hours

options(repr.plot.width= 10, repr.plot.height= 6)
hist(datout$fscore, breaks=12, xlab="f-score",
     main="Distribution of f-scores for rfc1f_best on traindat")

```

Time difference of 1.76 hours



```

In [17]: # Get summaries for rfc1f_best.

```



```
fn_avg <- mean(datout$FN)
fp_avg <- mean(datout$FP)

c(round(mean(datout$fscore), 4), round(mean(datout$Acc), 4),
  round(fn_avg, 2), round(fp_avg, 2))
# f-score: 0.8322
# accuracy: 0.8954
# false negatives: 613.33
# false positives: 449.24


# Type2: 0.8575

0.8322 · 0.8954 · 613.33 · 449.24
```

```
In [18]: round(mean(datout$Type2), 4)
# 0.8575

0.8575
```

```
In [19]: options(repr.plot.width= 10, repr.plot.height= 6)
hist(datout$Acc, breaks=16, xlab="accuracy",
      main="Distribution of accuracy scores for rfclf_best on traindat")
```



```
In [20]: # Identify seeds with an accuracy score between 0.89525
# and 0.8955. When constructing rfclf_best, I may want
# to use a seed which has an accuracy in the center
# of the distribution of accuracy scores. (It may be that
# rfclf_best will generalize better to the test data if
# I choose a seed for it that is at the midpoint of the
# accuracy distribution.)

rf_candidate_seeds <- datout[which((datout$Acc > 0.89525) & (datout$Acc < 0.8955)),]$seed
length(rf_candidate_seeds)
# 190
head(rf_candidate_seeds)
# 667 2700 381 4104 1493 5802

190

667 · 2700 · 381 · 4104 · 1493 · 5802
```

```
In [22]: # Identify seeds with the highest accuracy scores.

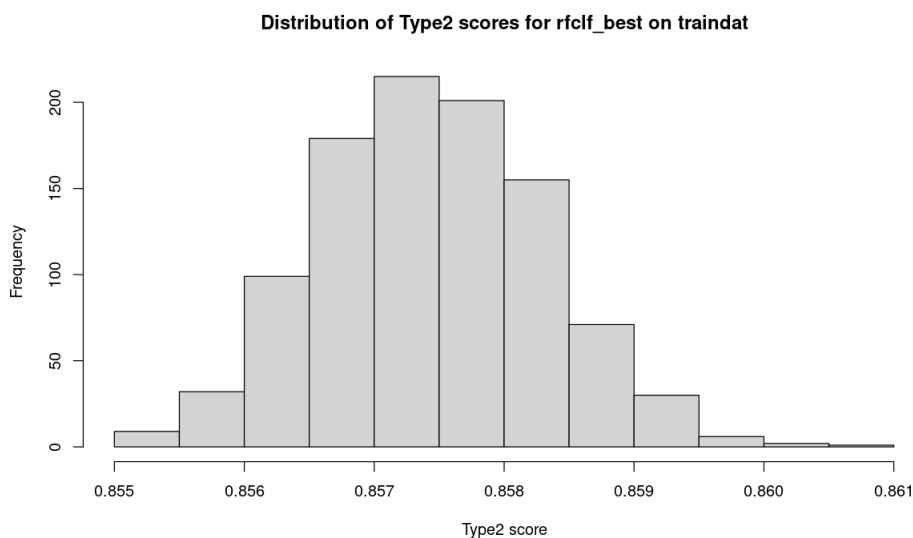
rf_highAcc_seeds <- datout[which(datout$Acc > 0.897),]$seed
length(rf_highAcc_seeds)
```

```
# 5
head(rf_highAcc_seeds)
# 287 9137 9192 2191 8458
```

5

```
287 · 9137 · 9192 · 2191 · 8458
```

```
In [21]: options(repr.plot.width= 10, repr.plot.height= 6)
hist(datout$Type2, breaks=16, xlab="Type2 score",
      main="Distribution of Type2 scores for rfclf_best on traindat")
```



```
In [15]: # Get feature importances for rfclf_best.
```

```
set.seed(1493)
(rfclf_best <- randomForest(I(as.factor(Inland)) ~ .,
                           data= traindat[, rfclf_columns],
                           ntree=900,
                           mtry= 3, nodesize= 1, importance=TRUE))
```

Call:

```
randomForest(formula = I(as.factor(Inland)) ~ ., data = traindat[, rfclf_columns], n
tree = 900, mtry = 3, nodesize = 1, importance = TRUE)
```

Type of random forest: classification

Number of trees: 900

No. of variables tried at each split: 3

OOB estimate of error rate: 10.39%

Confusion matrix:

```
      0      1 class.error
0 6497  450    0.064776
1  609 2640    0.187442
```

```
In [24]: print(round(rfclf_best$importance, 3))
```

	0	1	MeanDecreaseAccuracy	MeanDecreaseGini
housing_median_age	0.015	0.036	0.022	310.43
median_income	0.045	0.025	0.038	367.88
median_house_value	0.146	0.260	0.182	1666.85
HHdens_ln	0.035	0.137	0.067	947.54
rooms_per_hh	0.034	0.051	0.039	447.76
bdrms_per_room	0.033	0.039	0.035	350.40
pop_per_hh	0.017	0.018	0.018	335.81

Get scores for rfc1f_best on testdat

```
In [25]: # Function for obtaining a set of scores on the testset data
# using rfc1f_best as the classifier.

get_testdatScores_rf <- function(seedv, dat) {

  seedv_len <- length(seedv)
  datout <- rep(NA, 5 * seedv_len)
  dim(datout) <- c(seedv_len, 5)
  datout <- as.data.frame(datout)
  colnames(datout) <- c("fscore", "Acc", "Type2", "FN", "FP")
  rownames(datout) <- as.character(seedv)

  for(h in 1:seedv_len) {
    # shuffle dat
    cur_seed <- seedv[h]
    set.seed(cur_seed)
    # It is expected that dat is testdat, which has 10K rcds
    smp <- sample(rownames(dat), 4000, replace= FALSE)
    df <- dat[smp,]

    preds <- predict(rfc1f_best, newdata= df, type="response")
    ans <- get_confusion(preds, df[, "Inland", drop=FALSE])

    mat <- as.matrix(ans[[1]])
    fscore <- round(as.numeric(ans[[2]]), 4)
    acc <- round(sum(diag(mat))/floor(sum(mat)), 4)
    type2 <- round((0.4 * acc + 0.6 * ans[[2]]), 4)
    FN <- as.numeric(mat[2,1])
    FP <- as.numeric(mat[1,2])
    datout[as.character(cur_seed), 1:5] <- c(fscore, acc, type2, FN, FP)
  }
  return(datout)
}
```

```
In [26]: # Get rfc1f_best's scores on the testset data. For each of
# the 250 seeds, I sample 4K (no replacement) from the 10K
# set of testdat records.
```

```
set.seed(1821)
seed_vector <- sample(1:9999, 250, replace=FALSE)

start <- Sys.time()
paste("Start time: ", start, sep="")
dat_result <- get_testdatScores_rf(seed_vector, rf1test)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 1.12 mins
```

'Start time: 2021-04-23 18:19:47'

Time difference of 1.12 mins

```
In [27]: dim(dat_result)
head(dat_result)
```

250 · 5

A data.frame: 6 × 5

	fscore	Acc	Type2	FN	FP
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
5934	0.8386	0.8995	0.8630	233	169

	fscore	Acc	Type2	FN	FP
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1953	0.8443	0.9008	0.8669	224	173
7591	0.8291	0.8915	0.8541	272	162

```
In [28]: fscore_mean <- round(mean(dat_result$fscore), 4)
fscore_sd <- round(sd(dat_result$fscore), 4)
paste0("fscore mean: ", as.character(fscore_mean))
# 0.8366
paste0("fscore StdDev: ", as.character(fscore_sd))
# 0.006
""
summary(dat_result$fscore)
```

```
'fscore mean: 0.8366'
```

```
'fscore StdDev: 0.006'
```

```
"
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	0.820	0.833	0.836	0.837	0.841	0.851

```
In [29]: Acc_mean <- round(mean(dat_result$Acc), 4)
Acc_sd <- round(sd(dat_result$Acc), 4)
paste0("accuracy mean: ", as.character(Acc_mean))
# 0.8979
paste0("accuracy StdDev: ", as.character(Acc_sd))
# 0.0037
""
summary(dat_result$Acc)
```

```
'accuracy mean: 0.8979'
```

```
'accuracy StdDev: 0.0037'
```

```
"
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	0.887	0.895	0.898	0.898	0.900	0.907

```
In [30]: Type2_mean <- round(mean(dat_result$Type2), 4)
Type2_sd <- round(sd(dat_result$Type2), 4)
paste0("Type2 mean: ", as.character(Type2_mean))
# 0.8611
paste0("Type2 StdDev: ", as.character(Type2_sd))
# 0.0049
""
summary(dat_result$Type2)
```

```
'Type2 mean: 0.8611'
```

```
'Type2 StdDev: 0.0049'
```

```
"
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	0.847	0.858	0.861	0.861	0.864	0.873

```
In [31]: FN_mean <- round(mean(dat_result$FN), 2)
FN_sd <- round(sd(dat_result$FN), 2)
paste0("FN mean: ", as.character(FN_mean))
# 242.34
paste0("FN StdDev: ", as.character(FN_sd))
# 11.62
""
summary(dat_result$FN)
```

```
'FN mean: 242.34'
```

```
'FN StdDev: 11.62'
```

```
"
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
207	236	242	242	250	272

```
In [32]: FP_mean <- round(mean(dat_result$FP), 2)
FP_sd <- round(sd(dat_result$FP), 2)
paste0("FP mean: ", as.character(FP_mean))
# 165.95
paste0("FP StdDev: ", as.character(FP_sd))
# 9.76
""
summary(dat_result$FP)
```

```
'FP mean: 165.95'
```

```
'FP StdDev: 9.76'
```

```
"
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
142	160	167	166	173	194

Find best gradient boosting model


Variable selection for gradient boosting model

```
In [16]: options(repr.plot.width= 6, repr.plot.height= 4)

set.seed(123)
gbmod <- gbm(Inland ~ ., data= traindat, n.trees= 1000,
              distribution= "bernoulli", shrinkage= 0.1)
summary(gbmod)
```

A data.frame: 11 × 2

	var	rel.inf
	<chr>	<dbl>
median_house_value	median_house_value	53.22173
HHdens_In	HHdens_In	25.06241
rooms_per_hh	rooms_per_hh	5.16927
housing_median_age	housing_median_age	3.57757
bdrms_per_room	bdrms_per_room	3.13607
pop_per_hh	pop_per_hh	2.75024
total_rooms	total_rooms	1.84534
median_income	median_income	1.62417
population	population	1.47985
total_bedrooms	total_bedrooms	1.13530
households	households	0.99805



```
In [18]: preds <- suppressMessages(predict(gbmod, newdata= traindat, type="response"))
preds_transf <- preds
names(preds_transf) <- rownames(traindat)
preds_transf[which(preds_transf >= 0.5)] <- 1
preds_transf[which(preds_transf < 0.5)] <- 0
preds_transf <- as.factor(preds_transf)
ans <- get_confusion(preds_transf, traindat[, "Inland", drop=FALSE])
print(ans$matrix)
''
print(paste("f-score for gbmod: ", as.character(ans[[2]]), sep=""))
# 0.8372

# Accuracy is 0.8994

# Type2 is 0.8621
```

```
      0      1 class.error
0 6531  416      0.0599
1  610 2639      0.1878
```

"

```
[1] "f-score for gbmod: 0.8372"
```

```
In [19]: # Remove households and total_bedrooms.


newcols <- colnames(traindat)[which(!(colnames(traindat) %in%
                                     c("households", "total_bedrooms")))]
df <- traindat[, newcols]
```

```
In [20]: options(repr.plot.width= 6, repr.plot.height= 4)

set.seed(123)
gbmod <- gbm(Inland ~ ., data= df, n.trees= 1000,
              distribution= "bernoulli", shrinkage= 0.1)
summary(gbmod)
```

A data.frame: 9 × 2

	var	rel.inf
	<chr>	<dbl>
median_house_value	median_house_value	53.4198
HHdens_In	HHdens_In	25.3147
rooms_per_hh	rooms_per_hh	5.2292
housing_median_age	housing_median_age	3.6702
bdrms_per_room	bdrms_per_room	3.1768
pop_per_hh	pop_per_hh	3.0683
total_rooms	total_rooms	2.4803
population	population	1.9066
median_income	median_income	1.7341



```

In [21]: preds <- suppressMessages(predict(gbmmod, newdata= df, type="response"))
preds_transf <- preds
names(preds_transf) <- rownames(df)
preds_transf[which(preds_transf >= 0.5)] <- 1
preds_transf[which(preds_transf < 0.5)] <- 0
preds_transf <- as.factor(preds_transf)
ans <- get_confusion(preds_transf, df[, "Inland", drop=FALSE])
print(ans$matrix)
''

print(paste("f-score for gbmmod: ", as.character(ans[[2]]), sep=""))
# 0.8365

# Accuracy is 0.8989

# Type2 is 0.8615

      0      1 class.error
0 6524  423      0.0609
1  609 2640      0.1874

"

[1] "f-score for gbmmod: 0.8365"

```

```

In [22]: # See what results are when we remove population. We already
# have pop_per_hh.

newcols <- colnames(traindat)[which(!(colnames(traindat) %in%
                                     c("households", "total_bedrooms",
                                       "population")))]

df <- traindat[, newcols]

```




```
In [23]: options(repr.plot.width= 6, repr.plot.height= 4)

set.seed(123)
gbmod <- gbm(Inland ~ ., data= df, n.trees= 1000,
              distribution= "bernoulli", shrinkage= 0.1)
summary(gbmod)
```

A data.frame: 8 × 2

	var	rel.inf
	<chr>	<dbl>
median_house_value	median_house_value	53.7036
HHdens_In	HHdens_In	25.5192
rooms_per_hh	rooms_per_hh	5.3179
housing_median_age	housing_median_age	3.7408
total_rooms	total_rooms	3.5693
bdrms_per_room	bdrms_per_room	3.2124
pop_per_hh	pop_per_hh	3.1650
median_income	median_income	1.7718



```
In [24]: preds <- suppressMessages(predict(gbmod, newdata= df, type="response"))
preds_transf <- preds
names(preds_transf) <- rownames(df)
preds_transf[which(preds_transf >= 0.5)] <- 1
preds_transf[which(preds_transf < 0.5)] <- 0
preds_transf <- as.factor(preds_transf)
ans <- get_confusion(preds_transf, df[, "Inland", drop=FALSE])
print(ans$matrix)
''
print(paste("f-score for gbmod: ", as.character(ans[[2]]), sep=""))
# 0.8346

# Accuracy is 0.8974

# Type2 is 0.8597
```

```
      0      1 class.error
0 6511  436      0.0628
1  610 2639      0.1878

[1] "f-score for gbmod: 0.8346"
```

```
In [25]: # See what results are when we remove median_income
# instead of population.

newcols <- colnames(traindat)[which(!(colnames(traindat) %in%
```

```

c("households", "total_bedrooms",
  "median_income")))]

df <- traindat[, newcols]

```

```


In [26]: options(repr.plot.width= 6, repr.plot.height= 4)

set.seed(123)
gbmod <- gbm(Inland ~ ., data= df, n.trees= 1000,
              distribution= "bernoulli", shrinkage= 0.1)
summary(gbmod)

```

A data.frame: 8 × 2

	var	rel.inf
	<chr>	<dbl>
median_house_value	median_house_value	54.3413
HHdens_In	HHdens_In	25.7256
rooms_per_hh	rooms_per_hh	5.3308
housing_median_age	housing_median_age	3.7384
pop_per_hh	pop_per_hh	3.2143
bdrms_per_room	bdrms_per_room	2.8958
total_rooms	total_rooms	2.6660
population	population	2.0879



```

In [27]: preds <- suppressMessages(predict(gbmod, newdata= df, type="response"))
preds_transf <- preds
names(preds_transf) <- rownames(df)
preds_transf[which(preds_transf >= 0.5)] <- 1
preds_transf[which(preds_transf < 0.5)] <- 0
preds_transf <- as.factor(preds_transf)
ans <- get_confusion(preds_transf, df[, "Inland", drop=FALSE])
print(ans$matrix)
''
print(paste("f-score for gbmod: ", as.character(ans[[2]]), sep=""))
# 0.8355

# Accuracy is 0.8984

# Type2 is 0.8607

```

```

      0      1 class.error
0 6530  417      0.0600
1  619 2630      0.1905

```

"

```
[1] "f-score for gbmod: 0.8355"
```

```
In [17]: # Identify the columns we will use for our gradient boosting
# model. There are 8 predictors.

gbclf_columns <- c('Inland', 'housing_median_age', 'population',
                  'median_house_value', 'HHdens_ln', 'rooms_per_hh',
                  'bdrms_per_room', 'pop_per_hh', 'total_rooms')
```

```
In [34]: # Restrict variables in traindat and testdat for gradient
# boosting modeling.

gbtrain <- traindat[, gbclf_columns]
gbtest <- testdat[, gbclf_columns]
```

Tune the gradient boosting model

```
In [35]: # This function is called from get_cvScore_gb.

get_Type2_gb <- function(traindat, valdat, ntrees, shrinkage) {

  gbmod <- gbm(Inland ~ ., data= traindat, n.trees=ntrees,
               distribution= "bernoulli", shrinkage=shrinkage)

  preds <- suppressMessages(predict(gbmod, newdata= valdat, type="response"))
  names(preds) <- rownames(valdat)
  preds[which(preds >= 0.5)] <- 1
  preds[which(preds < 0.5)] <- 0
  preds <- as.factor(preds)
  ans <- get_confusion(preds, valdat[, "Inland", drop=FALSE])

  # Type2 score is average of accuracy and f-score.
  mat <- as.matrix(ans[[1]])
  percent_correct <- sum(diag(mat))/floor(sum(mat))
  result <- round((0.4 * percent_correct + 0.6 * ans[[2]]), 4)
  return(result)
}
```

```
In [31]: # Function to obtain a cross-validation score, averaging the
# Type2 scores of the folds. This function is called from
# avg_seedScores_gb.

get_cvScore_gb <- function(seed, dat, ntrees, shrinkage,
                           folds= 5) {

  # divide dat by the number of folds
  segment_size <- round(dim(dat)[1]/folds)
  diff <- nrow(dat) - folds * segment_size
  last_seg_size <- segment_size + diff
  segmentsv <- c(rep(segment_size, (folds - 1)), last_seg_size)
  stopifnot(sum(segmentsv) == nrow(dat))

  # shuffle dat
  set.seed(seed)
  smp <- sample(rownames(dat), nrow(dat), replace= FALSE)
  dat <- dat[smp,]

  # split the data into the folds
  row_list <- vector("list", length= folds)
  names(row_list) <- as.character(1:folds)
  startpt <- 1
  for(i in 1:folds) {
    endpt <- startpt + segmentsv[i] - 1
    stopifnot(endpt <= dim(dat)[1])
    row_list[[i]] <- rownames(dat)[startpt:endpt]
    startpt <- endpt + 1
  }
}
```

```

train_list <- test_list <- vector("list", length= folds)
for(j in 1:folds) {
  testdat <- dat[row_list[[j]],]
  traindat <- dat[which(!(rownames(dat) %in% rownames(testdat))),]
  stopifnot((length(rownames(traindat)) + length(rownames(testdat))) == dim(dat)[1])
  test_list[[j]] <- testdat
  train_list[[j]] <- traindat
}

scores <- mcmapply(get_Type2_gb, train_list, test_list,
                  MoreArgs= list(ntrees= ntrees, shrinkage=shrinkage),
                  SIMPLIFY= TRUE, mc.cores=5)

# The average is of Type2 scores.
return(round(mean(scores), 5))
}

```

In [32]: *# Since the seed value can have a big effect on the results,
I take the average over a number of seeds. This ftn is
called from gridSearch03.*

```

avg_seedScores_gb <- function(seed_vector, traindat, n_trees,
                              shrinkage, folds= 5) {

  seed_len <- length(seed_vector)
  outv <- rep(NA, seed_len)
  for(i in 1:seed_len) {
    seed <- seed_vector[i]
    outv[i] <- get_cvScore_gb(seed, traindat, n_trees,
                              shrinkage, folds=folds)
  }
  return(round(mean(outv), 5))
}

```

In [33]: *# This grid search is specific to finding the best random forest
classifier for traindat.*

```

gridSearch03 <- function(seed_vector, traindat, ntree_vector,
                          shrinkagev, folds=5) {

  tree_len <- length(ntree_vector)
  shrink_len <- length(shrinkagev)
  # We need to capture the gridSearch parameters as well as
  # the cross-val scores.
  datout <- rep(NA, 2 * tree_len * shrink_len)
  dim(datout) <- c((tree_len * shrink_len), 2)
  datout <- as.data.frame(datout)
  colnames(datout) <- c("params", "Type2")
  datout$params <- ""

  index <- 0
  for(i in 1:tree_len) {
    n_trees <- ntree_vector[i]
    for(j in 1:shrink_len) {
      shrinkage <- shrinkagev[j]
      index <- index + 1
      param_string <- paste(as.character(n_trees),
                            as.character(shrinkage), sep= "--")
      datout$params[index] <- param_string
      datout$Type2[index] <- avg_seedScores_gb(seed_vector, traindat, n_trees,
                                                folds=folds, shrinkage=shrinkage)
    }
  }
  return(datout)
}

```

In [35]: *# Run grid search to get better parameters for the
random forest model. Test with 21 seeds. For each*

```
# seed, an average is taken over 5 folds.

set.seed(7543)
seed_smp <- sample(1:9999, 21, replace=FALSE)
tree_vector <- c(750, 900, 1000, 1200)
shrinkage_v <- c(0.05, 0.08, 0.1, 0.15, 0.2)

start <- Sys.time()
paste("Start time: ", start, sep="")
ans <- gridSearch03(seed_smp, gbtrain, tree_vector, shrinkage_v)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 11 mins

(best_params <- ans[which(ans$Type2 == max(ans$Type2)),]$params)
# '750--0.1'

(best_rf_Type2 <- ans[which(ans$Type2 == max(ans$Type2)),]$Type2)
# 0.8477

'Start time: 2021-04-23 09:10:16'

Time difference of 11.05 mins

'750--0.1'

0.84767
```

In [36]: # Refine the search.

```
set.seed(7541)
seed_smp <- sample(1:9999, 21, replace=FALSE)
tree_vector <- c(400, 600, 750)
shrinkage_v <- c(0.08, 0.1, 0.12)

start <- Sys.time()
paste("Start time: ", start, sep="")
ans <- gridSearch03(seed_smp, gbtrain, tree_vector, shrinkage_v)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 3 mins

(best_params <- ans[which(ans$Type2 == max(ans$Type2)),]$params)
# '750--0.1'

(best_rf_Type2 <- ans[which(ans$Type2 == max(ans$Type2)),]$Type2)
# 0.8476

'Start time: 2021-04-23 09:24:16'

Time difference of 3 mins

'750--0.1'

0.84761
```

In [37]: ans

A data.frame: 9 × 2

params	Type2
<chr>	<dbl>
400--0.08	0.84525
400--0.1	0.84582
400--0.12	0.84678

params	Type2
<chr>	<dbl>
600--0.08	0.84696
600--0.1	0.84738
600--0.12	0.84715
750--0.08	0.84723

Get scores for gbc1f_best on traindat

```

In [36]: # Get stable scores for the best gradient boosting model.
# I will refer to this model as gbclf_best. Note that
# 250 seeds are being used.

set.seed(1433)
seed_smp <- sample(1:9999, 250, replace=FALSE)

datout <- rep(NA, 6 * length(seed_smp))
dim(datout) <- c(length(seed_smp), 6)
datout <- as.data.frame(datout)
colnames(datout) <- c("seed", "fscore", "Acc", "Type2", "FN", "FP")
datout$seed <- seed_smp

start <- Sys.time()
for(i in 1:length(seed_smp)) {

  set.seed(seed_smp[i])
  gbmod <- gbm(Inland ~ ., data= gbtrain, n.trees=750,
               distribution= "bernoulli", shrinkage=0.1)

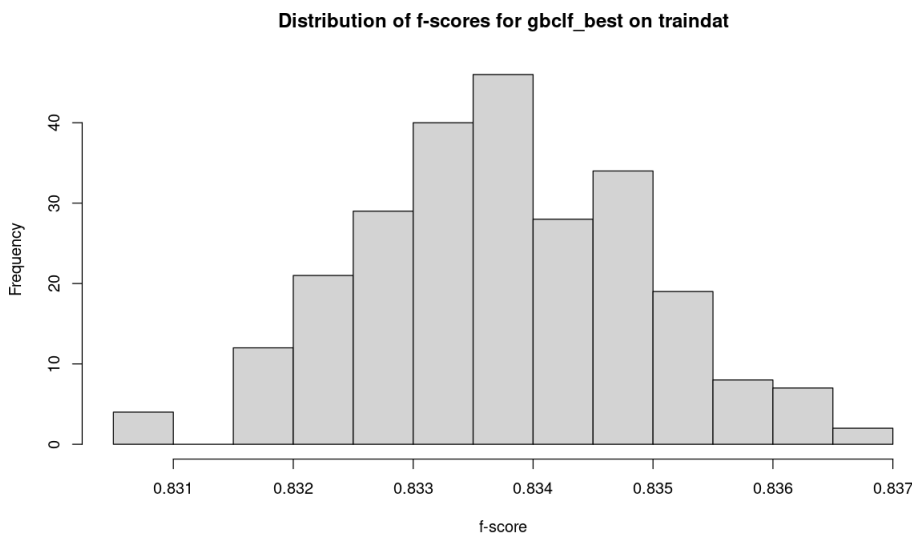
  preds <- suppressMessages(predict(gbmod, newdata= gbtrain, type="response"))
  names(preds) <- rownames(gbtrain)
  preds[which(preds >= 0.5)] <- 1
  preds[which(preds < 0.5)] <- 0
  preds <- as.factor(preds)
  ans <- get_confusion(preds, gbtrain[, "Inland", drop=FALSE])

  mat <- as.matrix(ans[[1]])
  percent_correct <- sum(diag(mat))/floor(sum(mat))
  datout[i, c("Acc")] <- acc <- round(percent_correct, 4)
  datout[i, c("fscore")] <- fscore <- round(ans[[2]], 4)
  datout[i, c("Type2")] <- round(0.4*acc + 0.6*fscore, 4)
  datout[i, c("FN")] <- as.numeric(mat[2,1])
  datout[i, c("FP")] <- as.numeric(mat[1,2])
}
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 5.19 mins

options(repr.plot.width= 10, repr.plot.height= 6)
hist(datout$fscore, breaks=12, xlab="f-score",
     main="Distribution of f-scores for gbclf_best on traindat")

```

Time difference of 5.19 mins



```

In [37]: # Get summaries for gbclf_best.

```

```
fn_avg <- mean(datout$FN)
fp_avg <- mean(datout$FP)

c(round(mean(datout$fscore), 4), round(mean(datout$Acc), 4),
  round(fn_avg, 2), round(fp_avg, 2))
# f-score: 0.8338
# accuracy: 0.8974
# false negatives: 625.56
# false positives: 420.32


# Type2: 0.8592
```

```
0.8338 · 0.8974 · 625.56 · 420.32
```

```
In [38]: round(mean(datout$Type2), 4)
# 0.8592
```

```
0.8592
```

```
In [39]: options(repr.plot.width=10, repr.plot.height=6)
hist(datout$Acc, breaks=16, xlab="accuracy",
      main="Distribution of accuracy scores for gbclf_best on traindat")
```



```
In [40]: # Identify seeds with an accuracy score between 0.8974
# and 0.8976. When constructing gbclf_best, I may want
# to use a seed which has an accuracy in the center
# of the distribution of accuracy scores. (It may be that
# gbclf_best will generalize better to the test data if
# I choose a seed for it that is at the midpoint of the
# accuracy distribution.)

gb_candidate_seeds <- datout[which((datout$Acc > 0.8974) & (datout$Acc < 0.8976)),]$seed
length(gb_candidate_seeds)
# 16
head(gb_candidate_seeds)
# 4621 9214 5711 7584 6932 2254
```

```
16
```

```
4621 · 9214 · 5711 · 7584 · 6932 · 2254
```

```
In [42]: # Identify seeds with the highest accuracy scores.

gb_highAcc_seeds <- datout[which(datout$Acc > 0.8985),]$seed
length(gb_highAcc_seeds)
# 12
```

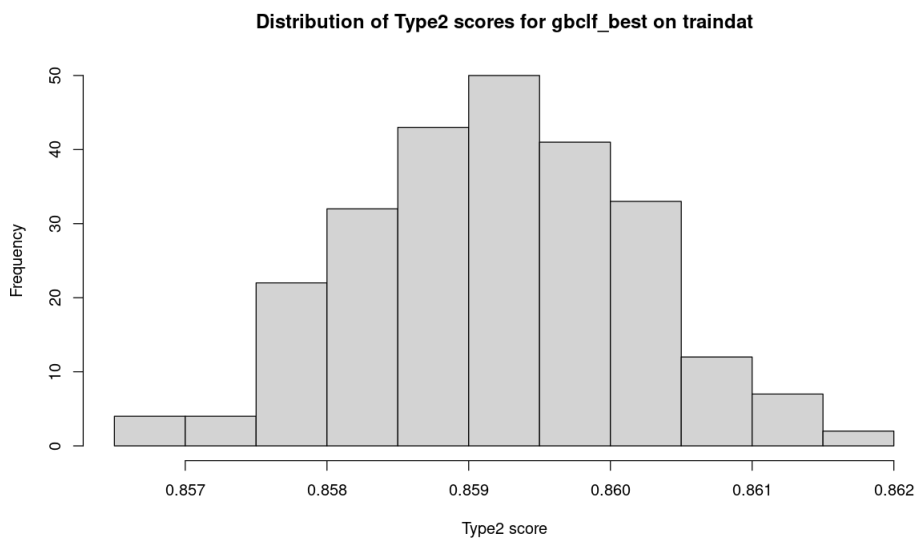


```
head.gb_highAcc_seeds)
# 6937 5021 64 2650 3393 3057
```

12

6937 · 5021 · 64 · 2650 · 3393 · 3057

```
In [43]: options(repr.plot.width= 10, repr.plot.height= 6)
hist(datout$Type2, breaks=16, xlab="Type2 score",
     main="Distribution of Type2 scores for gbclf_best on traindat")
```




```
In [44]: # Get feature importances for gbclf_best.

set.seed(5711)
gbclf_best <- gbm(Inland ~ ., data= gbtrain, n.trees=750,
                  distribution= "bernoulli", shrinkage=0.1)

options(repr.plot.width= 6, repr.plot.height= 4)
summary(gbclf_best)
```

A data.frame: 8 × 2

	var	rel.inf
	<chr>	<dbl>
median_house_value	median_house_value	55.9842
HHdens_In	HHdens_In	25.8345
rooms_per_hh	rooms_per_hh	5.7162
housing_median_age	housing_median_age	3.3490
pop_per_hh	pop_per_hh	2.8637
total_rooms	total_rooms	2.2532
bdrms_per_room	bdrms_per_room	2.1195
population	population	1.8798



Get scores for gbclf_best on testdat

```
In [45]: # Function for obtaining a set of scores on the testset data
# using gbclf_best as the classifier.

get_testdatScores_gb <- function(seedv, dat) {

  seedv_len <- length(seedv)
  datout <- rep(NA, 5 * seedv_len)
  dim(datout) <- c(seedv_len, 5)
  datout <- as.data.frame(datout)
  colnames(datout) <- c("fscore", "Acc", "Type2", "FN", "FP")
  rownames(datout) <- as.character(seedv)

  for(h in 1:seedv_len) {
    # shuffle dat
    cur_seed <- seedv[h]
    set.seed(cur_seed)
    # It is expected that dat is testdat, which has 10K rcds
    smp <- sample(rownames(dat), 4000, replace= FALSE)
    df <- dat[smp,]

    preds <- suppressMessages(predict(gbclf_best, newdata= df, type="response"))
    names(preds) <- rownames(df)
    preds[which(preds >= 0.5)] <- 1
  }
}
```

```

    preds[which(preds < 0.5)] <- 0
    preds <- as.factor(preds)
    ans <- get_confusion(preds, df[, "Inland", drop=FALSE])

    mat <- as.matrix(ans[[1]])
    fscore <- round(as.numeric(ans[[2]]), 4)
    acc <- round(sum(diag(mat))/floor(sum(mat)), 4)
    type2 <- round((0.4 * acc + 0.6 * ans[[2]]), 4)
    FN <- as.numeric(mat[2,1])
    FP <- as.numeric(mat[1,2])
    datout[as.character(cur_seed), 1:5] <- c(fscore, acc, type2, FN, FP)
  }
  return(datout)
}

```

In [46]: *# Get gbm_best's scores on the testset data. For each of
the 250 seeds, I sample 4K (no replacement) from the 10K
set of testdat records.*

```

set.seed(1821)
seed_vector <- sample(1:9999, 250, replace=FALSE)

start <- Sys.time()
paste("Start time: ", start, sep="")
dat_result <- get_testdatScores_gb(seed_vector, gbtest)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 9.28 secs

```

'Start time: 2021-04-23 18:42:11'

Time difference of 9.28 secs

In [47]: `dim(dat_result)`
`head(dat_result)`

250 × 5

A data.frame: 6 × 5

	fscore	Acc	Type2	FN	FP
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
5934	0.8274	0.8925	0.8534	246	184
1953	0.8268	0.8902	0.8522	252	187
7591	0.8116	0.8810	0.8394	300	176
1038	0.8216	0.8908	0.8493	263	174
49	0.8251	0.8890	0.8507	254	190
3203	0.8216	0.8895	0.8488	261	181

In [48]: `fscore_mean <- round(mean(dat_result$fscore), 4)`
`fscore_sd <- round(sd(dat_result$fscore), 4)`
`paste0("fscore mean: ", as.character(fscore_mean))`
0.8204
`paste0("fscore StdDev: ", as.character(fscore_sd))`
0.0065
`" "`
`summary(dat_result$fscore)`

'fscore mean: 0.8204'

'fscore StdDev: 0.0065'

" "

```

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
      ~~~~~ ~~~~~ ~~~~~ ~~~~~ ~~~~~ ~~~~~
In [49]: Acc_mean <- round(mean(dat_result$Acc), 4)
Acc_sd <- round(sd(dat_result$Acc), 4)
paste0("accuracy mean: ", as.character(Acc_mean))
# 0.8882
paste0("accuracy StdDev: ", as.character(Acc_sd))
# 0.0039
""
summary(dat_result$Acc)

```

'accuracy mean: 0.8882'

'accuracy StdDev: 0.0039'

"

```

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
      ~~~~~ ~~~~~ ~~~~~ ~~~~~ ~~~~~ ~~~~~
0.877   0.886   0.888   0.888   0.891   0.899

```

```

In [50]: Type2_mean <- round(mean(dat_result$Type2), 4)
Type2_sd <- round(sd(dat_result$Type2), 4)
paste0("Type2 mean: ", as.character(Type2_mean))
# 0.8475
paste0("Type2 StdDev: ", as.character(Type2_sd))
# 0.0054
""
summary(dat_result$Type2)

```

'Type2 mean: 0.8475'

'Type2 StdDev: 0.0054'

"

```

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
      ~~~~~ ~~~~~ ~~~~~ ~~~~~ ~~~~~ ~~~~~
0.832   0.844   0.848   0.848   0.851   0.862

```

```

In [51]: FN_mean <- round(mean(dat_result$FN), 2)
FN_sd <- round(sd(dat_result$FN), 2)
paste0("FN mean: ", as.character(FN_mean))
# 265.96
paste0("FN StdDev: ", as.character(FN_sd))
# 11.92
""
summary(dat_result$FN)

```

'FN mean: 265.96'

'FN StdDev: 11.92'

"

```

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
      ~~~~~ ~~~~~ ~~~~~ ~~~~~ ~~~~~ ~~~~~
233      257      266      266      273      302

```

```
In [52]: FP_mean <- round(mean(dat_result$FP), 2)
FP_sd <- round(sd(dat_result$FP), 2)
paste0("FP mean: ", as.character(FP_mean))
# 181.21
paste0("FP StdDev: ", as.character(FP_sd))
# 10.49
""
summary(dat_result$FP)
```

'FP mean: 181.21'

'FP StdDev: 10.49'

"

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
149	174	182	181	188	209

```
In [ ]: ### COMMENT:

# rfclf_best looks to be a slightly better model than
# gbclf_best.
```

Find best logistic regression model

Variable selection

```
In [14]: f01 <- glm(Inland ~ ., data=trainat, family=binomial())
summary(f01)
```

Call:

glm(formula = Inland ~ ., family = binomial(), data = trainat)

Deviance Residuals:

Min	1Q	Median	3Q	Max
-3.148	-0.499	-0.163	0.405	4.859

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	6.07e+00	7.71e-01	7.87	3.6e-15
housing_median_age	-2.78e-02	2.92e-03	-9.54	< 2e-16
total_rooms	1.53e-04	7.64e-05	2.01	0.04447
total_bedrooms	-2.55e-03	6.73e-04	-3.78	0.00016
population	-2.33e-04	1.29e-04	-1.80	0.07125
households	2.72e-03	8.19e-04	3.33	0.00088
median_income	-2.39e-01	4.31e-02	-5.54	2.9e-08
median_house_value	-2.32e-05	7.05e-07	-32.98	< 2e-16
HHdens_ln	-4.05e-01	2.47e-02	-16.41	< 2e-16
rooms_per_hh	7.67e-01	7.67e-02	10.00	< 2e-16
bdrms_per_room	-8.23e+00	1.81e+00	-4.56	5.2e-06
pop_per_hh	-2.86e-01	6.33e-02	-4.52	6.2e-06

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 12762.4 on 10195 degrees of freedom
 Residual deviance: 6725.3 on 10184 degrees of freedom
 AIC: 6749

Number of Fisher Scoring iterations: 6

```
In [15]: # Remove population as a predictor.

newcols <- colnames(trainat)[which(!(colnames(trainat) %in%
                                     c("population")))]
```

```
df <- traindat[, newcols]
```

```
In [16]: f01 <- glm(Inland ~ ., data=df, family=binomial())
summary(f01)
```

Call:

```
glm(formula = Inland ~ ., family = binomial(), data = df)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-3.161	-0.501	-0.163	0.405	4.856

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	6.37e+00	7.54e-01	8.45	< 2e-16
housing_median_age	-2.78e-02	2.92e-03	-9.52	< 2e-16
total_rooms	1.20e-04	7.39e-05	1.63	0.10321
total_bedrooms	-2.26e-03	6.56e-04	-3.44	0.00059
households	1.93e-03	6.94e-04	2.78	0.00537
median_income	-2.38e-01	4.31e-02	-5.53	3.3e-08
median_house_value	-2.32e-05	7.04e-07	-32.94	< 2e-16
HHdens_ln	-4.04e-01	2.47e-02	-16.36	< 2e-16
rooms_per_hh	7.65e-01	7.68e-02	9.97	< 2e-16
bdrms_per_room	-8.50e+00	1.80e+00	-4.71	2.4e-06
pop_per_hh	-3.75e-01	4.15e-02	-9.02	< 2e-16

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 12762.4 on 10195 degrees of freedom
 Residual deviance: 6728.5 on 10185 degrees of freedom
 AIC: 6750

Number of Fisher Scoring iterations: 6

```
In [17]: # Remove total_rooms as a predictor.

newcols <- colnames(traindat)[which(!(colnames(traindat) %in%
                                     c("population", "total_rooms")))]
df <- traindat[, newcols]
```

```
In [18]: f01 <- glm(Inland ~ ., data=df, family=binomial())
summary(f01)
```

```
Call:
glm(formula = Inland ~ ., family = binomial(), data = df)
```

```
In [19]: # Remove total_bedrooms.
```

```
newcols <- colnames(traindat)[which(!(colnames(traindat) %in%
                                     c("population", "total_rooms",
                                       "total_bedrooms")))]
df <- traindat[, newcols]
```

```
In [20]: f01 <- glm(Inland ~ ., data=df, family=binomial())
summary(f01)
```

```
Call:
glm(formula = Inland ~ ., family = binomial(), data = df)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-3.237	-0.503	-0.163	0.409	4.832

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	7.56e+00	6.37e-01	11.86	< 2e-16
housing_median_age	-2.74e-02	2.91e-03	-9.41	< 2e-16
households	1.77e-04	8.26e-05	2.14	0.032
median_income	-2.16e-01	4.27e-02	-5.06	4.3e-07
median_house_value	-2.31e-05	7.04e-07	-32.87	< 2e-16
HHdens_ln	-4.07e-01	2.47e-02	-16.48	< 2e-16
rooms_per_hh	6.57e-01	5.80e-02	11.33	< 2e-16
bdrms_per_room	-1.21e+01	1.50e+00	-8.04	9.1e-16
pop_per_hh	-3.57e-01	4.09e-02	-8.74	< 2e-16

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 12762.4 on 10195 degrees of freedom
 Residual deviance: 6739.5 on 10187 degrees of freedom
 AIC: 6758

Number of Fisher Scoring iterations: 6

```
In [21]: # Remove households.
```

```
newcols <- colnames(traindat)[which(!(colnames(traindat) %in%
                                     c("population", "total_rooms",
                                       "total_bedrooms", "households")))]
df <- traindat[, newcols]
```

```
In [23]: f01 <- glm(Inland ~ ., data=df, family=binomial())
summary(f01)
get_RsqrdDev(f01)
```

```
Call:
glm(formula = Inland ~ ., family = binomial(), data = df)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-3.235  -0.505  -0.163   0.407   4.826

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept)    7.68e+00   6.33e-01  12.13  < 2e-16
housing_median_age -2.98e-02  2.69e-03 -11.06  < 2e-16
0.6244
```

```
In [25]: # Add power transformation.

f01 <- glm(Inland ~ housing_median_age + median_income +
            I(log(median_house_value)) +
            HHdens_ln +
            rooms_per_hh +
            bdrms_per_room +
            pop_per_hh,
            data=df, family=binomial())

summary(f01)
get_RsqrdDev(f01)
```

```
Call:
glm(formula = Inland ~ housing_median_age + median_income + I(log(median_house_value)) +
    HHdens_ln + rooms_per_hh + bdrms_per_room + pop_per_hh, family = binomial(),
    data = df)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-3.355  -0.474  -0.176   0.323   3.674

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept)    54.64413    1.39729   39.11  < 2e-16
housing_median_age -0.03526    0.00278  -12.69  < 2e-16
median_income    -0.15859    0.04024   -3.94  8.1e-05
I(log(median_house_value)) -4.12299    0.11195  -36.83  < 2e-16
HHdens_ln       -0.36888    0.02438  -15.13  < 2e-16
rooms_per_hh      0.50661    0.04965   10.20  < 2e-16
bdrms_per_room   -16.69153    1.44595  -11.54  < 2e-16
pop_per_hh       -0.42148    0.04425   -9.52  < 2e-16

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 12762.4  on 10195  degrees of freedom
Residual deviance:  6366.5  on 10188  degrees of freedom
AIC: 6383

Number of Fisher Scoring iterations: 6

0.6526
```

```
In [37]: # Tweak power transformation.

f02 <- glm(Inland ~ housing_median_age + median_income +
            I(median_house_value^(-0.07)) +
            HHdens_ln +
            rooms_per_hh +
            bdrms_per_room +
            pop_per_hh,
            data=df, family=binomial())

summary(f02)
```



```
get_RsqrdDev(f02)
```

Call:

```
glm(formula = Inland ~ housing_median_age + median_income + I(median_house_value^-0.07) +
     HHdens_ln + rooms_per_hh + bdrms_per_room + pop_per_hh, family = binomial(),
     data = df)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-3.372	-0.473	-0.179	0.319	3.621

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-53.06331	1.75378	-30.26	< 2e-16
housing_median_age	-0.03582	0.00278	-12.89	< 2e-16
median_income	-0.16913	0.03980	-4.25	2.1e-05
I(median_house_value^-0.07)	135.29657	3.67098	36.86	< 2e-16
HHdens_ln	-0.36586	0.02432	-15.04	< 2e-16
rooms_per_hh	0.49415	0.04906	10.07	< 2e-16
bdrms_per_room	-17.18793	1.44128	-11.93	< 2e-16
pop_per_hh	-0.42245	0.04443	-9.51	< 2e-16

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 12762.4 on 10195 degrees of freedom
 Residual deviance: 6361.2 on 10188 degrees of freedom
 AIC: 6377

Number of Fisher Scoring iterations: 6

0.653

```
In [62]: # Remove median_income since it is highly correlated
         # with median_house_value.
```

```
f03 <- glm(Inland ~ housing_median_age +
            I(median_house_value^-0.07) +
            HHdens_ln +
            rooms_per_hh +
            bdrms_per_room +
            pop_per_hh,
            data=df, family=binomial())
```

```
summary(f03)
get_RsqrdDev(f03)
```

```
Call:
glm(formula = Inland ~ housing_median_age + I(median_house_value^-0.07) +
    HHdens_ln + rooms_per_hh + bdrms_per_room + pop_per_hh, family = binomial(),
    0.6517
```

f03 model diagnostics

```
In [59]: # Check for dispersion < 1.

# We want the following value to NOT be considerably larger
# than 1. If it is, then we have to use family= quasibinomial
# in our modeling. (See p.311 of Kabacoff's "R in Action".)

phi <- deviance(f03)/df.residual(f03)
print(round(phi, 3))
```

```
[1] 0.626
```


```
In [ ]: ### COMMENT:

# The f03 model has dispersion < 1. This under-
# dispersion does not affect the parameter estimates.
# But it means that the standard errors for our model
# coefficients, as seen in the above summary,
# are larger than what they should be.
```

```
In [63]: # Plot the residuals. Julian Faraway's half-normal plot can
# be used to check for outliers. See p.46 of Faraway's
# "Extending the Linear Model with R". The residuals plotted
# here are the deviance residuals.

options(repr.plot.width= 6, repr.plot.height= 6)

# Function halfnorm is from package faraway.
halfnorm(residuals(f03))
```



```
In [ ]: ### COMMENT:

# The diagnostics look ok for our purposes.
```

Get scores for f03 on testdat

```
In [15]: f03_columns <- colnames(traindat)[which(!(colnames(traindat) %in%
                                                    c("population", "total_rooms",
                                                      "total_bedrooms", "households",
                                                      "median_income")))]
```

```
In [64]: f03train <- traindat[, f03_columns]
f03test <- testdat[, f03_columns]
```

```
In [65]: # Function for obtaining a set of scores on the testset data
# using f03 as the classifier.
```

```
get_testdatScores_f03 <- function(seedv, dat) {

  seedv_len <- length(seedv)
  datout <- rep(NA, 5 * seedv_len)
  dim(datout) <- c(seedv_len, 5)
  datout <- as.data.frame(datout)
  colnames(datout) <- c("fscore", "Acc", "Type2", "FN", "FP")
  rownames(datout) <- as.character(seedv)

  for(h in 1:seedv_len) {
    # shuffle dat
    cur_seed <- seedv[h]
    set.seed(cur_seed)
    # It is expected that dat is testdat, which has 10K rcds
    smp <- sample(rownames(dat), 4000, replace= FALSE)
    df <- dat[smp,]

    preds <- predict(f03, newdata= df)
    names(preds) <- rownames(df)
    preds[which(preds >= 0.5)] <- 1
    preds[which(preds < 0.5)] <- 0
    preds <- as.factor(preds)
    ans <- get_confusion(preds, df[, "Inland", drop=FALSE])

    mat <- as.matrix(ans[[1]])
    fscore <- round(as.numeric(ans[[2]]), 4)
    acc <- round(sum(diag(mat))/floor(sum(mat)), 4)
    type2 <- round((0.4 * acc + 0.6 * ans[[2]]), 4)
    FN <- as.numeric(mat[2,1])
    FP <- as.numeric(mat[1,2])
    datout[as.character(cur_seed), 1:5] <- c(fscore, acc, type2, FN, FP)
  }
  return(datout)
}
```

```
In [66]: # Get f03's scores on the testset data. For each of
# the 250 seeds, I sample 4K (no replacement) from the 10K
# set of testdat records.
```

```
set.seed(1821)
seed_vector <- sample(1:9999, 250, replace=FALSE)

start <- Sys.time()
# paste("Start time: ", start, sep="")
dat_result <- get_testdatScores_f03(seed_vector, f03test)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 5 secs
```

Time difference of 4.99 secs

```
In [67]: dim(dat_result)
```

```
head(dat_result)
```

```
250 · 5
```

A data.frame: 6 × 5

	fscore	Acc	Type2	FN	FP
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
5934	0.7916	0.8805	0.8272	369	109
1953	0.7779	0.8722	0.8156	405	106
7591	0.7639	0.8628	0.8035	437	112
1038	0.7771	0.8745	0.8161	394	108
49	0.7831	0.8732	0.8191	386	121
3203	0.7759	0.8728	0.8147	398	111

```
In [68]: fscore_mean <- round(mean(dat_result$fscore), 4)
fscore_sd <- round(sd(dat_result$fscore), 4)
paste0("fscore mean: ", as.character(fscore_mean))
# 0.7765
paste0("fscore StdDev: ", as.character(fscore_sd))
# 0.0076
""
summary(dat_result$fscore)
```

```
'fscore mean: 0.7765'
```

```
'fscore StdDev: 0.0076'
```

```
"
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	0.760	0.771	0.776	0.776	0.781	0.799

```
In [69]: Acc_mean <- round(mean(dat_result$Acc), 4)
Acc_sd <- round(sd(dat_result$Acc), 4)
paste0("accuracy mean: ", as.character(Acc_mean))
# 0.8725
paste0("accuracy StdDev: ", as.character(Acc_sd))
# 0.0042
""
summary(dat_result$Acc)
```

```
'accuracy mean: 0.8725'
```

```
'accuracy StdDev: 0.0042'
```

```
"
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	0.863	0.870	0.873	0.872	0.875	0.883

```
In [70]: Type2_mean <- round(mean(dat_result$Type2), 4)
Type2_sd <- round(sd(dat_result$Type2), 4)
paste0("Type2 mean: ", as.character(Type2_mean))
# 0.8149
paste0("Type2 StdDev: ", as.character(Type2_sd))
# 0.006
""
summary(dat_result$Type2)
```

```
'Type2 mean: 0.8149'
```

```
'Type2 StdDev: 0.006'
```

```
"
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.801	0.811	0.815	0.815	0.819	0.832

```
In [71]: FN_mean <- round(mean(dat_result$FN), 2)
FN_sd <- round(sd(dat_result$FN), 2)
paste0("FN mean: ", as.character(FN_mean))
# 401.55
paste0("FN StdDev: ", as.character(FN_sd))
# 14.29
""
summary(dat_result$FN)
```

'FN mean: 401.55'

'FN StdDev: 14.29'

"

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
369	393	401	402	411	437

```
In [72]: FP_mean <- round(mean(dat_result$FP), 2)
FP_sd <- round(sd(dat_result$FP), 2)
paste0("FP mean: ", as.character(FP_mean))
# 108.58
paste0("FP StdDev: ", as.character(FP_sd))
# 8.36
""
summary(dat_result$FP)
```

'FP mean: 108.58'

'FP StdDev: 8.36'

"

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
89	102	109	109	114	134

```
In [ ]: ### COMMENT:

# f03 does a great job avoiding false positives.
# But it is very high on the false negatives. It's
# mean accuracy is more than 2.5 percentage points lower
# than the mean accuracy for rfclf_best. On average,
# f03 mis-classifies 200 more districts (of the 10K) than
# does rfclf_best.
```

Add median_income back to the model; get scores on testset

```
In [73]: f04 <- glm(Inland ~ housing_median_age + median_income +
  I(median_house_value^-0.07) +
  HHdens_lh +
  rooms_per_hh +
  bdrms_per_room +
  pop_per_hh,
  data=trainat, family=binomial())
get_RsqrDev(f04)
```

0.653

```
In [74]: # Function for obtaining a set of scores on the testset data
# using f03 as the classifier.

get_testdatScores_f04 <- function(seedv, dat) {
  seedv_len <- length(seedv)
```

```

datout <- rep(NA, 5 * seedv_len)
dim(datout) <- c(seedv_len, 5)
datout <- as.data.frame(datout)
colnames(datout) <- c("fscore", "Acc", "Type2", "FN", "FP")
rownames(datout) <- as.character(seedv)

for(h in 1:seedv_len) {
  # shuffle dat
  cur_seed <- seedv[h]
  set.seed(cur_seed)
  # It is expected that dat is testdat, which has 10K rcds
  smp <- sample(rownames(dat), 4000, replace= FALSE)
  df <- dat[smp,]

  preds <- predict(f04, newdata= df)
  names(preds) <- rownames(df)
  preds[which(preds >= 0.5)] <- 1
  preds[which(preds < 0.5)] <- 0
  preds <- as.factor(preds)
  ans <- get_confusion(preds, df[, "Inland", drop=FALSE])

  mat <- as.matrix(ans[[1]])
  fscore <- round(as.numeric(ans[[2]]), 4)
  acc <- round(sum(diag(mat))/floor(sum(mat)), 4)
  type2 <- round((0.4 * acc + 0.6 * ans[[2]]), 4)
  FN <- as.numeric(mat[2,1])
  FP <- as.numeric(mat[1,2])
  datout[as.character(cur_seed), 1:5] <- c(fscore, acc, type2, FN, FP)
}
return(datout)
}

```

In [75]: # Get f04's scores on the testset data. For each of
 # the 250 seeds, I sample 4K (no replacement) from the 10K
 # set of testdat records.

```

set.seed(1821)
seed_vector <- sample(1:9999, 250, replace=FALSE)

start <- Sys.time()
# paste("Start time: ", start, sep="")
dat_result <- get_testdatScores_f04(seed_vector, testdat)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 5 secs

```

Time difference of 4.98 secs

In [76]: dim(dat_result)
 head(dat_result)

250 · 5

A data.frame: 6 × 5

	fscore	Acc	Type2	FN	FP
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
5934	0.7888	0.8792	0.8250	375	108
1953	0.7732	0.8698	0.8118	412	109
7591	0.7657	0.8635	0.8048	433	113
1038	0.7767	0.8742	0.8157	394	109
49	0.7794	0.8715	0.8162	393	121
3203	0.7717	0.8708	0.8113	405	112

```
In [77]: fscore_mean <- round(mean(dat_result$fscore), 4)
fscore_sd <- round(sd(dat_result$fscore), 4)
paste0("fscore mean: ", as.character(fscore_mean))
# 0.7748
paste0("fscore StdDev: ", as.character(fscore_sd))
# 0.0077
""
summary(dat_result$fscore)
```

'fscore mean: 0.7748'

'fscore StdDev: 0.0077'

"

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	0.757	0.770	0.775	0.775	0.780	0.799

```
In [78]: Acc_mean <- round(mean(dat_result$Acc), 4)
Acc_sd <- round(sd(dat_result$Acc), 4)
paste0("accuracy mean: ", as.character(Acc_mean))
# 0.8717
paste0("accuracy StdDev: ", as.character(Acc_sd))
# 0.0043
""
summary(dat_result$Acc)
```

'accuracy mean: 0.8717'

'accuracy StdDev: 0.0043'

"

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	0.862	0.869	0.872	0.872	0.875	0.883

```
In [ ]: ### COMMENT:

# The scores on the testset are slightly worse when
# we include median_income in the model Thus, I will
# stick with the f03 model.
```

```
In [80]: rm(f04)
```

```
In [87]: colnames(f03train)
```

'Inland' · 'housing_median_age' · 'median_house_value' · 'HHdens_ln' · 'rooms_per_hh' · 'bdrms_per_room' · 'pop_per_hh'

Find best SVM

```
In [84]: colnames(rftrain)
```

'Inland' · 'housing_median_age' · 'median_income' · 'median_house_value' · 'HHdens_ln' · 'rooms_per_hh' · 'bdrms_per_room' · 'pop_per_hh'

```
In [86]: gbclf_columns
```

'Inland' · 'housing_median_age' · 'population' · 'median_house_value' · 'HHdens_ln' · 'rooms_per_hh' · 'bdrms_per_room' · 'pop_per_hh' · 'total_rooms'

```
In [16]: svm02_columns <- f03_columns
```

```
In [88]: # We need a dataframe for the svm modeling. Use only
# the f03 predictors; this will help with computation time.
```

```
svmtrain <- traindat[, colnames(f03train)]
svmtrain$median_house_value <- log(svmtrain$median_house_value)

svm_scaled <- scale(svmtrain[, -1])
svm_centers <- attr(svm_scaled, "scaled:center")
svm_scales <- attr(svm_scaled, "scaled:scale")
svm_scaled <- as.data.frame(cbind(as.numeric(svmtrain$Inland), svm_scaled),
                             row.names=rownames(svmtrain))
colnames(svm_scaled) <- colnames(svmtrain)
```

```
In [89]: # Test out a support vector machine. Try kernel =
# radial basis function.
```

```
svm01 <- svm(I(as.factor(Inland)) ~ ., data=svm_scaled, kernel="radial",
             gamma= 1.0, cost= 700, scale=FALSE)

pred <- fitted(svm01)
(ans <- table(pred, as.factor(svmtrain$Inland)))
get_fscore(as.matrix(ans))
```

```
pred    0    1
0 6904 120
1   43 3129

0.9746
```

```
In [92]: # Prepare the testset data.
```

```
svmtest <- testdat[, colnames(f03train)]
svmtest$median_house_value <- log(svmtest$median_house_value)

svmtest_scaled <- scale(svmtest[, -1], center=svm_centers,
                       scale=svm_scales)
svmtest_scaled <- as.data.frame(cbind(as.numeric(svmtest$Inland), svmtest_scaled),
                                row.names=rownames(svmtest))
colnames(svmtest_scaled) <- colnames(svmtest)
```

```
In [93]: # Function to compute a Type2 score for an svm cv-fold.
```

```
get_Type2_svm <- function(traindat, valdat, gamma, cost) {

  # traindat and valdat need to be scaled
  train_scaled <- scale(traindat[, -1])
  train_centers <- attr(train_scaled, "scaled:center")
  train_scales <- attr(train_scaled, "scaled:scale")
  train_scaled <- as.data.frame(cbind(as.numeric(traindat$Inland), train_scaled),
                                row.names=rownames(traindat))
  colnames(train_scaled) <- colnames(traindat)

  svmmod <- svm(I(as.factor(Inland)) ~ ., data= train_scaled, gamma=gamma,
               cost=cost, scale=FALSE, kernel="radial")

  # Scale valdat.
  valdat_scaled <- scale(valdat[, -1], center=train_centers,
                        scale=train_scales)
  valdat_scaled <- as.data.frame(cbind(as.numeric(valdat$Inland), valdat_scaled),
                                row.names=rownames(valdat))
  colnames(valdat_scaled) <- colnames(valdat)

  preds <- predict(svmmod, newdata= valdat_scaled)
  ans <- table(preds, as.factor(valdat$Inland))

  mat <- as.matrix(ans)
```



```

percent_correct <- sum(diag(mat))/floor(sum(mat))
result <- round((0.4 * percent_correct + 0.6 * get_fscore(mat)), 4)
return(result)
}

```

In [94]: *# This grid search searches for the best parameters for svm
modeling of the data. It takes a vector of seeds.*

```

gridSearch_svm <- function(seedv, dat, gammav, costv, folds=5) {

  gamma_len <- length(gammav)
  cost_len <- length(costv)
  # We need to capture the gridSearch parameters as well as  
# the cross-val scores.
  datout <- rep(NA, 2 * gamma_len * cost_len)
  dim(datout) <- c((gamma_len * cost_len), 2)
  datout <- as.data.frame(datout)
  colnames(datout) <- c("params", "Type2")
  datout$params <- ""

  # Divide dat by the number of folds to get a  
# size for each fold.
  segment_size <- round(nrow(dat)/folds)
  diff <- nrow(dat) - folds * segment_size
  last_seg_size <- segment_size + diff
  segmentsv <- c(rep(segment_size, (folds - 1)), last_seg_size)
  stopifnot(sum(segmentsv) == nrow(dat))

  index <- 0
  for(i in 1:gamma_len) {
    gamma <- gammav[i]
    for(j in 1:cost_len) {
      index <- index + 1
      cost <- costv[j]
      param_string <- paste(as.character(gamma),
                           as.character(cost), sep= "--")
      datout$params[index] <- param_string

      # Each set of parameters gets tested over many folds.  
# The different folds are created using different seeds.

      # Create a vector to store the Type2 score for each seed.
      seedv_len <- length(seedv)
      seed_scores <- rep(NA, seedv_len)

      for(h in 1:seedv_len) {
        # shuffle dat
        cur_seed <- seedv[h]
        set.seed(cur_seed)
        smp <- sample(rownames(dat), nrow(dat), replace= FALSE)
        dat <- dat[smp,]

        # Each element of row_list will be the rows we pick  
# out for one of the folds. E.g., the first element  
# of row_list will contain the rows we want for the  
# first fold, the second element of row_list will  
# contain the rows we want for the second fold, and  
# so forth.
        row_list <- vector("list", length=folds)
        names(row_list) <- as.character(1:folds)
        startpt <- 1
        for(k in 1:folds) {
          endpt <- startpt + segmentsv[k] - 1
          stopifnot(endpt <= nrow(dat))
          row_list[[k]] <- rownames(dat)[startpt:endpt]
          startpt <- endpt + 1
        }

        train_list <- test_list <- vector("list", length= folds)
        for(k in 1:folds) {

```

```

        testdat <- dat[row_list[[k]],]
        traindat <- dat[which(!(rownames(dat) %in% rownames(testdat))),]
        stopifnot((length(rownames(traindat)) + length(rownames(testdat))) == n)
        test_list[[k]] <- testdat
        train_list[[k]] <- traindat
      }
      # When there are only 5 folds, only 5 cores get used.
      scores <- mcmapply(get_Type2_svm, train_list, test_list,
                        MoreArgs= list(gamma=gamma, cost=cost),
                        SIMPLIFY= TRUE, mc.cores=5)
      # For the current seed, store the average of the Type2
      # scores, the average taken over the folds.
      seed_scores[h] <- round(mean(scores), 5)

    } ## end of for-loop, index h

    # Here I am taking an average of average scores. This
    # could be improved by simply taking a single average.
    datout$Type2[index] <- round(mean(seed_scores), 5)

  } ## end of for-loop, index j
} ## end of for-loop, index i

return(datout)
}

```

In [95]: # Run grid search to get better parameters for the
svm classifier.

```

set.seed(7543)
seed_vector <- sample(1:9999, 1, replace=FALSE)

gamma_v <- seq(0.1, 1.0, by=0.1)
cost_v <- seq(100, 700, by=100)

start <- Sys.time()
paste("Start time: ", start, sep="")
ans <- gridSearch_svm(seed_vector, svmtrain, gamma_v, cost_v)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 13.81 mins (with 1 seed)

(best_params <- ans[which(ans$Type2 == max(ans$Type2)),]$params)
# '0.1--100'

(best_Type2 <- ans[which(ans$Type2 == max(ans$Type2)),]$Type2)
# 0.8509

```

'Start time: 2021-04-24 15:03:18'

Time difference of 13.81 mins

'0.1--100'

0.8509

In [97]: # Refine the grid search.

```

set.seed(7541)
seed_vector <- sample(1:9999, 5, replace=FALSE)

gamma_v <- seq(0.04, 0.1, by=0.02)
cost_v <- seq(80, 120, by=20)

start <- Sys.time()
paste("Start time: ", start, sep="")
ans <- gridSearch_svm(seed_vector, svmtrain, gamma_v, cost_v)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 3 mins (with 5 seeds)

```

```
(best_params <- ans[which(ans$Type2 == max(ans$Type2)),]$params)
# '0.06--80'

(best_Type2 <- ans[which(ans$Type2 == max(ans$Type2)),]$Type2)
# 0.8523

'Start time: 2021-04-24 15:24:11'

Time difference of 2.96 mins

'0.06--80'

0.85229
```

In []: *#&* Bookmark*

```
In [98]: # Construct an svm with the identified parameters.

# We need to get probability estimates from the output.
## Note the use of svm_scaled (which is just svmtrain, scaled)
svm02 <- svm(I(as.factor(Inland)) ~ ., data=svm_scaled, kernel="radial",
             gamma=0.06, cost=80, scale=FALSE, probability=TRUE)

pred <- fitted(svm02)
(ans <- table(pred, as.factor(svm_scaled$Inland)))
print(paste("f-score for 'best' svm classifier, trainset: ",
            as.character(get_fscore(as.matrix(ans))), sep=""))
# 0.8368

# Accuracy is 0.9004

# Type2 is 0.8622
```

```
pred    0    1
  0 6575  644
  1  372 2605

[1] "f-score for 'best' svm classifier, trainset: 0.8368"
```

Get scores for svm02 on testdat

```
In [103]: # Function for obtaining a set of scores on the testset data
# using svm02 as the classifier.

get_testdatScores_svm02 <- function(seedv, dat) {

  seedv_len <- length(seedv)
  datout <- rep(NA, 5 * seedv_len)
  dim(datout) <- c(seedv_len, 5)
  datout <- as.data.frame(datout)
  colnames(datout) <- c("fscore", "Acc", "Type2", "FN", "FP")
  rownames(datout) <- as.character(seedv)

  for(h in 1:seedv_len) {
    # shuffle dat
    cur_seed <- seedv[h]
    set.seed(cur_seed)
    # It is expected that dat is testdat, which has 10K rcds
    smp <- sample(rownames(dat), 4000, replace= FALSE)
    df <- dat[smp,]

    df_scaled <- scale(df[, -1], center=svm_centers,
                      scale=svm_scales)
    df_scaled <- as.data.frame(cbind(as.numeric(df$Inland), df_scaled),
                              row.names=rownames(df))
    colnames(df_scaled) <- colnames(df)
```

```

preds <- predict(svm02, newdata=df_scaled, scale=FALSE, probability=TRUE)
preds_transf <- as.numeric(attr(preds, "probabilities")[, 2])
names(preds_transf) <- rownames(df)
preds_transf[which(preds_transf >= 0.5)] <- 1
preds_transf[which(preds_transf < 0.5)] <- 0
preds_transf <- as.factor(preds_transf)
ans <- get_confusion(preds_transf, df[, "Inland", drop=FALSE])

mat <- as.matrix(ans[[1]])
fscore <- round(as.numeric(ans[[2]]), 4)
acc <- round(sum(diag(mat))/floor(sum(mat)), 4)
type2 <- round((0.4 * acc + 0.6 * ans[[2]]), 4)
FN <- as.numeric(mat[2,1])
FP <- as.numeric(mat[1,2])
datout[as.character(cur_seed), 1:5] <- c(fscore, acc, type2, FN, FP)
}
return(datout)
}

```

In [106]: *# Get svm02's scores on the testset data. For each of
the 250 seeds, I sample 4K (no replacement) from the 10K
set of testdat records.*

```

set.seed(1821)
seed_vector <- sample(1:9999, 250, replace=FALSE)

start <- Sys.time()
# paste("Start time: ", start, sep="")
dat_result <- get_testdatScores_svm02(seed_vector, svmtest)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 49 secs

```

Time difference of 49.18 secs

In [107]: `dim(dat_result)`
`head(dat_result)`

250 · 5

A data.frame: 6 × 5

	fscore	Acc	Type2	FN	FP
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
5934	0.8372	0.9010	0.8627	259	137
1953	0.8392	0.9000	0.8635	256	144
7591	0.8182	0.8870	0.8457	308	144
1038	0.8250	0.8945	0.8528	274	148
49	0.8338	0.8962	0.8588	260	155
3203	0.8317	0.8975	0.8580	266	144

```
In [108]: fscore_mean <- round(mean(dat_result$fscore), 4)
fscore_sd <- round(sd(dat_result$fscore), 4)
paste0("fscore mean: ", as.character(fscore_mean))
# 0.8265
paste0("fscore StdDev: ", as.character(fscore_sd))
# 0.0065
""
summary(dat_result$fscore)
```

'fscore mean: 0.8265'

'fscore StdDev: 0.0065'

"

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	0.810	0.823	0.827	0.827	0.831	0.842

```
In [109]: Acc_mean <- round(mean(dat_result$Acc), 4)
Acc_sd <- round(sd(dat_result$Acc), 4)
paste0("accuracy mean: ", as.character(Acc_mean))
# 0.8943
paste0("accuracy StdDev: ", as.character(Acc_sd))
# 0.0038
""
summary(dat_result$Acc)
```

'accuracy mean: 0.8943'

'accuracy StdDev: 0.0038'

"

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	0.883	0.892	0.894	0.894	0.897	0.905

```
In [110]: Type2_mean <- round(mean(dat_result$Type2), 4)
Type2_sd <- round(sd(dat_result$Type2), 4)
paste0("Type2 mean: ", as.character(Type2_mean))
# 0.8536
paste0("Type2 StdDev: ", as.character(Type2_sd))
# 0.0053
""
summary(dat_result$Type2)
```

'Type2 mean: 0.8536'

'Type2 StdDev: 0.0053'

"

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	0.839	0.850	0.854	0.854	0.857	0.868

```
In [111]: FN_mean <- round(mean(dat_result$FN), 2)
FN_sd <- round(sd(dat_result$FN), 2)
paste0("FN mean: ", as.character(FN_mean))
# 280.3
paste0("FN StdDev: ", as.character(FN_sd))
# 11.79
""
summary(dat_result$FN)
```

'FN mean: 280.3'

'FN StdDev: 11.79'

"

```

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
      244     272     280     280     280     317

In [112]: FP_mean <- round(mean(dat_result$FP), 2)
FP_sd <- round(sd(dat_result$FP), 2)
paste0("FP mean: ", as.character(FP_mean))
# 142.58
paste0("FP StdDev: ", as.character(FP_sd))
# 8.91
""
summary(dat_result$FP)

```

'FP mean: 142.58'

'FP StdDev: 8.91'

"

```

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
      118     137     143     143     149     170

```

```

In [ ]: ### COMMENT:

# After rfclf_best, svm02 is the next best model. It has almost
# the same average accuracy score as rfclf_best. On average,
# svm02 has fewer false positives (143 vs 166). svm02 has more
# false negatives than rfclf_best: 280 vs 242.

```

Section 2: Model data with k-means

```
In [20]: svm02_columns
```

```
'Inland' · 'housing_median_age' · 'median_house_value' · 'HHdens_ln' · 'rooms_per_hh' · 'bdrms_per_room' ·
'pop_per_hh'
```

```
In [17]: # For k-means, use the same predictors as we use for the
# svm02 and f03 models.
```

```
km_columns <- svm02_columns
km_predictors <- km_columns[-1]
```

```
In [50]: # Use PCA to reduce the number of variables we are
# working with. This will make finding weights
# quite a bit easier.
```

```
pca <- prcomp(traindat[, km_predictors], center=TRUE, scale.=TRUE,
              rank.=4, retx=TRUE)
summary(pca)
```

Importance of first k=4 (out of 6) components:

	PC1	PC2	PC3	PC4
Standard deviation	1.415	1.203	0.974	0.845
Proportion of Variance	0.334	0.241	0.158	0.119
Cumulative Proportion	0.334	0.575	0.733	0.852

```
In [25]: head(pca$x)
```

A matrix: 6 × 4 of type dbl

	PC1	PC2	PC3	PC4
11174	0.62155	-0.40828	0.78395	0.02940
1542	1.55471	-0.36770	0.39013	1.04186
3537	-1.06642	0.42280	-1.37004	0.98605

	PC1	PC2	PC3	PC4
12747	0.23624	0.76784	-0.98573	1.35681
13348	-0.87807	0.12463	0.52692	-0.97772

```
In [51]: # Construct training set data for k-means modeling
# with PCA.

kmtrain <- cbind(as.numeric(traindat$Inland), as.data.frame(pca$x))
rownames(kmtrain) <- rownames(traindat)
colnames(kmtrain) <- c("Inland", "pc1", "pc2", "pc3", "pc4")
head(kmtrain)
```

A data.frame: 6 × 5

	Inland	pc1	pc2	pc3	pc4
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
11174	0	0.62155	-0.40828	0.78395	0.02940
1542	0	1.55471	-0.36770	0.39013	1.04186
3537	0	-1.06642	0.42280	-1.37004	0.98605
12747	1	0.23624	0.76784	-0.98573	1.35681
13348	1	-0.87807	0.12463	0.52692	-0.97772
6628	0	-1.51037	-0.42055	-1.56922	0.36121

Construct initial k-means model

```
In [30]: summary(kmtrain[, -1])
```

pc1	pc2	pc3	pc4
Min. : -6.1240	Min. : -4.7373	Min. : -3.5855	Min. : -3.5949
1st Qu.: -0.9157	1st Qu.: -0.6568	1st Qu.: -0.5856	1st Qu.: -0.5298
Median : -0.0184	Median : 0.0818	Median : 0.0152	Median : 0.0435
Mean : 0.0000	Mean : 0.0000	Mean : 0.0000	Mean : 0.0000
3rd Qu.: 0.8990	3rd Qu.: 0.7622	3rd Qu.: 0.5407	3rd Qu.: 0.5846
Max. : 12.4876	Max. : 9.5072	Max. : 14.4457	Max. : 4.4018

```
In [52]: # Apply min-max scaling to reduce the effects of the
# more extreme values.

kmtrain_scaled <- apply(kmtrain[, -1], MARGIN=2, range=01)
kmtrain_scaled <- as.data.frame(cbind(as.numeric(kmtrain$Inland), kmtrain_scaled),
                                row.names=row.names(kmtrain))
colnames(kmtrain_scaled) <- colnames(kmtrain)
head(kmtrain_scaled)
```

A data.frame: 6 × 5

	Inland	pc1	pc2	pc3	pc4
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
11174	0	0.36244	0.30391	0.24233	0.45322
1542	0	0.41258	0.30676	0.22049	0.57983
3537	0	0.27174	0.36225	0.12287	0.57285
12747	1	0.34174	0.38647	0.14418	0.61922
13348	1	0.28186	0.34132	0.22807	0.32728
6628	0	0.24789	0.30305	0.11182	0.49472

```
In [32]: summary(kmtrain_scaled[, -1])
```

pc1	pc2	pc3	pc4
Min. :0.000	Min. :0.000	Min. :0.000	Min. :0.000
1st Qu.:0.280	1st Qu.:0.286	1st Qu.:0.166	1st Qu.:0.383
Median :0.328	Median :0.338	Median :0.200	Median :0.455
Mean :0.329	Mean :0.333	Mean :0.199	Mean :0.450
3rd Qu.:0.377	3rd Qu.:0.386	3rd Qu.:0.229	3rd Qu.:0.523
Max. :1.000	Max. :1.000	Max. :1.000	Max. :1.000

In [53]: *# Run k-means algorithm with number of clusters set to 2.*

```
set.seed(1233)
start <- Sys.time()
km_mod <- kmeans(kmtrain_scaled[, -1], 2, iter.max = 50, nstart = 15)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 0.04 secs

print(km_mod$size)
```

Time difference of 0.06 secs

```
[1] 5503 4693
```

In [54]: `datout <- as.data.frame(cbind(kmtrain_scaled$Inland, km_mod$cluster))`
`colnames(datout) <- c("Inland", "cluster")`
`rownames(datout) <- rownames(kmtrain_scaled)`
`head(datout)`

A data.frame: 6 × 2

	Inland	cluster
	<dbl>	<dbl>
11174	0	1
1542	0	1
3537	0	1
12747	1	1
13348	1	2
6628	0	1

In [55]: `table(as.factor(datout$cluster))`

```
 1    2
5503 4693
```

In [56]: `table(as.factor(datout$Inland))`

```
 0    1
6947 3249
```

In [57]: `dfc1 <- datout[which(datout$cluster== 1),]`
`nrow(dfc1)`
`(ans <- table(as.factor(dfc1$Inland)))`

```
5503
```

```
 0    1
4118 1385
```

In [58]: `dfc2 <- datout[which(datout$cluster== 2),]`


```
table(as.factor(df2$Inland))
```

```

  0    1
2829 1864

```

In [59]: *# Get percent of Inland districts in kmtrain.*

```
kmtrain_Inland_percent <- mean(kmtrain$Inland)
round(kmtrain_Inland_percent, 4)
```

```
0.3187
```

In [60]: *# See how the clusters are associated with Inland.*

```
dfout <- as.data.frame(cbind(kmtrain_scaled$Inland, km_mod$cluster))
colnames(dfout) <- c("Inland", "cluster")
rownames(dfout) <- rownames(kmtrain_scaled)
dat_c1 <- dfout[which(dfout$cluster== 1),]
datc1_Inland_percent <- mean(dat_c1$Inland)

tmpdat <- dfout

c1_to_InlandYES <- FALSE
if(datc1_Inland_percent >= kmtrain_Inland_percent) { c1_to_InlandYES <- TRUE }

if(c1_to_InlandYES) {
  # cluster 1 is associated with the Inland districts
  tmpdat[which(tmpdat$cluster== 1),]$Inland <- 1
  tmpdat[which(tmpdat$cluster== 2),]$Inland <- 0
} else {
  # cluster 2 is associated with the Inland districts
  tmpdat[which(tmpdat$cluster== 2),]$Inland <- 1
  tmpdat[which(tmpdat$cluster== 1),]$Inland <- 0
}

# Generate confusion matrix for the k-means clusters and
# the corresponding f-score.
preds <- as.factor(tmpdat$Inland)
names(preds) <- rownames(tmpdat)
ans <- get_confusion(preds, kmtrain_scaled[, "Inland", drop=FALSE])

print(ans$matrix)
''

print(paste("initial f-score for kmeans, train set: ", as.character(ans[[2]]), sep=""))
# [1] "initial f-score for kmeans, train set: 0.4694"

# The result returned is a Type2 score (which is a mixture
# of accuracy and f-score).
mat <- as.matrix(ans[[1]])
percent_correct <- sum(diag(mat))/floor(sum(mat))
result <- round((0.4 * percent_correct + 0.6 * ans[[2]]), 4)

''

print(paste("Initial Type2 score for kmeans, train set: ", as.character(result), sep=""))
# [1] "Initial Type2 score for kmeans, train set: 0.5163"

""

print(paste("Accuracy: ", as.character(round(percent_correct,4)), sep=""))
# [1] "Accuracy: 0.5867"
```

```

      0    1 class.error
0 4118 2829      0.4072
1 1385 1864      0.4263

```

```
"
```

```
[1] "initial f-score for kmeans, train set: 0.4694"
```

```
"
```

```
[1] "Initial Type2 score for kmeans, train set: 0.5163"
```

```
"
```

```
[1] "Accuracy: 0.5867"
```

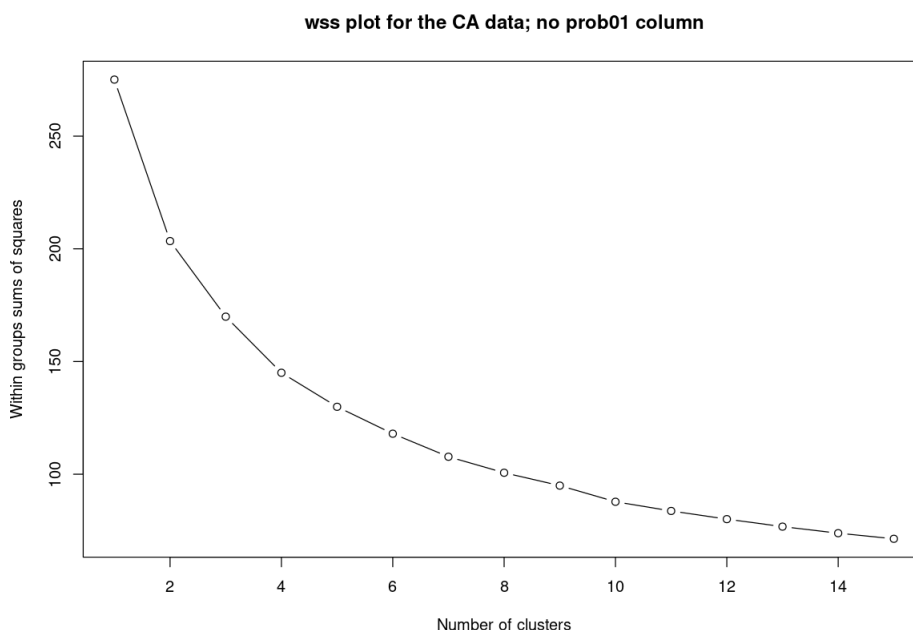
```
In [ ]: ### COMMENT:
```

```
# The k-means algorithm has difficulty with this dataset.
# We see from the following wss plot that k-means would
# do somewhat better if we asked it to find 3 clusters, not
# 2.
```

```
In [44]: # See whether 2 clusters makes sense for this dataframe.
```

```
options(repr.plot.width= 10, repr.plot.height= 7)
```

```
wssplot(kmtrain_scaled[, -1], title= "wss plot for the CA data; no prob01 column")
```



```
In [ ]: ### COMMENT:
```

```
# The above plot shows no obvious "kink" in the curve.
```

```
# It is not at all clear that 2 subgroups
# are optimal. 3 would seem to work just as well.
# This is interesting because, as seen above, op_transf
# (a variable removed from the current dataset) has 3
# levels: OCEAN, INLAND, and NEAR BAY. op_transf is a
# categorical variable that is a transformation of
# ocean_proximity, which has 4 levels: OCEAN, NEAR OCEAN,
# NEAR BAY, and INLAND. In the AnalyzeCAhousingData set
# of Jupyter notebooks, I found that NEAR OCEAN and OCEAN
# are enough alike (in terms of predicting median_house_value)
# that they ought to be reduced to a single level.
```

```
# In any case, for this dataset, it is clear that,
# unlike the downer cow data, there is not a natural
# ordering into 2 subgroups. Recall that with the cow data
# the specific purpose of collecting variables such as
# AST and CK was to see if they could help predict surviving
# cows from non-survivors. By contrast, the CA housing
# variables I am working with were not collected with the
```

```
# purpose of distinguishing Inland districts from districts
# closer to the ocean. So it is reasonable to expect that
# there would be no obvious kink in the above curve.
```

Construct k-means model with prob01 from rfclf_best

```
In [61]: # Add a prob01 column to kmtrain.

preds <- predict(rfclf_best, newdata=trainat[, rfclf_columns],
                  type="prob")


# preds is a matrix with 2 columns.
kmtrain_scaled$prob01 <- preds[, 2]
summary(kmtrain_scaled$prob01)

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
0.00000 0.00889 0.04778 0.31911 0.83778 1.00000
```

```
In [54]: # See if a 2-cluster solution now makes more sense.

options(repr.plot.width= 10, repr.plot.height= 7)

suppressMessages(wssplot(kmtrain_scaled[, -1],
                          title= "wss plot for the CA data with prob01 column"))
```



```
In [ ]: ### COMMENT:

# This is exactly the kind of curve we want to see if
# we want to partition the data into 2 clusters.
```

Get scores on trainset for k-means with prob01, no weights

```
In [62]: # Run k-means algorithm with number of clusters set to 2.

set.seed(1233)
start <- Sys.time()
kmp1 <- kmeans(kmtrain_scaled[, -1], 2, iter.max = 50, nstart = 15)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 0.03 secs
```

```
print(kmpl$size)
```

Time difference of 0.03 secs

```
[1] 6947 3249
```

In [63]: *# See how the clusters are associated with Inland.*

```
dfout <- as.data.frame(cbind(kmtrain_scaled$Inland, kmpl$cluster))
colnames(dfout) <- c("Inland", "cluster")
rownames(dfout) <- rownames(kmtrain_scaled)
dat_c1 <- dfout[which(dfout$cluster== 1),]
datc1_Inland_percent <- mean(dat_c1$Inland)

tmpdat <- dfout

c1_to_InlandYES <- FALSE
if(datc1_Inland_percent >= kmtrain_Inland_percent) { c1_to_InlandYES <- TRUE }

if(c1_to_InlandYES) {
  # cluster 1 is associated with the Inland districts
  tmpdat[which(tmpdat$cluster== 1),]$Inland <- 1
  tmpdat[which(tmpdat$cluster== 2),]$Inland <- 0
} else {
  # cluster 2 is associated with the Inland districts
  tmpdat[which(tmpdat$cluster== 2),]$Inland <- 1
  tmpdat[which(tmpdat$cluster== 1),]$Inland <- 0
}

preds <- as.factor(tmpdat$Inland)
names(preds) <- rownames(tmpdat)
ans <- get_confusion(preds, kmtrain_scaled[, "Inland", drop=FALSE])

print(ans$matrix)
''
print(paste("f-score for kmeans with prob01, trainset: ", as.character(ans[[2]]), sep=""))

mat <- as.matrix(ans[[1]])
percent_correct <- sum(diag(mat))/floor(sum(mat))
result <- round((0.4 * percent_correct + 0.6 * ans[[2]]), 4)

''
print(paste("Type2 score for kmeans with prob01, trainset: ", as.character(result), sep=""))
""
print(paste("Accuracy: ", as.character(round(percent_correct,4)), sep=""))
```

```
      0      1 class.error
0 6947      0          0
1      0 3249          0
```

"

```
[1] "f-score for kmeans with prob01, trainset: 1"
```

"

```
[1] "Type2 score for kmeans with prob01, trainset: 1"
```

"

```
[1] "Accuracy: 1"
```

Get scores on testset for k-means with prob01, no weights

In [64]: *# Add a prob01 column to testdat.*

```
kmtest <- testdat
```

```

preds <- predict(rfclf_best, newdata=testdat[, rfclf_columns],
                 type="prob")

# preds is a matrix with 2 columns.
kmtest$prob01 <- preds[, 2]

```

In [65]: kmp1\$centers

A matrix: 2 × 5 of type dbl

	pc1	pc2	pc3	pc4	prob01
1	0.31507	0.30647	0.20443	0.46747	0.045566
2	0.35892	0.38838	0.18692	0.41122	0.903993

In [66]: *# Function for obtaining a set of scores on the testset data
using kmp1 as the classifier.*

```

### NOTE: the value of c1_to_InlandYES from above is being
### used in the following function.

get_testdatScores_kmp1 <- function(seedv, dat) {

  n_smp <- 4000

  seedv_len <- length(seedv)
  datout <- rep(NA, 5 * seedv_len)
  dim(datout) <- c(seedv_len, 5)
  datout <- as.data.frame(datout)
  colnames(datout) <- c("fscore", "Acc", "Type2", "FN", "FP")
  rownames(datout) <- as.character(seedv)

  # Using model kmp1 from above.
  ctr_list <- vector("list", length= n_smp)
  for(i in 1:n_smp) {
    ctr_list[[i]] <- kmp1$centers
  }

  for(h in 1:seedv_len) {
    # shuffle dat
    cur_seed <- seedv[h]
    set.seed(cur_seed)
    # It is expected that dat is testdat, which has 10K rcds
    smp <- sample(rownames(dat), n_smp, replace= FALSE)
    df <- dat[smp,]

    # CAUTION: df has the prob01 column.

    # Using model pca (constructed from the training set).
    df_pca <- predict(pca, df[, km_predictors])
    df02 <- cbind(as.data.frame(df_pca), df$prob01)

    # Apply min-max scaling to df02.
    df02_scaled <- apply(df02, MARGIN=2, range01)
    df02_scaled <- as.data.frame(df02_scaled,
                                row.names=rownames(df))
    colnames(df02_scaled) <- colnames(kmp1$centers)

    # Each element of the following list is a row of df.
    df02_asList <- split(df02_scaled[, colnames(kmp1$centers)], seq(n_smp))
    names(ctr_list) <- rownames(df)

    # Get the predictions for df.
    preds <- mcmapply(getCluster, df02_asList, ctr_list,
                      SIMPLIFY=TRUE, mc.cores=6)
    df$cluster <- as.numeric(preds)
  }
}

```

```

df$pred_Inland <- NA
if(c1_to_InlandYES) {
  df[which(df$cluster==1),]$pred_Inland <- 1
  df[which(df$cluster==2),]$pred_Inland <- 0
} else {
  df[which(df$cluster==1),]$pred_Inland <- 0
  df[which(df$cluster==2),]$pred_Inland <- 1
}

# Generate confusion matrix.
preds <- as.factor(df$pred_Inland)
names(preds) <- rownames(df)
ans <- get_confusion(preds, df[, "Inland", drop=FALSE])

mat <- as.matrix(ans[[1]])
fscore <- round(as.numeric(ans[[2]]), 4)
acc <- round(sum(diag(mat))/floor(sum(mat)), 4)
type2 <- round((0.4 * acc + 0.6 * ans[[2]]), 4)
FN <- as.numeric(mat[2,1])
FP <- as.numeric(mat[1,2])
datout[as.character(cur_seed), 1:5] <- c(fscore,acc,type2,FN,FP)
}
return(datout)
}

```

In [67]: *# Get kmp1's scores on the testset data. For each of
the 250 seeds, I sample 4K (no replacement) from the 10K
set of testdat records.*

```

set.seed(1821)
seed_vector <- sample(1:9999, 250, replace=FALSE)

start <- Sys.time()
paste("Start time: ", start, sep="")
dat_result <- get_testdatScores_kmp1(seed_vector, kmtest)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 7.81 mins

```

'Start time: 2021-04-27 22:10:36'

Time difference of 7.81 mins

In [68]: `dim(dat_result)`
`head(dat_result)`

250 · 5

A data.frame: 6 × 5

	fscore	Acc	Type2	FN	FP
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
5934	0.8348	0.8960	0.8593	226	190
1953	0.8437	0.8995	0.8660	215	187
7591	0.8325	0.8928	0.8566	259	170
1038	0.8364	0.8985	0.8612	231	175
49	0.8326	0.8922	0.8564	229	202
3203	0.8362	0.8968	0.8604	225	188

In [69]: `fscore_mean <- round(mean(dat_result$fscore), 4)`
`fscore_sd <- round(sd(dat_result$fscore), 4)`
`paste0("fscore mean: ", as.character(fscore_mean))`
0.8348
`paste0("fscore StdDev: ", as.character(fscore_sd))`

```
# 0.0061
""
summary(dat_result$fscore)
```

```
'fscore mean: 0.8348'
```

```
'fscore StdDev: 0.0061'
```

```
"
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.820	0.831	0.835	0.835	0.839	0.852

```
In [70]: Acc_mean <- round(mean(dat_result$Acc), 4)
Acc_sd <- round(sd(dat_result$Acc), 4)
paste0("accuracy mean: ", as.character(Acc_mean))
# 0.8957
paste0("accuracy StdDev: ", as.character(Acc_sd))
# 0.0037
""
summary(dat_result$Acc)
```

```
'accuracy mean: 0.8957'
```

```
'accuracy StdDev: 0.0037'
```

```
"
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.886	0.893	0.895	0.896	0.898	0.908

```
In [71]: Type2_mean <- round(mean(dat_result$Type2), 4)
Type2_sd <- round(sd(dat_result$Type2), 4)
paste0("Type2 mean: ", as.character(Type2_mean))
# 0.8592
paste0("Type2 StdDev: ", as.character(Type2_sd))
# 0.005
""
summary(dat_result$Type2)
```

```
'Type2 mean: 0.8592'
```

```
'Type2 StdDev: 0.005'
```

```
"
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.847	0.856	0.859	0.859	0.862	0.874

```
In [72]: FN_mean <- round(mean(dat_result$FN), 2)
FN_sd <- round(sd(dat_result$FN), 2)
paste0("FN mean: ", as.character(FN_mean))
# 233.07
paste0("FN StdDev: ", as.character(FN_sd))
# 11.34
""
summary(dat_result$FN)
```

```
'FN mean: 233.07'
```

```
'FN StdDev: 11.34'
```

```
"
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
195	225	233	233	240	262

```
In [73]: FP_mean <- round(mean(dat_result$FP), 2)
FP_sd <- round(sd(dat_result$FP), 2)
paste0("FP mean: ", as.character(FP_mean))
```

```
# 184.21
paste0("FP StdDev: ", as.character(FP_sd))
# 10.59
""
summary(dat_result$FP)

'FP mean: 184.21'

'FP StdDev: 10.59'

"

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
      157     175     185     184     192     210
```

Comments on kmp1, no weights

The kmp1 model is competitive with rfclf_best, our current best model. On the testset data, rfclf_best has an accuracy score of 0.8979; for kmp1, this score is 0.8957. rfclf_best has a Type2 score of 0.8611; for kmp1, this score is 0.8592.

We should be able to improve the kmp1 model using weights obtained through cross-validation.

Section 3: Use 5 principal components with prob01, no weights

See if we can improve the k-means model by adding a fifth principal component. The downside of doing this is the additional effort involved in finding optimal weights. However, if we get a noticeable improvement over the rfclf_best model, we can perhaps forego the weights.

```
In [19]: pca5 <- prcomp(traindat[, km_predictors], center=TRUE, scale.=TRUE,
                    rank.=5, retx=TRUE)
summary(pca5)
```

```
Importance of first k=5 (out of 6) components:
              PC1  PC2  PC3  PC4  PC5
Standard deviation    1.415 1.203 0.974 0.845 0.7200
Proportion of Variance 0.334 0.241 0.158 0.119 0.0864
Cumulative Proportion 0.334 0.575 0.733 0.852 0.9383
```

```
In [21]: kmtrain <- cbind(traindat$Inland, as.data.frame(pca5$x))
rownames(kmtrain) <- rownames(traindat)
colnames(kmtrain) <- c("Inland", "pc1", "pc2", "pc3", "pc4", "pc5")
head(kmtrain)
```

A data.frame: 6 × 6

	Inland	pc1	pc2	pc3	pc4	pc5
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
11174	0	0.62155	-0.40828	0.78395	0.02940	-0.606426
1542	0	1.55471	-0.36770	0.39013	1.04186	-0.262446
3537	0	-1.06642	0.42280	-1.37004	0.98605	-0.119789
12747	1	0.23624	0.76784	-0.98573	1.35681	-0.857062
13348	1	-0.87807	0.12463	0.52692	-0.97772	-0.747739
6628	0	-1.51037	-0.42055	-1.56922	0.36121	0.092308

Construct initial k-means model


```
In [22]: # Apply min-max scaling to reduce the effects of the
# more extreme values.

kmtrain_scaled <- apply(kmtrain[, -1], MARGIN=2, range01)
kmtrain_scaled <- as.data.frame(cbind(as.numeric(kmtrain$Inland), kmtrain_scaled),
                                row.names=row.names(kmtrain))
colnames(kmtrain_scaled) <- colnames(kmtrain)
```

```
In [24]: # Get percent of Inland districts in kmtrain.

kmtrain_Inland_percent <- mean(kmtrain$Inland)
round(kmtrain_Inland_percent, 4)

0.3187
```

```
In [26]: # Add a prob01 column to kmtrain.

preds <- predict(rfc1f_best, newdata=train1dat[, rfc1f_columns],
                 type="prob")

# preds is a matrix with 2 columns.
kmtrain_scaled$prob01 <- preds[, 2]
summary(kmtrain_scaled$prob01)

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
0.000000 0.008889 0.04778 0.31911 0.83778 1.000000
```

Get scores on trainset for k-means with prob01, no weights

```
In [28]: # Run k-means algorithm with number of clusters set to 2.

set.seed(1233)
start <- Sys.time()
kmp1 <- kmeans(kmtrain_scaled[, -1], 2, iter.max = 50, nstart = 15)
stop <- Sys.time()
# round(stop - start, 2)
# Time difference of 0.03 secs

print(kmp1$size)

[1] 6947 3249
```

```
In [29]: # See how the clusters are associated with Inland.

dfout <- as.data.frame(cbind(kmtrain_scaled$Inland, kmp1$cluster))
colnames(dfout) <- c("Inland", "cluster")
rownames(dfout) <- rownames(kmtrain_scaled)
dat_c1 <- dfout[which(dfout$cluster == 1),]
datc1_Inland_percent <- mean(dat_c1$Inland)

tmpdat <- dfout

c1_to_InlandYES <- FALSE
if(datc1_Inland_percent >= kmtrain_Inland_percent) { c1_to_InlandYES <- TRUE }

if(c1_to_InlandYES) {
  # cluster 1 is associated with the Inland districts
  tmpdat[which(tmpdat$cluster == 1),]$Inland <- 1
  tmpdat[which(tmpdat$cluster == 2),]$Inland <- 0
} else {
  # cluster 2 is associated with the Inland districts
  tmpdat[which(tmpdat$cluster == 2),]$Inland <- 1
  tmpdat[which(tmpdat$cluster == 1),]$Inland <- 0
}
```

```

preds <- as.factor(tmpdat$Inland)
names(preds) <- rownames(tmpdat)
ans <- get_confusion(preds, kmtrain_scaled[, "Inland", drop=FALSE])

print(ans$matrix)
''
print(paste("f-score for kmeans with prob01, trainset: ", as.character(ans[[2]]), sep=""))

mat <- as.matrix(ans[[1]])
percent_correct <- sum(diag(mat))/floor(sum(mat))
result <- round((0.4 * percent_correct + 0.6 * ans[[2]]), 4)

''
print(paste("Type2 score for kmeans with prob01, trainset: ", as.character(result), sep=""))

""
print(paste("Accuracy: ", as.character(round(percent_correct,4)), sep=""))

      0      1 class.error
0 6947      0           0
1      0 3249           0

"

[1] "f-score for kmeans with prob01, trainset: 1"

"

[1] "Type2 score for kmeans with prob01, trainset: 1"

"

[1] "Accuracy: 1"

```

Get scores on testset for k-means with prob01, no weights

In [30]: *# Add a prob01 column to testdat.*

```

kmtest <- testdat

preds <- predict(rfclf_best, newdata=testdat[, rfclf_columns],
                 type="prob")

# preds is a matrix with 2 columns.
kmtest$prob01 <- preds[, 2]

```

In [31]: kmpl\$centers

A matrix: 2 × 6 of type dbl

	pc1	pc2	pc3	pc4	pc5	prob01
1	0.31507	0.30647	0.20443	0.46747	0.41426	0.045566
2	0.35892	0.38838	0.18692	0.41122	0.38720	0.903993

In [34]: *# Function for obtaining a set of scores on the testset data*
using kmpl as the classifier. Here 5 principal components
are being used.

NOTE: the value of c1_to_InlandYES from above is being
used in the following function.

```

get_testdatScores_kmplb <- function(seedv, dat) {

  n_smp <- 4000

  seedv_len <- length(seedv)
  datout <- rep(NA, 5 * seedv_len)

```

```

dim(datout) <- c(seedv_len, 5)
datout <- as.data.frame(datout)
colnames(datout) <- c("fscore", "Acc", "Type2", "FN", "FP")
rownames(datout) <- as.character(seedv)

# Using model kmp1 from above.
ctr_list <- vector("list", length= n_smp)
for(i in 1:n_smp) {
  ctr_list[[i]] <- kmp1$centers
}

for(h in 1:seedv_len) {
  # shuffle dat
  cur_seed <- seedv[h]
  set.seed(cur_seed)
  # It is expected that dat is testdat, which has 10K rcds
  smp <- sample(rownames(dat), n_smp, replace= FALSE)
  df <- dat[smp,]

  # CAUTION: df has the prob01 column.

  # Using model pca5 (constructed from the training set).
  df_pca <- predict(pca5, df[, km_predictors])
  df02 <- cbind(as.data.frame(df_pca), df$prob01)

  # Apply min-max scaling to df02.
  df02_scaled <- apply(df02, MARGIN=2, range01)
  df02_scaled <- as.data.frame(df02_scaled,
                                row.names=rownames(df))
  colnames(df02_scaled) <- colnames(kmp1$centers)

  # Each element of the following list is a row of df.
  df02_asList <- split(df02_scaled[, colnames(kmp1$centers)], seq(n_smp))
  names(ctr_list) <- rownames(df)

  # Get the predictions for df.
  preds <- mcmapply(getCluster, df02_asList, ctr_list,
                    SIMPLIFY=TRUE, mc.cores=6)
  df$cluster <- as.numeric(preds)

  df$pred_Inland <- NA
  if(c1 to InlandYES) {
    df[which(df$cluster==1),]$pred_Inland <- 1
    df[which(df$cluster==2),]$pred_Inland <- 0
  } else {
    df[which(df$cluster==1),]$pred_Inland <- 0
    df[which(df$cluster==2),]$pred_Inland <- 1
  }

  # Generate confusion matrix.
  preds <- as.factor(df$pred_Inland)
  names(preds) <- rownames(df)
  ans <- get_confusion(preds, df[, "Inland", drop=FALSE])

  mat <- as.matrix(ans[[1]])
  fscore <- round(as.numeric(ans[[2]]), 4)
  acc <- round(sum(diag(mat))/floor(sum(mat)), 4)
  type2 <- round((0.4 * acc + 0.6 * ans[[2]]), 4)
  FN <- as.numeric(mat[2,1])
  FP <- as.numeric(mat[1,2])
  datout[as.character(cur_seed), 1:5] <- c(fscore, acc, type2, FN, FP)
}
return(datout)
}

```

In [35]: # Get kmp1's scores on the testset data. For each of
 # the 250 seeds, I sample 4K (no replacement) from the 10K
 # set of testdat records.

```

set.seed(1821)
seed_vector <- sample(1:9999, 250, replace=FALSE)

start <- Sys.time()
paste("Start time: ", start, sep="")
dat_result <- get_testdatScores_kmplb(seed_vector, kmtest)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 8.47 mins

```

'Start time: 2021-04-25 18:04:45'

Time difference of 8.47 mins

```

In [37]: fscore_mean <- round(mean(dat_result$fscore), 4)
fscore_sd <- round(sd(dat_result$fscore), 4)
paste0("fscore mean: ", as.character(fscore_mean))
# 0.8348
paste0("fscore StdDev: ", as.character(fscore_sd))
# 0.006
""
summary(dat_result$fscore)

```

'fscore mean: 0.8348'

'fscore StdDev: 0.006'

"

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.820	0.831	0.835	0.835	0.839	0.852

```

In [38]: Acc_mean <- round(mean(dat_result$Acc), 4)
Acc_sd <- round(sd(dat_result$Acc), 4)
paste0("accuracy mean: ", as.character(Acc_mean))
# 0.8957
paste0("accuracy StdDev: ", as.character(Acc_sd))
# 0.0037
""
summary(dat_result$Acc)

```

'accuracy mean: 0.8957'

'accuracy StdDev: 0.0037'

"

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.887	0.893	0.896	0.896	0.898	0.908

```

In [39]: Type2_mean <- round(mean(dat_result$Type2), 4)
Type2_sd <- round(sd(dat_result$Type2), 4)
paste0("Type2 mean: ", as.character(Type2_mean))
# 0.8592
paste0("Type2 StdDev: ", as.character(Type2_sd))
# 0.005
""
summary(dat_result$Type2)

```

'Type2 mean: 0.8592'

'Type2 StdDev: 0.005'

"

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.847	0.856	0.859	0.859	0.862	0.874

```

In [40]: FN_mean <- round(mean(dat_result$FN), 2)
FN_sd <- round(sd(dat_result$FN), 2)

```

```
paste0("FN mean: ", as.character(FN_mean))
# 233.9
paste0("FN StdDev: ", as.character(FN_sd))
# 11.4
""
summary(dat_result$FN)
```

```
'FN mean: 233.9'
```

```
'FN StdDev: 11.4'
```

```
"
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
196	226	234	234	241	262

```
In [41]: FP_mean <- round(mean(dat_result$FP), 2)
FP_sd <- round(sd(dat_result$FP), 2)
paste0("FP mean: ", as.character(FP_mean))
# 183.13
paste0("FP StdDev: ", as.character(FP_sd))
# 10.38
""
summary(dat_result$FP)
```

```
'FP mean: 183.13'
```

```
'FP StdDev: 10.38'
```

```
"
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
156	175	184	183	190	210

```
In [ ]: ### COMMENT:

# There is NO improvement when using the 5th
# principal component.
```

Section 4: Find optimal weights for kmp1 (4 principal components)

In what follows I first make use of the method worked out in the Addendum of Part 2, using `tot.withinss` to help find optimal weights (since relying on this surrogate score takes less time). I find out, however, that this method does not work well for this dataset. The problem might not be with the method itself, since I am not able to get any boost at all from the weights I identify and/or test below. Perhaps the scaling that I am using (centering, scaling, then applying min-max) on the columns, in combination with the large number of records, makes weights superfluous.

```
In [47]: # Function for computing the tot.withinss for each set of
# weights in df_params (a dataframe, each row of which is
# a candidate set of weights). The optimal set of weights
# will be the set that yields the smallest average (over
# the folds) for tot.withinss.
# This function is called from gridSearch07.

get_tot.withinss_rfclf <- function(traindat, valdat, wghts) {

  # Apply pca to traindat (4 components).
  pca <- prcomp(traindat[, km_predictors], center=TRUE, scale.=TRUE,
                rank.=4, retx=FALSE)

  # Construct random forest model.
  set.seed(1493)
  rfclf <- randomForest(I(as.factor(Inland)) ~ .,
                        data= traindat[, rfclf_columns],
```

```

ntree=900, mtry= 3, nodesize= 1)

#####
# Prepare valdat.
#####

# Apply pca.
tmpdat <- predict(pca, valdat[, km_predictors])

# Apply min-max scaling.
valdat_scaled <- apply(tmpdat, MARGIN=2, range01)
valdat_scaled <- as.data.frame(valdat_scaled)
colnames(valdat_scaled) <- paste0("pc", 1:4)

# Add prob01 column.
preds_val <- predict(rfc1f, newdata= valdat[, rfc1f_columns], type="prob")
valdat_scaled$prob01 <- preds_val[, 2]

# Apply weights.
cols <- names(wgths)
valdat_wgths <- t(t(valdat_scaled[, cols]) * as.numeric(wgths[cols]))

# Construct k-means model on valdat to get tot.withinss.
kmod <- suppressWarnings(kmeans(valdat_wgths, 2, iter.max = 50, nstart=5))

return(kmod$tot.withinss)
}

```

```

In [48]: # This grid search searches for the best set of weights to use
# in our k-means clustering model. The best weights are those
# which generalize best to the validation set. So we look for
# the best cross-validation score.

gridSearch07 <- function(seed_vector, dat, df_params, folds=7) {

  datout <- rep(NA, 2*nrow(df_params))
  dim(datout) <- c(nrow(df_params), 2)
  datout <- as.data.frame(datout)
  colnames(datout) <- c("row", "tot.withinss")
  datout$row <- params_rows <- rownames(df_params)

  # We want the sqrt of the weights.
  df_params <- df_params^0.5

  #####
  # Partition the data into folds.

  segment_size <- round(dim(dat)[1]/folds)
  diff <- nrow(dat) - folds * segment_size
  last_seg_size <- segment_size + diff
  segmentsv <- c(rep(segment_size, (folds - 1)), last_seg_size)
  stopifnot(sum(segmentsv) == nrow(dat))

  # Create a dataframe, each row for a distinct seed.
  # Each column of the dataframe is for a distinct set
  # of weights. The entries in the cells are tot.withinss
  # scores.
  seedv_len <- length(seed_vector)
  df_scores <- rep(NA, seedv_len*nrow(df_params))
  dim(df_scores) <- c(seedv_len, nrow(df_params))
  df_scores <- as.data.frame(df_scores)
  colnames(df_scores) <- rownames(df_params)
  rownames(df_scores) <- as.character(seed_vector)

  for(h in 1:seedv_len) {
    # shuffle dat
    cur_seed <- seed_vector[h]
    set.seed(cur_seed)
    smp <- sample(rownames(dat), nrow(dat), replace= FALSE)

```

```

dat <- dat[smp,]

# Each element of row_list will be the rows we pick
# out for one of the folds. E.g., the first element
# of row_list will contain the rows we want for the
# first fold, the second element of row_list will
# contain the rows we want for the second fold, and
# so forth.
row_list <- vector("list", length=folds)
names(row_list) <- as.character(1:folds)
startpt <- 1
for(k in 1:folds) {
  endpt <- startpt + segmentsv[k] - 1
  stopifnot(endpt <= nrow(dat))
  row_list[[k]] <- rownames(dat)[startpt:endpt]
  startpt <- endpt + 1
}

for(i in 1:nrow(df_params)) {

  cur_row <- params_rows[i]
  wghts <- as.numeric(df_params[i,])
  names(wghts) <- colnames(df_params)

  train_list <- test_list <- vector("list", length= folds)
  for(j in 1:folds) {
    testdat <- dat[row_list[[j]],]
    traindat <- dat[which(!(rownames(dat) %in% rownames(testdat))),]
    stopifnot((length(rownames(traindat)) + length(rownames(testdat))) == nrow(dat))
    test_list[[j]] <- testdat
    train_list[[j]] <- traindat
  }
  # When there are only 5 folds, only 5 cores get used.
  scores <- mcmapply(get_tot.withinss_rfclf, train_list, test_list,
                    MoreArgs= list(wghts=wghts),
                    SIMPLIFY= TRUE, mc.cores=7)

  # For the current seed, store the average of the tot.withinss
  # scores, the average taken over the folds.
  df_scores[as.character(cur_seed), cur_row] <- round(mean(scores), 5)

} # end of for-loop, index i
} ## end of for-loop, index h

# Compute the average over the seeds of the tot.withinss scores
# obtained for each set of parameters in df_params.
datout$tot.withinss <- round(apply(df_scores, MARGIN=2, mean), 5)
return(datout)
}

```

```

In [49]: # There are 5 parameter lists to work with. The best
# approach, perhaps, is to start by exploring the
# region around the space where all parameters have an
# equal weight--in this case, a weight of 0.20.

lst <- vector("list", length= 5)
names(lst) <- c(paste0("pc", 1:4), "prob01")

lst[[1]] <- lst[[2]] <- lst[[3]] <- lst[[4]] <- lst[[5]] <- seq(0.12, 0.28, by=0.02)

start <- Sys.time()
dfc01 <- generate_combs(lst)
stop <- Sys.time()
# round(stop - start, 2)

dim(dfc01)
# 3951 5

```

3951 · 5

In [42]: *# Test on a sample of 10.*

```
set.seed(42)
smp <- sample(rownames(dfc01), 10, replace=FALSE)
tst_params <- dfc01[smp,]
head(tst_params)
```

A data.frame: 6 × 5

	pc1	pc2	pc3	pc4	prob01
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
37613	0.14	0.18	0.22	0.24	0.22
34757	0.26	0.12	0.24	0.16	0.22
18669	0.16	0.20	0.22	0.26	0.16
17805	0.16	0.26	0.18	0.24	0.16
19845	0.28	0.28	0.14	0.12	0.18
46709	0.26	0.22	0.12	0.14	0.26

In [37]: *# Find the best weights of those in tst_params.*

```
set.seed(1221)
seed_vector <- sample(1:9999, 3, replace=FALSE)

start <- Sys.time()
dat_result <- gridSearch07(seed_vector, traindat, tst_params)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 3.42 mins (for 10 rows, 3 seeds, 5 folds per)
```

Time difference of 3.42 mins

In [43]: *# Run the test again with only 1 seed, but 10 folds and
using 10 cores.*

```
set.seed(1221)
seed_vector <- sample(1:9999, 1, replace=FALSE)

start <- Sys.time()
dat_result <- gridSearch07(seed_vector, traindat, tst_params)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 2.13 mins
```

Time difference of 2.13 mins

In [50]: *# Run the test again with only 1 seed, but 7 folds and
using 7 cores.*

```
set.seed(1221)
seed_vector <- sample(1:9999, 1, replace=FALSE)

start <- Sys.time()
dat_result <- gridSearch07(seed_vector, traindat, tst_params)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 1.7 mins
```

Time difference of 1.69 mins

```
In [51]: best_params <- dat_result[which(dat_result$tot.withinss ==
                                         min(dat_result$tot.withinss, na.rm=TRUE)),]$row
length(best_params)
```



```
best_tot.withinss <- round(dat_result[which(dat_result$tot.withinss ==
                                             min(dat_result$tot.withinss, na.rm=TRUE)),$tot.withinss, 2
1
```

```
In [52]: dfc01[best_params,]

best_tot.withinss
```

A data.frame: 1 × 5

	pc1	pc2	pc3	pc4	prob01
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
21093	0.22	0.18	0.28	0.14	0.18

13.47

```
In [53]: # Find the best weights of those in dfc01 (3951 rows,
# 1 seed, 7 folds).

set.seed(1221)
seed_vector <- sample(1:9999, 1, replace=FALSE)

start <- Sys.time()
paste("Start time: ", start, sep="")
dat_result <- gridSearch07(seed_vector, traindat, dfc01)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 11.29 hours.
```

'Start time: 2021-04-25 22:20:20'

Time difference of 11.29 hours

```
In [54]: best_params <- dat_result[which(dat_result$tot.withinss ==
                                         min(dat_result$tot.withinss, na.rm=TRUE)),$row
length(best_params)

best_tot.withinss <- round(dat_result[which(dat_result$tot.withinss ==
                                         min(dat_result$tot.withinss, na.rm=TRUE)),$tot.withinss, 2
1
```

```
In [55]: dfc01[best_params,]
# 693    0.28    0.20    0.28    0.12    0.12

best_tot.withinss
# 12.61
```

A data.frame: 1 × 5

	pc1	pc2	pc3	pc4	prob01
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
693	0.28	0.2	0.28	0.12	0.12

12.61

```
In [56]: # Refine the search.

lst <- vector("list", length= 5)
names(lst) <- c(paste0("pc", 1:4), "prob01")

lst[[1]] <- seq(0.26, 0.36, by=0.02)
lst[[2]] <- seq(0.16, 0.24, by=0.02)
```

```

lst[[3]] <- seq(0.26, 0.36, by=0.02)
lst[[4]] <- seq(0.06, 0.14, by=0.02)
lst[[5]] <- seq(0.06, 0.14, by=0.02)

start <- Sys.time()
dfc02 <- generate_combs(lst)
stop <- Sys.time()
# round(stop - start, 2)

dim(dfc02)
# 486 5

```

486 · 5

```

In [57]: # Find the best weights of those in dfc01 (486 rows,
# 1 seed, 7 folds).

set.seed(1233)
seed_vector <- sample(1:9999, 1, replace=FALSE)

start <- Sys.time()
paste("Start time: ", start, sep="")
dat_result <- gridSearch07(seed_vector, traindat, dfc02)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 1.39 hours.

```

'Start time: 2021-04-26 09:45:09'

Time difference of 1.39 hours

```

In [58]: best_params <- dat_result[which(dat_result$tot.withinss ==
min(dat_result$tot.withinss, na.rm=TRUE)),]$row
length(best_params)

best_tot.withinss <- round(dat_result[which(dat_result$tot.withinss ==
min(dat_result$tot.withinss, na.rm=TRUE)),]$tot.withinss, 2)

1

```

```

In [59]: dfc02[best_params,]
# 156 0.36 0.16 0.36 0.06 0.06

best_tot.withinss
# 10.59

```

A data.frame: 1 × 5

	pc1	pc2	pc3	pc4	prob01
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
156	0.36	0.16	0.36	0.06	0.06

10.59

```

In [60]: # Refine the search.

lst <- vector("list", length= 5)
names(lst) <- c(paste0("pc", 1:4), "prob01")

lst[[1]] <- seq(0.36, 0.40, by=0.02)
lst[[2]] <- seq(0.10, 0.16, by=0.02)
lst[[3]] <- seq(0.36, 0.40, by=0.02)
lst[[4]] <- seq(0.04, 0.06, by=0.02)
lst[[5]] <- seq(0.04, 0.06, by=0.02)

start <- Sys.time()
dfc03 <- generate_combs(lst)

```

```
stop <- Sys.time()
# round(stop - start, 2)

dim(dfc03)
# 30      5

30 * 5
```

```
In [61]: # Find the best weights of those in dfc01 (30 rows,
# 1 seed, 7 folds).

set.seed(1235)
seed_vector <- sample(1:9999, 1, replace=FALSE)

start <- Sys.time()
paste("Start time: ", start, sep="")
dat_result <- gridSearch07(seed_vector, traindat, dfc03)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 5.12 mins.
```

'Start time: 2021-04-26 11:13:16'

Time difference of 5.12 mins

```
In [62]: best_params <- dat_result[which(dat_result$tot.withinss ==
                                         min(dat_result$tot.withinss, na.rm=TRUE)),]$row
length(best_params)

best_tot.withinss <- round(dat_result[which(dat_result$tot.withinss ==
                                             min(dat_result$tot.withinss, na.rm=TRUE)),]$tot.withinss, 2)

1
```

```
In [63]: dfc03[best_params,]
# 34      0.36      0.16      0.40      0.04      0.04

best_tot.withinss
# 11.43
```

A data.frame: 1 × 5

	pc1	pc2	pc3	pc4	prob01
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
34	0.36	0.16	0.4	0.04	0.04

11.43

In [64]: *# Refine the search.*

```

lst <- vector("list", length= 5)
names(lst) <- c(paste0("pc", 1:4), "prob01")

lst[[1]] <- seq(0.34, 0.37, by=0.01)
lst[[2]] <- seq(0.15, 0.17, by=0.01)
lst[[3]] <- seq(0.38, 0.41, by=0.01)
lst[[4]] <- seq(0.04, 0.06, by=0.01)
lst[[5]] <- seq(0.04, 0.06, by=0.01)

start <- Sys.time()
dfc04 <- generate_combs(lst)
stop <- Sys.time()
# round(stop - start, 2)

dim(dfc04)
# 71      5

```

71 · 5

In [65]: *# Find the best weights of those in dfc01 (71 rows,
1 seed, 7 folds).*

```

set.seed(245)
seed_vector <- sample(1:9999, 1, replace=FALSE)

start <- Sys.time()
paste("Start time: ", start, sep="")
dat_result <- gridSearch07(seed_vector, traindat, dfc04)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 12.33 mins.

```

'Start time: 2021-04-26 11:23:39'

Time difference of 12.33 mins

In [66]:

```

best_params <- dat_result[which(dat_result$tot.withinss ==
                                min(dat_result$tot.withinss, na.rm=TRUE)),]$row
length(best_params)

best_tot.withinss <- round(dat_result[which(dat_result$tot.withinss ==
                                min(dat_result$tot.withinss, na.rm=TRUE)),]$tot.withinss, 2)

```

1

In [67]:

```

dfc04[best_params,]
# 39      0.36      0.15      0.41      0.04      0.04

best_tot.withinss
# 11.5

```

A data.frame: 1 × 5

	pc1	pc2	pc3	pc4	prob01
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
39	0.36	0.15	0.41	0.04	0.04

11.5

Test the weights using cross-validation

Here gridSearch06, rather than gridSearch07, is used. Up to this point, we have been relying on a surrogate score to obtain

optimal weights. gridSearch06 relies on the Type2 score.

```
In [71]: lst <- vector("list", length= 5)
names(lst) <- c(paste0("pc", 1:4), "prob01")

lst[[1]] <- c(0.36)
lst[[2]] <- c(0.15)
lst[[3]] <- c(0.41)
lst[[4]] <- c(0.04)
lst[[5]] <- c(0.04)

start <- Sys.time()
dfc05 <- generate_combs(lst)
stop <- Sys.time()
# round(stop - start, 2)

# Try other nearby weights.
dfc05 <- rbind(dfc05, c(0.36, 0.16, 0.36, 0.06, 0.06))
dfc05 <- rbind(dfc05, c(0.28, 0.20, 0.28, 0.12, 0.12))
dfc05 <- rbind(dfc05, c(0.32, 0.18, 0.32, 0.09, 0.09))
dfc05
```

A data.frame: 4 × 5

pc1	pc2	pc3	pc4	prob01
<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
0.36	0.15	0.41	0.04	0.04
0.36	0.16	0.36	0.06	0.06
0.28	0.20	0.28	0.12	0.12
0.32	0.18	0.32	0.09	0.09

```
In [86]: # Find the best weights of those in dfc05,
# using 5 seeds.

set.seed(1933)
seed_vector <- sample(1:9999, 5, replace=FALSE)

start <- Sys.time()
paste("Start time: ", start, sep="")
dat_result <- gridSearch06(seed_vector, traindat, dfc05)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 2.61 mins.
```

'Start time: 2021-04-26 13:22:22'

Time difference of 2.61 mins

```
In [87]: datout <- cbind(dfc05, dat_result$Type2)
colnames(datout) <- c(colnames(dfc05), "Type2")
datout
```

A data.frame: 4 × 6

pc1	pc2	pc3	pc4	prob01	Type2
<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
0.36	0.15	0.41	0.04	0.04	0.85494
0.36	0.16	0.36	0.06	0.06	0.85562
0.28	0.20	0.28	0.12	0.12	0.85566
0.32	0.18	0.32	0.09	0.09	0.85571

```
In [88]: # Find the best weights of those in dfc05,
# using 11 seeds.

set.seed(17)
seed_vector <- sample(1:9999, 11, replace=FALSE)

start <- Sys.time()
paste("Start time: ", start, sep="")
dat_result <- gridSearch06(seed_vector, traindat, dfc05)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 5.84 mins.
```

'Start time: 2021-04-26 13:29:11'

Time difference of 5.84 mins

```
In [89]: datout <- cbind(dfc05, dat_result$Type2)
colnames(datout) <- c(colnames(dfc05), "Type2")
datout
```

A data.frame: 4 × 6

pc1	pc2	pc3	pc4	prob01	Type2
<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
0.36	0.15	0.41	0.04	0.04	0.85537
0.36	0.16	0.36	0.06	0.06	0.85618
0.28	0.20	0.28	0.12	0.12	0.85634
0.32	0.18	0.32	0.09	0.09	0.85625

```
In [90]: # Remove the first set of weights since it
# is getting the lowest score.
```

```
dfc05 <- dfc05[2:4,]
```

```
In [91]: # Find the best weights of those in dfc05,
# using 51 seeds.
```

```
set.seed(17)
seed_vector <- sample(1:9999, 51, replace=FALSE)

start <- Sys.time()
paste("Start time: ", start, sep="")
dat_result <- gridSearch06(seed_vector, traindat, dfc05)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 20.58 mins.
```

'Start time: 2021-04-26 13:49:25'

Time difference of 20.58 mins

```
In [92]: datout <- cbind(dfc05, dat_result$Type2)
colnames(datout) <- c(colnames(dfc05), "Type2")
datout
```

A data.frame: 3 × 6

	pc1	pc2	pc3	pc4	prob01	Type2
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
2	0.36	0.16	0.36	0.06	0.06	0.85625

```

      pc1    pc2    pc3    pc4  prob01   Type2
      <dbl> <dbl> <dbl> <dbl>   <dbl>   <dbl>

In [ ]: ### COMMENT:

# There is not much difference in average Type2 score
# among the different sets of weights. It appears
# that row 4 might be the best set.

# It is not at all clear that using tot.withinss works
# with this larger dataset. It might also be the case
# that getting a cross-val score from 7 folds is not
# sufficient.

# In order to have a model that performs better than
# our current best model, rfclf_best, I likely need
# an average cross-val Type2 score over 0.8610.

# wghts <- c(0.32, 0.18, 0.32, 0.09, 0.09)

```

```

In [17]: lst <- vector("list", length= 5)
names(lst) <- c(paste0("pc", 1:4), "prob01")

lst[[1]] <- c(0.31)
lst[[2]] <- c(0.18)
lst[[3]] <- c(0.31)
lst[[4]] <- c(0.10)
lst[[5]] <- c(0.10)

start <- Sys.time()
dfc06 <- generate_combs(lst)
stop <- Sys.time()
# round(stop - start, 2)

# Check that output from rep(0.20, 5) is same, or
# nearly the same, as output from rep(1, 5).
dfc06 <- rbind(dfc06, rep(1, 5))
dfc06 <- rbind(dfc06, rep(0.20, 5))

# Check the degree of difference in the Type2 score
# when we use weights which are NOT optimal.
dfc06 <- rbind(dfc06, c(0.24, 0.16, 0.20, 0.12, 0.28))
dfc06 <- rbind(dfc06, c(0.22, 0.14, 0.22, 0.14, 0.28))
dfc06

```

A data.frame: 5 × 5

pc1	pc2	pc3	pc4	prob01
<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
0.31	0.18	0.31	0.10	0.10
1.00	1.00	1.00	1.00	1.00
0.20	0.20	0.20	0.20	0.20
0.24	0.16	0.20	0.12	0.28
0.22	0.14	0.22	0.14	0.28

```

In [18]: # Find the best weights of those in dfc06,
# using 51 seeds.

set.seed(17)
seed_vector <- sample(1:9999, 51, replace=FALSE)

start <- Sys.time()
paste("Start time: ", start, sep="")
dat_result <- gridSearch06(seed_vector, traindat, dfc06)
stop <- Sys.time()

```

```
round(stop - start, 2)
# Time difference of 36 mins.
```

```
'Start time: 2021-04-27 19:06:53'
```

```
Time difference of 36.03 mins
```

```
In [19]: datout <- cbind(dfc06, dat_result$Type2)
colnames(datout) <- c(colnames(dfc06), "Type2")
datout
```

```
A data.frame: 5 × 6
```

pc1	pc2	pc3	pc4	prob01	Type2
<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
0.31	0.18	0.31	0.10	0.10	0.85637
1.00	1.00	1.00	1.00	1.00	0.85632
0.20	0.20	0.20	0.20	0.20	0.85632
0.24	0.16	0.20	0.12	0.28	0.85643
0.22	0.14	0.22	0.14	0.28	0.85647

```
In [ ]: ### COMMENTS:

# The method of finding weights using tot.withinss
# clearly did not work for this dataset. Part of
# the issue might be that 5 folds are simply not
# enough by which to judge the best set of weights.

# We can also see from the above table that a
# significant change in the weights makes very
# little difference in the average Type2 score. So
# from the results I have thus far, it appears as
# though using weights will not do much to improve
# our k-means hybrid model.
```

```
In [21]: # Try one more set of weights, giving more weight to
# the prob01 column.
```

```
lst <- vector("list", length= 5)
names(lst) <- c(paste0("pc", 1:4), "prob01")
```

```
lst[[1]] <- c(0.20)
lst[[2]] <- c(0.14)
lst[[3]] <- c(0.20)
lst[[4]] <- c(0.14)
lst[[5]] <- c(0.32)
```

```
start <- Sys.time()
dfc06 <- generate_combs(lst)
stop <- Sys.time()
```

```
dfc06
```

```
A data.frame: 1 × 5
```

	pc1	pc2	pc3	pc4	prob01
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	0.2	0.14	0.2	0.14	0.32

```
In [22]: # Find the best weights of those in dfc06,
# using 51 seeds.
```

```
set.seed(17)
```



```
seed_vector <- sample(1:9999, 51, replace=FALSE)

start <- Sys.time()
paste("Start time: ", start, sep="")
dat_result <- gridSearch06(seed_vector, traindat, dfc06)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 7.2 mins.
```

'Start time: 2021-04-27 20:05:20'

Time difference of 7.19 mins

```
In [23]: datout <- cbind(dfc06, dat_result$Type2)
colnames(datout) <- c(colnames(dfc06), "Type2")
datout
```

A data.frame: 1 × 6

	pc1	pc2	pc3	pc4	prob01	Type2
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	0.2	0.14	0.2	0.14	0.32	0.85644

In []:

Get scores on trainset for hybrid model, using weights

```
In [25]: pca4 <- prcomp(traindat[, km_predictors], center=TRUE, scale.=TRUE,
rank.=4, retx=TRUE)
summary(pca4)
```

Importance of first k=4 (out of 6) components:

	PC1	PC2	PC3	PC4
Standard deviation	1.415	1.203	0.974	0.845
Proportion of Variance	0.334	0.241	0.158	0.119
Cumulative Proportion	0.334	0.575	0.733	0.852

```
In [74]: kmtrain <- cbind(traindat$Inland, as.data.frame(pca4$x))
rownames(kmtrain) <- rownames(traindat)
colnames(kmtrain) <- c("Inland", "pc1", "pc2", "pc3", "pc4")
head(kmtrain)
```

A data.frame: 6 × 5

	Inland	pc1	pc2	pc3	pc4
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
11174	0	0.62155	-0.40828	0.78395	0.02940
1542	0	1.55471	-0.36770	0.39013	1.04186
3537	0	-1.06642	0.42280	-1.37004	0.98605
12747	1	0.23624	0.76784	-0.98573	1.35681
13348	1	-0.87807	0.12463	0.52692	-0.97772
6628	0	-1.51037	-0.42055	-1.56922	0.36121

```
In [75]: # Apply min-max scaling to reduce the effects of the
# more extreme values.

kmtrain_scaled <- apply(kmtrain[, -1], MARGIN=2, range01)
kmtrain_scaled <- as.data.frame(cbind(as.numeric(kmtrain$Inland), kmtrain_scaled),
row.names=row.names(kmtrain))
colnames(kmtrain_scaled) <- colnames(kmtrain)
```

In [76]: *# Get percent of Inland districts in kmtrain.*

```
kmtrain_Inland_percent <- mean(kmtrain$Inland)
round(kmtrain_Inland_percent, 4)
```

0.3187

In [77]: *# Construct random forest model with specific seed.*

*# As we have seen previously, we get a slightly better
model when we ask for the importances.*

```
set.seed(1493)
(rfclf_best <- randomForest(I(as.factor(Inland)) ~ .,
                           data= traindat[, rfclf_columns],
                           ntree=900, mtry= 3, nodesize= 1,
                           importance= TRUE))
```

Call:

```
randomForest(formula = I(as.factor(Inland)) ~ ., data = traindat[, rfclf_columns], n
tree = 900, mtry = 3, nodesize = 1, importance = TRUE)
```

Type of random forest: classification

Number of trees: 900

No. of variables tried at each split: 3

OOB estimate of error rate: 10.39%

Confusion matrix:

	0	1	class.error
0	6497	450	0.064776
1	609	2640	0.187442

In [78]: *# Add a prob01 column to kmtrain.*

```
preds <- predict(rfclf_best, newdata=traindat[, rfclf_columns],
                 type="prob")
```

preds is a matrix with 2 columns.

```
kmtrain_scaled$prob01 <- preds[, 2]
summary(kmtrain_scaled$prob01)
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.00000	0.00889	0.04778	0.31911	0.83778	1.00000	

In [79]: *# Apply our best weights.*

```
wgths <- c(0.22, 0.14, 0.22, 0.14, 0.28)
names(wgths) <- cols <- c(paste0("pc", 1:4), "prob01")

kmtrain_wgths <- t(t(kmtrain_scaled[, cols]) * as.numeric(wgths[cols]))
```

In [80]: *# Run k-means algorithm with number of clusters set to 2.*

```
set.seed(1233)
start <- Sys.time()
kmp1 <- kmeans(kmtrain_wgths, 2, iter.max = 50, nstart = 15)
stop <- Sys.time()
# round(stop - start, 2)
# Time difference of 0.03 secs

print(kmp1$size)
```

[1] 6947 3249

In [81]: *# See how the clusters are associated with Inland.*

```

dfout <- as.data.frame(cbind(kmtrain_scaled$Inland, kmp1$cluster))
colnames(dfout) <- c("Inland", "cluster")
rownames(dfout) <- rownames(kmtrain_scaled)
dat_c1 <- dfout[which(dfout$cluster== 1),]
datc1_Inland_percent <- mean(dat_c1$Inland)

tmpdat <- dfout

c1_to_InlandYES <- FALSE
if(datc1_Inland_percent >= kmtrain_Inland_percent) { c1_to_InlandYES <- TRUE }

if(c1_to_InlandYES) {
  # cluster 1 is associated with the Inland districts
  tmpdat[which(tmpdat$cluster== 1),]$Inland <- 1
  tmpdat[which(tmpdat$cluster== 2),]$Inland <- 0
} else {
  # cluster 2 is associated with the Inland districts
  tmpdat[which(tmpdat$cluster== 2),]$Inland <- 1
  tmpdat[which(tmpdat$cluster== 1),]$Inland <- 0
}

preds <- as.factor(tmpdat$Inland)
names(preds) <- rownames(tmpdat)
ans <- get_confusion(preds, kmtrain_scaled[, "Inland", drop=FALSE])

print(ans$matrix)
''
print(paste("f-score for hybrid model with weights, trainset: ", as.character(ans[[2]]), se

mat <- as.matrix(ans[[1]])
percent_correct <- sum(diag(mat))/floor(sum(mat))
result <- round((0.4 * percent_correct + 0.6 * ans[[2]]), 4)

''
print(paste("Type2 score for hybrid model with weights, trainset: ", as.character(result),
""
print(paste("Accuracy: ", as.character(round(percent_correct,4)), sep=""))

      0      1 class.error
0 6947      0          0
1      0 3249          0

"

[1] "f-score for hybrid model with weights, trainset: 1"

"

[1] "Type2 score for hybrid model with weights, trainset: 1"

"

[1] "Accuracy: 1"

```

Get scores on testset for hybrid model with weights

In [82]: # Add a prob01 column to testdat.

```

kmtest <- testdat

preds <- predict(rfclf_best, newdata=testdat[, rfclf_columns],
                type="prob")

# preds is a matrix with 2 columns.
kmtest$prob01 <- preds[, 2]

```

In [83]: kmp1\$centers

A matrix: 2 × 5 of type dbl

	pc1	pc2	pc3	pc4	prob01
1	0.069316	0.042905	0.044974	0.065446	0.012758
2	0.078962	0.054373	0.041123	0.057571	0.253118

```
In [84]: # Function for obtaining a set of scores on the testset data
# using kmp1 as the classifier.

### NOTE: the value of c1_to_InlandYES from above is being
### used in the following function.

get_testdatScores_kmp1b <- function(seedv, dat) {

  n_smp <- 4000

  seedv_len <- length(seedv)
  datout <- rep(NA, 5 * seedv_len)
  dim(datout) <- c(seedv_len, 5)
  datout <- as.data.frame(datout)
  colnames(datout) <- c("fscore", "Acc", "Type2", "FN", "FP")
  rownames(datout) <- as.character(seedv)

  # Using model kmp1 from above.
  ctr_list <- vector("list", length= n_smp)
  for(i in 1:n_smp) {
    ctr_list[[i]] <- kmp1$centers
  }

  for(h in 1:seedv_len) {
    # shuffle dat
    cur_seed <- seedv[h]
    set.seed(cur_seed)
    # It is expected that dat is testdat, which has 10K rcids
    smp <- sample(rownames(dat), n_smp, replace= FALSE)
    df <- dat[smp,]

    # CAUTION: df has the prob01 column.

    # Using model pca4 (constructed from the training set).
    df_pca <- predict(pca4, df[, km_predictors])
    df02 <- cbind(as.data.frame(df_pca), df$prob01)

    # Apply min-max scaling to df02.
    df02_scaled <- apply(df02, MARGIN=2, range01)
    df02_scaled <- as.data.frame(df02_scaled,
                                row.names=rownames(df))
    colnames(df02_scaled) <- colnames(kmp1$centers)

    # Apply weights to valdat. wgths and cols are defined above.
    df02_wgths <- t(t(df02_scaled[, cols]) * as.numeric(wgths[cols]))

    # Each element of the following list is a row of df.
    df02_asList <- split(df02_wgths[, colnames(kmp1$centers)], seq(n_smp))
    names(ctr_list) <- rownames(df)

    # Get the predictions for df.
    preds <- mcmapply(getCluster, df02_asList, ctr_list,
                     SIMPLIFY=TRUE, mc.cores=6)
    df$cluster <- as.numeric(preds)

    df$pred_Inland <- NA
    if(c1_to_InlandYES) {
      df[which(df$cluster==1),]$pred_Inland <- 1
    }
  }
}
```

```

        df[which(df$cluster==2),]$pred_Inland <- 0
    } else {
        df[which(df$cluster==1),]$pred_Inland <- 0
        df[which(df$cluster==2),]$pred_Inland <- 1
    }

    # Generate confusion matrix.
    preds <- as.factor(df$pred_Inland)
    names(preds) <- rownames(df)
    ans <- get_confusion(preds, df[, "Inland", drop=FALSE])

    mat <- as.matrix(ans[[1]])
    fscore <- round(as.numeric(ans[[2]]), 4)
    acc <- round(sum(diag(mat))/floor(sum(mat)), 4)
    type2 <- round((0.4 * acc + 0.6 * ans[[2]]), 4)
    FN <- as.numeric(mat[2,1])
    FP <- as.numeric(mat[1,2])
    datout[as.character(cur_seed), 1:5] <- c(fscore, acc, type2, FN, FP)
  }
  return(datout)
}

```

In [85]: *# Get kmp1's scores on the testset data. For each of
the 250 seeds, I sample 4K (no replacement) from the 10K
set of testdat records.*

```

set.seed(1821)
seed_vector <- sample(1:9999, 250, replace=FALSE)

start <- Sys.time()
paste("Start time: ", start, sep="")
dat_result <- get_testdatScores_kmplb(seed_vector, kmtest)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 27.6 secs

```

'Start time: 2021-04-27 22:28:57'

Time difference of 27.6 secs

In [86]: `dim(dat_result)`
`head(dat_result)`

250 · 5

A data.frame: 6 × 5

	fscore	Acc	Type2	FN	FP
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
5934	0.8357	0.8968	0.8601	227	186
1953	0.8436	0.8995	0.8660	216	186
7591	0.8320	0.8925	0.8562	260	170
1038	0.8360	0.8985	0.8610	234	172
49	0.8320	0.8920	0.8560	231	201
3203	0.8359	0.8968	0.8603	227	186

```
In [87]: fscore_mean <- round(mean(dat_result$fscore), 4)
fscore_sd <- round(sd(dat_result$fscore), 4)
paste0("fscore mean: ", as.character(fscore_mean))
# 0.8349
paste0("fscore StdDev: ", as.character(fscore_sd))
# 0.006
""
summary(dat_result$fscore)
```

'fscore mean: 0.8349'

'fscore StdDev: 0.006'

"

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.820	0.831	0.835	0.835	0.839	0.852

```
In [88]: Acc_mean <- round(mean(dat_result$Acc), 4)
Acc_sd <- round(sd(dat_result$Acc), 4)
paste0("accuracy mean: ", as.character(Acc_mean))
# 0.8959
paste0("accuracy StdDev: ", as.character(Acc_sd))
# 0.0037
""
summary(dat_result$Acc)
```

'accuracy mean: 0.8959'

'accuracy StdDev: 0.0037'

"

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.886	0.893	0.896	0.896	0.898	0.908

```
In [89]: Type2_mean <- round(mean(dat_result$Type2), 4)
Type2_sd <- round(sd(dat_result$Type2), 4)
paste0("Type2 mean: ", as.character(Type2_mean))
# 0.8593
paste0("Type2 StdDev: ", as.character(Type2_sd))
# 0.005
""
summary(dat_result$Type2)
```

'Type2 mean: 0.8593'

'Type2 StdDev: 0.005'

"

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.847	0.856	0.859	0.859	0.862	0.874

```
In [90]: FN_mean <- round(mean(dat_result$FN), 2)
FN_sd <- round(sd(dat_result$FN), 2)
paste0("FN mean: ", as.character(FN_mean))
# 234.66
paste0("FN StdDev: ", as.character(FN_sd))
# 11.27
""
summary(dat_result$FN)
```

'FN mean: 234.66'

'FN StdDev: 11.27'

"

```

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
      107     137     155     175     190     210

In [91]: FP_mean <- round(mean(dat_result$FP), 2)
FP_sd <- round(sd(dat_result$FP), 2)
paste0("FP mean: ", as.character(FP_mean))
# 181.9
paste0("FP StdDev: ", as.character(FP_sd))
# 10.16
""
summary(dat_result$FP)

```

'FP mean: 181.9'

'FP StdDev: 10.16'

"

```

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
      157     174     183     182     189     210

```

Section 4 Comments

Adding weights to the model terms is not helping. The model with weights is no better than the model without weights (this is the k-means model with prob01 constructed from rfc1f_best and using 4 principal components).

It may be that weights are not helping to improve the model because we have so much data to work with. The scaling, of course, also helps to make the weights less important.

If weights are not helpful, we do not need PCA. In other words, we can include more of the original variables in the model without having to worry much about increasing the computational cost. Similarly, we can also see if adding a second probability column improves performance. (It may be that we get a better model by applying PCA first.)

* * * * *

Section 5: k-means with prob1 & prob2; no pca, no weights

Here I use rfc1f_best for the prob01 column and svm02 for the prob02 column. These are our current best models, as measured by accuracy and f-score. rfc1f_best has the highest numbers. If the prob02 column is too highly correlated with the prob01 column, it will not help to improve the scores.

```

In [29]: kmtrain <- traindat[, km_columns]
kmtest <- testdat[, km_columns]
colnames(kmtrain)

```

'Inland' · 'housing_median_age' · 'median_house_value' · 'HHdens_In' · 'rooms_per_hh' · 'bdrms_per_room' · 'pop_per_hh'

```

In [22]: # Random forest model for the prob01 column:

```

```

set.seed(1493)
(rfc1f_best <- randomForest(I(as.factor(Inland)) ~ .,
                             data= traindat[, rfc1f_columns],
                             ntree=900,
                             mtry= 3, nodesize= 1, importance=TRUE))

```

```
Call:
randomForest(formula = I(as.factor(Inland)) ~ ., data = traindat[,      rfclf_columns], n
tree = 900, mtry = 3, nodesize = 1, importance = TRUE)
Type of random forest: classification
Number of trees: 900
No. of variables tried at each split: 3
```

In [30]: # Add prob01 column.

```
preds01 <- predict(rfclf_best, newdata= traindat[, rfclf_columns], type="prob")

# preds01 is a matrix with 2 columns.
kmtrain$prob01 <- preds01[, 2]
summary(kmtrain$prob01)

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
0.000000 0.00889 0.04778 0.31911 0.83778 1.00000
```

In [31]: # Scale data for svm model.

```
svmtrain <- traindat[, svm02_columns]
svmtrain$median_house_value <- log(svmtrain$median_house_value)

svm_scaled <- scale(svmtrain[, -1])
svm_centers <- attr(svm_scaled, "scaled:center")
svm_scales <- attr(svm_scaled, "scaled:scale")
svm_scaled <- as.data.frame(cbind(as.numeric(svmtrain$Inland), svm_scaled),
                             row.names=rownames(svmtrain))
colnames(svm_scaled) <- colnames(svmtrain)
```

In [32]: # SVM model for the prob02 column:

```
svm02 <- svm(I(as.factor(Inland)) ~ ., data=svm_scaled, kernel="radial",
             gamma=0.06, cost=80, scale=FALSE, probability=TRUE)
```

In [33]: # Construct prob02 column.

```
preds02 <- predict(svm02, newdata=svm_scaled, probability=TRUE)
kmtrain$prob02 <- as.numeric(attr(preds02, "probabilities")[, 2])
summary(kmtrain$prob02)

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
0.00000 0.0475 0.1017 0.3192 0.6861 0.9999
```

In [34]: # Initial scaling of kmtrain:

```
kmtrain_scaled <- scale(kmtrain[, -1])
centers <- attr(kmtrain_scaled, "scaled:center")
scales <- attr(kmtrain_scaled, "scaled:scale")
```

In [35]: # Apply min-max scaling to reduce the effects of the
more extreme values.

```
kmtrain_minmax <- apply(kmtrain_scaled, MARGIN=2, range01)
kmtrain_minmax <- as.data.frame(cbind(as.numeric(kmtrain$Inland), kmtrain_minmax),
                                row.names=rownames(kmtrain))
colnames(kmtrain_minmax) <- colnames(kmtrain)
```

In [36]: # Run k-means algorithm with number of clusters set to 2.

```
set.seed(1233)
start <- Sys.time()
kmp1p2 <- kmeans(kmtrain_minmax[, -1], 2, iter.max = 50, nstart = 15)
stop <- Sys.time()
```



```
round(stop - start, 2)
# Time difference of 0.15 secs
```

```
print(kmplp2$size)
```

```
Time difference of 0.05 secs
```

```
[1] 7008 3188
```

In [37]: *# Get percent of Inland districts in kmtrain.*

```
kmtrain_Inland_percent <- mean(kmtrain$Inland)
round(kmtrain_Inland_percent, 4)
```

```
0.3187
```

```

In [38]: # See how the clusters are associated with Inland.

dfout <- as.data.frame(cbind(kmtrain$Inland, kmp1p2$cluster))
colnames(dfout) <- c("Inland", "cluster")
rownames(dfout) <- rownames(kmtrain)
dat_c1 <- dfout[which(dfout$cluster== 1),]
datc1_Inland_percent <- mean(dat_c1$Inland)

tmpdat <- dfout

c1_to_InlandYES <- FALSE
if(datc1_Inland_percent >= kmtrain_Inland_percent) { c1_to_InlandYES <- TRUE }

if(c1_to_InlandYES) {
  # cluster 1 is associated with the Inland districts
  tmpdat[which(tmpdat$cluster== 1),]$Inland <- 1
  tmpdat[which(tmpdat$cluster== 2),]$Inland <- 0
} else {
  # cluster 2 is associated with the Inland districts
  tmpdat[which(tmpdat$cluster== 2),]$Inland <- 1
  tmpdat[which(tmpdat$cluster== 1),]$Inland <- 0
}

# Generate confusion matrix for the k-means clusters and
# the corresponding f-score.
preds <- as.factor(tmpdat$Inland)
names(preds) <- rownames(tmpdat)
ans <- get_confusion(preds, kmtrain[, "Inland", drop=FALSE])

print(ans$matrix)
''

print(paste("f-score for kmp1p2, train set: ", as.character(ans[[2]]), sep=""))

mat <- as.matrix(ans[[1]])
percent_correct <- sum(diag(mat))/floor(sum(mat))
result <- round((0.4 * percent_correct + 0.6 * ans[[2]]), 4)

''

print(paste("Type2 score for kmp1p2, train set: ", as.character(result), sep=""))

""

print(paste("Accuracy: ", as.character(round(percent_correct,4)), sep=""))
# [1] "Accuracy: 0.5867"

      0      1 class.error
0 6722  225      0.0324
1  286 2963      0.0880

"

[1] "f-score for kmp1p2, train set: 0.9206"

"

[1] "Type2 score for kmp1p2, train set: 0.9323"

"

[1] "Accuracy: 0.9499"

```

Get scores on testset for kmp1p2

```

In [39]: # SVM scaling for testdat records.

svmtest <- testdat[, svm02_columns]
svmtest$median_house_value <- log(svmtest$median_house_value)

svmtest_scaled <- scale(svmtest[, -1], center= svm_centers, scale= svm_scales)
svmtest_scaled <- as.data.frame(cbind(as.numeric(svmtest$Inland), svmtest_scaled),

```

```

row.names=rownames(svmtest))
colnames(svmtest_scaled) <- colnames(svmtest)

```

```

In [40]: # Add the prob01 and prob02 columns to testdat. For prob02
# I am cheating a bit because the svm scaling is over all of
# testdat. A more rigorous treatment would require separate
# scaling and predictions for each set of 4K records sampled
# from testdat in the get_testdatScores_kmplp2 function below.

kmtest <- testdat[, km_columns]

preds01 <- predict(rfclf_best, newdata=testdat[, rfclf_columns],
                  type="prob")
preds02 <- predict(svm02, newdata=svmtest_scaled, probability=TRUE)

kmtest$prob01 <- preds01[, 2]
kmtest$prob02 <- as.numeric(attr(preds02, "probabilities")[, 2])

```

```

In [41]: kmplp2$centers

```

A matrix: 2 × 8 of type dbl

	housing_median_age	median_house_value	HHdens_In	rooms_per_hh	bdrms_per_room	pop_per_hh	prob01	prob02
1	0.52657	0.34810	0.70255	0.11979	0.24111	0.12341	0.063744	0.095232
2	0.39732	0.14558	0.54147	0.13894	0.20998	0.12386	0.880459	0.811724

```

In [42]: # Function for obtaining a set of scores on the testset data
# using kmplp2 as the classifier.

### NOTE: the value of c1_to_InlandYES from above is being
### used in the following function.

get_testdatScores_kmplp2 <- function(seedv, dat) {

  n_smp <- 4000

  seedv_len <- length(seedv)
  datout <- rep(NA, 5 * seedv_len)
  dim(datout) <- c(seedv_len, 5)
  datout <- as.data.frame(datout)
  colnames(datout) <- c("fscore", "Acc", "Type2", "FN", "FP")
  rownames(datout) <- as.character(seedv)

  # Using model kmplp2 from above.
  ctr_list <- vector("list", length= n_smp)
  for(i in 1:n_smp) {
    ctr_list[[i]] <- kmplp2$centers
  }

  for(h in 1:seedv_len) {
    # shuffle dat
    cur_seed <- seedv[h]
    set.seed(cur_seed)
    # It is expected that dat is testdat, which has 10K rcds
    smp <- sample(rownames(dat), n_smp, replace= FALSE)
    df <- dat[smp,]

    # CAUTION: df has the prob01 & prob02 columns.

    # Scale df using the centers and scales from kmtrain_scaled.
    df_scaled <- scale(df[, -1], center=centers, scale=scales)

    # Apply min-max scaling to df_scaled.
    df02_scaled <- apply(df_scaled, MARGIN=2, range01)
    df02_scaled <- as.data.frame(df02_scaled,

```

```

                                row.names=rownames(df))
colnames(df02_scaled) <- colnames(kmplp2$centers)

# Each element of the following list is a row of df.
df02_asList <- split(df02_scaled[, colnames(kmplp2$centers)], seq(n_smp))
names(ctr_list) <- rownames(df)

# Get the predictions for df.
preds <- mcmapply(getCluster, df02_asList, ctr_list,
                  SIMPLIFY=TRUE, mc.cores=6)
df$cluster <- as.numeric(preds)

df$pred_Inland <- NA
if(c1_to_InlandYES) {
  df[which(df$cluster==1),]$pred_Inland <- 1
  df[which(df$cluster==2),]$pred_Inland <- 0
} else {
  df[which(df$cluster==1),]$pred_Inland <- 0
  df[which(df$cluster==2),]$pred_Inland <- 1
}

# Generate confusion matrix.
preds <- as.factor(df$pred_Inland)
names(preds) <- rownames(df)
ans <- get_confusion(preds, df[, "Inland", drop=FALSE])

mat <- as.matrix(ans[[1]])
fscore <- round(as.numeric(ans[[2]]), 4)
acc <- round(sum(diag(mat))/floor(sum(mat)), 4)
type2 <- round((0.4 * acc + 0.6 * ans[[2]]), 4)
FN <- as.numeric(mat[2,1])
FP <- as.numeric(mat[1,2])
datout[as.character(cur_seed), 1:5] <- c(fscore, acc, type2, FN, FP)
}
return(datout)
}

```

In [43]: # Get kmplp2's scores on the testset data. For each of
the 250 seeds, I sample 4K (no replacement) from the 10K
set of testdat records.

```

set.seed(1821)
seed_vector <- sample(1:9999, 250, replace=FALSE)

start <- Sys.time()
paste("Start time: ", start, sep="")
dat_result <- get_testdatScores_kmplp2(seed_vector, kmtest)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 10.41 mins

```

'Start time: 2021-04-30 08:58:41'

Time difference of 10.6 mins

In [44]: fscore_mean <- round(mean(dat_result\$fscore), 4)
fscore_sd <- round(sd(dat_result\$fscore), 4)
paste0("fscore mean: ", as.character(fscore_mean))
0.8346
paste0("fscore StdDev: ", as.character(fscore_sd))
0.0059
""
summary(dat_result\$fscore)

'fscore mean: 0.8346'

'fscore StdDev: 0.0059'

"

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.819	0.831	0.834	0.835	0.838	0.849

```
In [45]: Acc_mean <- round(mean(dat_result$Acc), 4)
Acc_sd <- round(sd(dat_result$Acc), 4)
paste0("accuracy mean: ", as.character(Acc_mean))
# 0.8962
paste0("accuracy StdDev: ", as.character(Acc_sd))
# 0.0037
""
summary(dat_result$Acc)
```

'accuracy mean: 0.8962'

'accuracy StdDev: 0.0037'

"

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.887	0.894	0.896	0.896	0.898	0.907

```
In [46]: Type2_mean <- round(mean(dat_result$Type2), 4)
Type2_sd <- round(sd(dat_result$Type2), 4)
paste0("Type2 mean: ", as.character(Type2_mean))
# 0.8593
paste0("Type2 StdDev: ", as.character(Type2_sd))
# 0.0049
""
summary(dat_result$Type2)
```

'Type2 mean: 0.8593'

'Type2 StdDev: 0.0049'

"

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.846	0.856	0.859	0.859	0.863	0.872

```
In [47]: FN_mean <- round(mean(dat_result$FN), 2)
FN_sd <- round(sd(dat_result$FN), 2)
paste0("FN mean: ", as.character(FN_mean))
# 240.4
paste0("FN StdDev: ", as.character(FN_sd))
# 11.4
""
summary(dat_result$FN)
```

'FN mean: 240.35'

'FN StdDev: 11.38'

"

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
206	233	241	240	248	272

```
In [48]: FP_mean <- round(mean(dat_result$FP), 2)
FP_sd <- round(sd(dat_result$FP), 2)
paste0("FP mean: ", as.character(FP_mean))
# 174.7
paste0("FP StdDev: ", as.character(FP_sd))
# 9.7
""
summary(dat_result$FP)
```

'FP mean: 174.72'

'FP StdDev: 9.7'

"

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
147	168	175	175	181	200

Section 5 Comments

The svm02 probabilities are neither helping nor hurting the k-means hybrid model. If I had used PCA on the 6 predictors I might have gotten slightly better results, or slightly worse results; I do not think the difference would have been great.

Section 6: different approach using 2 prob columns

Here I continue to use 2 probability columns, one from rfclf_best, the other from svm02. But I apply PCA to the 6 predictors and the svm02 probabilities. I then add the probabilities from rfclf_best to the model.

The reason for this approach: we saw in Part 2 that the svm02 probabilities did not mix well with those from the random forest model. Although the dataset in Part 2 is completely different, something similar might be going on here (any complementary information might be nullified by conflicting information). But the svm02 model is a better model than either f03 or gbclf_best, so if I am going to use a second set of probabilities, I want those from svm02. If we transform the prob02 column using PCA, we might reduce any conflict existing between the prob01 and prob02 probabilities.

```
In [49]: kmtrain <- traindat[, km_columns]
kmtest <- testdat[, km_columns]
colnames(kmtrain)
```

```
'Inland' · 'housing_median_age' · 'median_house_value' · 'HHdens_In' · 'rooms_per_hh' · 'bdrms_per_room' ·
'pop_per_hh'
```

```
In [19]: # Random forest model for the prob01 column:
```

```
set.seed(1493)
(rfclf_best <- randomForest(I(as.factor(Inland)) ~ .,
                             data= traindat[, rfclf_columns],
                             ntree=900,
                             mtry= 3, nodesize= 1, importance=TRUE))
```

Call:

```
randomForest(formula = I(as.factor(Inland)) ~ ., data = traindat[, rfclf_columns], n
tree = 900, mtry = 3, nodesize = 1, importance = TRUE)
```

Type of random forest: classification

Number of trees: 900

No. of variables tried at each split: 3

OOB estimate of error rate: 10.39%

Confusion matrix:

	0	1	class.error
0	6497	450	0.064776
1	609	2640	0.187442

```
In [50]: # Scale data for svm model.
```

```
svmtrain <- traindat[, svm02_columns]
svmtrain$median_house_value <- log(svmtrain$median_house_value)

svm_scaled <- scale(svmtrain[, -1])
svm_centers <- attr(svm_scaled, "scaled:center")
svm_scales <- attr(svm_scaled, "scaled:scale")
svm_scaled <- as.data.frame(cbind(as.numeric(svmtrain$Inland), svm_scaled),
                             row.names=rownames(svmtrain))
colnames(svm_scaled) <- colnames(svmtrain)
```

In [51]: *# SVM model for the prob02 column:*

```
svm02 <- svm(I(as.factor(Inland)) ~ ., data=svm_scaled, kernel="radial",
             gamma=0.06, cost=80, scale=FALSE, probability=TRUE)
```

In [52]: *# Construct prob02 column.*

```
preds02 <- predict(svm02, newdata=svm_scaled, probability=TRUE)
kmtrain$prob02 <- as.numeric(attr(preds02, "probabilities")[, 2])
summary(kmtrain$prob02)
```

```
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
0.0000  0.0479  0.1024  0.3196  0.6865  0.9999
```

In [53]: *# Apply PCA. Use 5 of the 7 components.*

```
pca <- prcomp(kmtrain[, c(km_predictors, "prob02")], center=TRUE, scale.=TRUE,
              rank.=5, retx=TRUE)
summary(pca)
```

```
Importance of first k=5 (out of 7) components:
              PC1  PC2  PC3  PC4  PC5
Standard deviation    1.517 1.321 0.993 0.896 0.7397
Proportion of Variance 0.329 0.249 0.141 0.115 0.0782
Cumulative Proportion 0.329 0.578 0.719 0.834 0.9118
```

In [54]: *# Construct training set data for k-means modeling*
with PCA.

```
kmtrain02 <- cbind(as.numeric(traindat$Inland), as.data.frame(pca$x))
rownames(kmtrain02) <- rownames(traindat)
colnames(kmtrain02) <- c("Inland", "pc1", "pc2", "pc3", "pc4", "pc5")
head(kmtrain02)
```

A data.frame: 6 × 6

	Inland	pc1	pc2	pc3	pc4	pc5
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
11174	0	-0.050533	-0.92857	-0.762667	-0.18391	0.521402
1542	0	0.599497	-1.52889	-0.716347	0.95981	0.414506
3537	0	-1.037601	0.70059	0.893106	1.49712	0.152263
12747	1	1.167845	0.87609	0.734401	1.19491	1.266115
13348	1	-0.256596	0.87814	-0.090126	-1.32222	0.581090
6628	0	-1.620032	0.42692	1.467853	0.77186	-0.096807

In [55]: *# Apply min-max scaling.*

```
kmtrain02_scaled <- apply(kmtrain02[, -1], MARGIN=2, range01)
kmtrain02_scaled <- as.data.frame(cbind(kmtrain02$Inland, kmtrain02_scaled),
                                  row.names=row.names(kmtrain02))
colnames(kmtrain02_scaled) <- colnames(kmtrain02)
head(kmtrain02_scaled)
```

A data.frame: 6 × 6

	Inland	pc1	pc2	pc3	pc4	pc5
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
11174	0	0.33699	0.53104	0.80673	0.45260	0.68716
1542	0	0.37827	0.48441	0.80902	0.63860	0.67339

	Inland	pc1	pc2	pc3	pc4	pc5
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
3537	0	0.27430	0.65758	0.88845	0.72598	0.63961
19747	1	0.41427	0.67121	0.88062	0.67682	0.78200

In [56]: # Add prob01 column.

```
preds01 <- predict(rfc1f_best, newdata= traindat[, rfc1f_columns], type="prob")

# preds01 is a matrix with 2 columns.
kmtrain02_scaled$prob01 <- preds01[, 2]
summary(kmtrain02_scaled$prob01)
```

```
      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
0.000000 0.00889  0.04778  0.31911  0.83778  1.00000
```

In [57]: # Run k-means algorithm with number of clusters set to 2.

```
set.seed(1233)
start <- Sys.time()
kmp1p2_b <- kmeans(kmtrain02_scaled[, -1], 2, iter.max = 50, nstart = 15)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 0.15 secs

print(kmp1p2_b$size)
```

Time difference of 0.03 secs

```
[1] 6947 3249
```

In [58]: # Get percent of Inland districts in kmtrain.

```
kmtrain_Inland_percent <- mean(kmtrain$Inland)
round(kmtrain_Inland_percent, 4)
```

```
0.3187
```

In [59]: # See how the clusters are associated with Inland.

```
dfout <- as.data.frame(cbind(kmtrain$Inland, kmp1p2_b$cluster))
colnames(dfout) <- c("Inland", "cluster")
rownames(dfout) <- rownames(kmtrain)
dat_c1 <- dfout[which(dfout$cluster== 1),]
datc1_Inland_percent <- mean(dat_c1$Inland)

tmpdat <- dfout

c1_to_InlandYES <- FALSE
if(datc1_Inland_percent >= kmtrain_Inland_percent) { c1_to_InlandYES <- TRUE }

if(c1_to_InlandYES) {
  # cluster 1 is associated with the Inland districts
  tmpdat[which(tmpdat$cluster== 1),]$Inland <- 1
  tmpdat[which(tmpdat$cluster== 2),]$Inland <- 0
} else {
  # cluster 2 is associated with the Inland districts
  tmpdat[which(tmpdat$cluster== 2),]$Inland <- 1
  tmpdat[which(tmpdat$cluster== 1),]$Inland <- 0
}

# Generate confusion matrix for the k-means clusters and
# the corresponding f-score.
preds <- as.factor(tmpdat$Inland)
names(preds) <- rownames(tmpdat)
ans <- get_confusion(preds, kmtrain[, "Inland", drop=FALSE])
```



```

print(ans$matrix)
''
print(paste("f-score for kmp1p2_b, train set: ", as.character(ans[[2]]), sep=""))

mat <- as.matrix(ans[[1]])
percent_correct <- sum(diag(mat))/floor(sum(mat))
result <- round((0.4 * percent_correct + 0.6 * ans[[2]]), 4)

''
print(paste("Type2 score for kmp1p2_b, train set: ", as.character(result), sep=""))

""
print(paste("Accuracy: ", as.character(round(percent_correct,4)), sep=""))

      0      1 class.error
0 6947      0           0
1      0 3249           0

"

[1] "f-score for kmp1p2_b, train set: 1"

"

[1] "Type2 score for kmp1p2_b, train set: 1"

"

[1] "Accuracy: 1"

```

Get scores on testset for kmp1p2_b

In [60]: *# SVM scaling for testdat records.*

```

svmtest <- testdat[, svm02_columns]
svmtest$median_house_value <- log(svmtest$median_house_value)

svmtest_scaled <- scale(svmtest[, -1], center= svm_centers, scale= svm_scales)
svmtest_scaled <- as.data.frame(cbind(as.numeric(svmtest$Inland), svmtest_scaled),
                                row.names=row.names(svmtest))
colnames(svmtest_scaled) <- colnames(svmtest)

```

In [61]: *# Add the prob02 column to kmtest.*

```

kmtest <- testdat[, km_columns]
preds02 <- predict(svm02, newdata=svmtest_scaled, probability=TRUE)
kmtest$prob02 <- as.numeric(attr(preds02, "probabilities")[, 2])

```

In [62]: *# Add the prob01 column to kmtest.*

```

preds01 <- predict(rfclf_best, newdata=testdat[, rfclf_columns],
                  type="prob")
kmtest$prob01 <- preds01[, 2]

```

In [63]: kmp1p2_b\$centers

A matrix: 2 x 6 of type dbl

	pc1	pc2	pc3	pc4	pc5	prob01
1	0.29971	0.57830	0.83920	0.50138	0.61212	0.045566
2	0.42677	0.65632	0.85543	0.44215	0.63683	0.903993

In [68]: *# Function for obtaining a set of scores on the testset data
using kmp1p2_b as the classifier.*

```

### NOTE: the value of c1_to_InlandYES from above is being
### used in the following function.

get_testdatScores_kmplp2_b <- function(seedv, dat) {

  n_smp <- 4000

  seedv_len <- length(seedv)
  datout <- rep(NA, 5 * seedv_len)
  dim(datout) <- c(seedv_len, 5)
  datout <- as.data.frame(datout)
  colnames(datout) <- c("fscore", "Acc", "Type2", "FN", "FP")
  rownames(datout) <- as.character(seedv)

  # Using model kmplp2_b from above.
  ctr_list <- vector("list", length= n_smp)
  for(i in 1:n_smp) {
    ctr_list[[i]] <- kmplp2_b$centers
  }

  for(h in 1:seedv_len) {
    # shuffle dat
    cur_seed <- seedv[h]
    set.seed(cur_seed)
    # It is expected that dat is testdat, which has 10K rcds
    smp <- sample(rownames(dat), n_smp, replace= FALSE)
    df <- dat[smp,]

    # CAUTION: df has the prob01 & prob02 columns.

    # Apply PCA.
    pca_columns <- c(km_predictors, "prob02")
    df_pca <- predict(pca, df[, pca_columns])

    # Apply min-max scaling to df_scaled.
    df_scaled <- apply(df_pca, MARGIN=2, range01)
    df_scaled <- as.data.frame(cbind(df_scaled, df$prob01),
                              row.names=rownames(df))
    colnames(df_scaled) <- c(paste0("pc", 1:5), "prob01")

    # Each element of the following list is a row of df.
    df_asList <- split(df_scaled[, colnames(kmplp2_b$centers)], seq(n_smp))
    names(ctr_list) <- rownames(df)

    # Get the predictions for df.
    preds <- mcmapply(getCluster, df_asList, ctr_list,
                      SIMPLIFY=TRUE, mc.cores=6)
    df$cluster <- as.numeric(preds)

    df$pred_Inland <- NA
    if(c1_to_InlandYES) {
      df[which(df$cluster==1),]$pred_Inland <- 1
      df[which(df$cluster==2),]$pred_Inland <- 0
    } else {
      df[which(df$cluster==1),]$pred_Inland <- 0
      df[which(df$cluster==2),]$pred_Inland <- 1
    }

    # Generate confusion matrix.
    preds <- as.factor(df$pred_Inland)
    names(preds) <- rownames(df)
    ans <- get_confusion(preds, df[, "Inland", drop=FALSE])

    mat <- as.matrix(ans[[1]])
    fscore <- round(as.numeric(ans[[2]]), 4)
    acc <- round(sum(diag(mat))/floor(sum(mat)), 4)
    type2 <- round((0.4 * acc + 0.6 * ans[[2]]), 4)
    FN <- as.numeric(mat[2,1])
    FP <- as.numeric(mat[1,2])
  }
}

```

```

        datout[as.character(cur_seed), 1:5] <- c(fscore,acc,type2,FN,FP)
      }
      return(datout)
    }
  }
}

```

In [71]: *# Get kmp1p2_b's scores on the testset data. For each of
the 250 seeds, I sample 4K (no replacement) from the 10K
set of testdat records.*

```

set.seed(1821)
seed_vector <- sample(1:9999, 250, replace=FALSE)

start <- Sys.time()
paste("Start time: ", start, sep="")
dat_result <- get_testdatScores_kmp1p2_b(seed_vector, kmtest)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 8.95 mins

```

'Start time: 2021-04-30 09:41:09'

Time difference of 8.95 mins

In [72]: `dim(dat_result)`
`head(dat_result)`

250 · 5

A data.frame: 6 × 5

	fscore	Acc	Type2	FN	FP
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
5934	0.8365	0.8972	0.8608	226	185
1953	0.8456	0.9010	0.8678	216	180
7591	0.8323	0.8930	0.8566	263	165
1038	0.8357	0.8982	0.8607	234	173
49	0.8340	0.8935	0.8578	231	195
3203	0.8360	0.8970	0.8604	229	183

In [73]: `fscore_mean <- round(mean(dat_result$fscore), 4)`
`fscore_sd <- round(sd(dat_result$fscore), 4)`
`paste0("fscore mean: ", as.character(fscore_mean))`
0.8357
`paste0("fscore StdDev: ", as.character(fscore_sd))`
0.0061
" "
`summary(dat_result$fscore)`

'fscore mean: 0.8357'

'fscore StdDev: 0.0061'

"

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.820	0.832	0.836	0.836	0.839	0.851

In [74]: `Acc_mean <- round(mean(dat_result$Acc), 4)`
`Acc_sd <- round(sd(dat_result$Acc), 4)`
`paste0("accuracy mean: ", as.character(Acc_mean))`
0.8965
`paste0("accuracy StdDev: ", as.character(Acc_sd))`
0.0037
" "
`summary(dat_result$Acc)`

```
'accuracy mean: 0.8965'
```

```
'accuracy StdDev: 0.0037'
```

```
"
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.887	0.894	0.896	0.896	0.899	0.907

```
In [75]: Type2_mean <- round(mean(dat_result$Type2), 4)
Type2_sd <- round(sd(dat_result$Type2), 4)
paste0("Type2 mean: ", as.character(Type2_mean))
# 0.8600
paste0("Type2 StdDev: ", as.character(Type2_sd))
# 0.005
""
summary(dat_result$Type2)
```

```
'Type2 mean: 0.86'
```

```
'Type2 StdDev: 0.005'
```

```
"
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.847	0.857	0.860	0.860	0.863	0.874

```
In [76]: FN_mean <- round(mean(dat_result$FN), 2)
FN_sd <- round(sd(dat_result$FN), 2)
paste0("FN mean: ", as.character(FN_mean))
# 234.6
paste0("FN StdDev: ", as.character(FN_sd))
# 11.5
""
summary(dat_result$FN)
```

```
'FN mean: 234.62'
```

```
'FN StdDev: 11.47'
```

```
"
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
196	227	235	235	243	263

```
In [77]: FP_mean <- round(mean(dat_result$FP), 2)
FP_sd <- round(sd(dat_result$FP), 2)
paste0("FP mean: ", as.character(FP_mean))
# 179.5
paste0("FP StdDev: ", as.character(FP_sd))
# 10.3
""
summary(dat_result$FP)
```

```
'FP mean: 179.48'
```

```
'FP StdDev: 10.28'
```

```
"
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
154	171	180	179	187	204

Section 6 Comments

The scores have improved with this approach, but we still do not have a better model than `rfclf_best`. The F-statistic for the difference in accuracy scores between this approach and `rfclf_best` is 4.23. So there is still a significant difference between

the two models, meaning that rfcf_best is a slightly better model.

Section 7: Use Section 6 approach, but with 4 principal components

With 4 components, rather than 5, there is less of a chance of overfitting the training set. A model with only 4 components therefore might generalize better to the testset data.

```
In [78]: kmtrain <- traindat[, km_columns]
kmtest <- testdat[, km_columns]
colnames(kmtrain)
```

```
'Inland' · 'housing_median_age' · 'median_house_value' · 'HHdens_In' · 'rooms_per_hh' · 'bdrms_per_room' ·
'pop_per_hh'
```

```
In [19]: # Random forest model for the prob01 column:
```

```
set.seed(1493)
(rfcf_best <- randomForest(I(as.factor(Inland)) ~ .,
                           data= traindat[, rfcf_columns],
                           ntree=900,
                           mtry= 3, nodesize= 1, importance=TRUE))
```

Call:

```
randomForest(formula = I(as.factor(Inland)) ~ ., data = traindat[, rfcf_columns], n
tree = 900, mtry = 3, nodesize = 1, importance = TRUE)
```

Type of random forest: classification

Number of trees: 900

No. of variables tried at each split: 3

OOB estimate of error rate: 10.39%

Confusion matrix:

	0	1	class.error
0	6497	450	0.064776
1	609	2640	0.187442

```
In [79]: # Scale data for svm model.
```

```
svmtrain <- traindat[, svm02_columns]
svmtrain$median_house_value <- log(svmtrain$median_house_value)

svm_scaled <- scale(svmtrain[, -1])
svm_centers <- attr(svm_scaled, "scaled:center")
svm_scales <- attr(svm_scaled, "scaled:scale")
svm_scaled <- as.data.frame(cbind(as.numeric(svmtrain$Inland), svm_scaled),
                           row.names=rownames(svmtrain))
colnames(svm_scaled) <- colnames(svmtrain)
```

```
In [80]: # SVM model for the prob02 column:
```

```
svm02 <- svm(I(as.factor(Inland)) ~ ., data=svm_scaled, kernel="radial",
             gamma=0.06, cost=80, scale=FALSE, probability=TRUE)
```

```
In [83]: # Construct prob02 column.
```

```
preds02 <- predict(svm02, newdata=svm_scaled, probability=TRUE)
kmtrain$prob02 <- as.numeric(attr(preds02, "probabilities")[, 2])
summary(kmtrain$prob02)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.0000	0.0479	0.1024	0.3196	0.6865	0.9999

In [84]: *# Apply PCA. Use 4 of the 7 components.*

```
pca4 <- prcomp(kmtrain[, c(km_predictors,"prob02")], center=TRUE, scale.=TRUE,
               rank.= 4, retx=TRUE)
summary(pca4)
```

Importance of first k=4 (out of 7) components:

	PC1	PC2	PC3	PC4
Standard deviation	1.517	1.321	0.993	0.896
Proportion of Variance	0.329	0.249	0.141	0.115
Cumulative Proportion	0.329	0.578	0.719	0.834

In [85]: *# Construct training set data for k-means modeling
with PCA.*

```
kmtrain02 <- cbind(as.numeric(traindat$Inland), as.data.frame(pca4$x))
rownames(kmtrain02) <- rownames(traindat)
colnames(kmtrain02) <- c("Inland", "pc1", "pc2", "pc3", "pc4")
head(kmtrain02)
```

A data.frame: 6 × 5

	Inland	pc1	pc2	pc3	pc4
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
11174	0	-0.050533	-0.92857	-0.762667	-0.18391
1542	0	0.599497	-1.52889	-0.716347	0.95981
3537	0	-1.037601	0.70059	0.893106	1.49712
12747	1	1.167845	0.87609	0.734401	1.19491
13348	1	-0.256596	0.87814	-0.090126	-1.32222
6628	0	-1.620032	0.42692	1.467853	0.77186

In [86]: *# Apply min-max scaling.*

```
kmtrain02_scaled <- apply(kmtrain02[, -1], MARGIN=2, range01)
kmtrain02_scaled <- as.data.frame(cbind(kmtrain02$Inland, kmtrain02_scaled),
                                row.names=row.names(kmtrain02))
colnames(kmtrain02_scaled) <- colnames(kmtrain02)
head(kmtrain02_scaled)
```

A data.frame: 6 × 5

	Inland	pc1	pc2	pc3	pc4
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
11174	0	0.33699	0.53104	0.80673	0.45260
1542	0	0.37827	0.48441	0.80902	0.63860
3537	0	0.27430	0.65758	0.88845	0.72598
12747	1	0.41437	0.67121	0.88062	0.67683
13348	1	0.32390	0.67137	0.83993	0.26748
6628	0	0.23731	0.63632	0.91682	0.60803

In [87]: *# Add prob01 column.*

```
preds01 <- predict(rfc1f_best, newdata= traindat[, rfc1f_columns], type="prob")
# preds01 is a matrix with 2 columns.
kmtrain02_scaled$prob01 <- preds01[, 2]
summary(kmtrain02_scaled$prob01)
```

```

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
a  aaaaa a  aaaaa a  aaaaa a  aaaaa a  aaaaa a  aaaaa

```

In [88]: *# Run k-means algorithm with number of clusters set to 2.*

```

set.seed(1233)
start <- Sys.time()
kmp1p2_b <- kmeans(kmtrain02_scaled[, -1], 2, iter.max = 50, nstart = 15)
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 0.15 secs

print(kmp1p2_b$size)

```

Time difference of 0.04 secs

```
[1] 6947 3249
```

In [89]: *# Get percent of Inland districts in kmtrain.*

```

kmtrain_Inland_percent <- mean(kmtrain$Inland)
round(kmtrain_Inland_percent, 4)

```

```
0.3187
```

In [90]: *# See how the clusters are associated with Inland.*

```

dfout <- as.data.frame(cbind(kmtrain$Inland, kmp1p2_b$cluster))
colnames(dfout) <- c("Inland", "cluster")
rownames(dfout) <- rownames(kmtrain)
dat_c1 <- dfout[which(dfout$cluster== 1),]
datc1_Inland_percent <- mean(dat_c1$Inland)

tmpdat <- dfout

c1_to_InlandYES <- FALSE
if(datc1_Inland_percent >= kmtrain_Inland_percent) { c1_to_InlandYES <- TRUE }

if(c1_to_InlandYES) {
  # cluster 1 is associated with the Inland districts
  tmpdat[which(tmpdat$cluster== 1),]$Inland <- 1
  tmpdat[which(tmpdat$cluster== 2),]$Inland <- 0
} else {
  # cluster 2 is associated with the Inland districts
  tmpdat[which(tmpdat$cluster== 2),]$Inland <- 1
  tmpdat[which(tmpdat$cluster== 1),]$Inland <- 0
}

# Generate confusion matrix for the k-means clusters and
# the corresponding f-score.
preds <- as.factor(tmpdat$Inland)
names(preds) <- rownames(tmpdat)
ans <- get_confusion(preds, kmtrain[, "Inland", drop=FALSE])

print(ans$matrix)
''
print(paste("f-score for kmp1p2_b, train set: ", as.character(ans[[2]]), sep=""))

mat <- as.matrix(ans[[1]])
percent_correct <- sum(diag(mat))/floor(sum(mat))
result <- round((0.4 * percent_correct + 0.6 * ans[[2]]), 4)

''
print(paste("Type2 score for kmp1p2_b, train set: ", as.character(result), sep=""))
'''
print(paste("Accuracy: ", as.character(round(percent_correct,4)), sep=""))

```

```

      0      1 class.error
A 6047      A      A

[1] "f-score for kmp1p2_b, train set: 1"
"

[1] "Type2 score for kmp1p2_b, train set: 1"
"

[1] "Accuracy: 1"

```

Get scores on testset for kmp1p2_b

In [91]: *# SVM scaling for testdat records.*

```

svmtest <- testdat[, svm02_columns]
svmtest$median_house_value <- log(svmtest$median_house_value)

svmtest_scaled <- scale(svmtest[, -1], center= svm_centers, scale= svm_scales)
svmtest_scaled <- as.data.frame(cbind(as.numeric(svmtest$Inland), svmtest_scaled),
                                row.names=row.names(svmtest))
colnames(svmtest_scaled) <- colnames(svmtest)

```

In [92]: *# Add the prob02 column to kmtest.*

```

kmtest <- testdat[, km_columns]
preds02 <- predict(svm02, newdata=svmtest_scaled, probability=TRUE)
kmtest$prob02 <- as.numeric(attr(preds02, "probabilities")[, 2])

```

In [93]: *# Add the prob01 column to kmtest.*

```

preds01 <- predict(rfc1f_best, newdata=testdat[, rfc1f_columns],
                  type="prob")
kmtest$prob01 <- preds01[, 2]

```

In [94]: kmp1p2_b\$centers

A matrix: 2 × 5 of type dbl

	pc1	pc2	pc3	pc4	prob01
1	0.29971	0.57830	0.83920	0.50138	0.045566
2	0.42677	0.65632	0.85543	0.44215	0.903993

In [97]: *# Function for obtaining a set of scores on the testset data
using kmp1p2_b as the classifier.*

```

### NOTE: the value of c1_to_InlandYES from above is being
### used in the following function.

get_testdatScores_kmp1p2_b <- function(seedv, dat) {

  n_smp <- 4000

  seedv_len <- length(seedv)
  datout <- rep(NA, 5 * seedv_len)
  dim(datout) <- c(seedv_len, 5)
  datout <- as.data.frame(datout)
  colnames(datout) <- c("fscore", "Acc", "Type2", "FN", "FP")
  rownames(datout) <- as.character(seedv)

  # Using model kmp1p2_b from above.

```



```

ctr_list <- vector("list", length= n_smp)
for(i in 1:n_smp) {
  ctr_list[[i]] <- km1p2_b$centers
}

for(h in 1:seedv_len) {
  # shuffle dat
  cur_seed <- seedv[h]
  set.seed(cur_seed)
  # It is expected that dat is testdat, which has 10K rcds
  smp <- sample(rownames(dat), n_smp, replace= FALSE)
  df <- dat[smp,]

  # CAUTION: df has the prob01 & prob02 columns.

  # Apply PCA.
  pca_columns <- c(km_predictors, "prob02")
  df_pca <- predict(pca4, df[, pca_columns])

  # Apply min-max scaling to df_scaled.
  df_scaled <- apply(df_pca, MARGIN=2, range01)
  df_scaled <- as.data.frame(cbind(df_scaled, df$prob01),
                             row.names=rownames(df))
  colnames(df_scaled) <- c(paste0("pc", 1:4), "prob01")

  # Each element of the following list is a row of df.
  df_asList <- split(df_scaled[, colnames(km1p2_b$centers)], seq(n_smp))
  names(ctr_list) <- rownames(df)

  # Get the predictions for df.
  preds <- mcmapply(getCluster, df_asList, ctr_list,
                    SIMPLIFY=TRUE, mc.cores=6)
  df$cluster <- as.numeric(preds)

  df$pred_Inland <- NA
  if(c1_to_InlandYES) {
    df[which(df$cluster==1),]$pred_Inland <- 1
    df[which(df$cluster==2),]$pred_Inland <- 0
  } else {
    df[which(df$cluster==1),]$pred_Inland <- 0
    df[which(df$cluster==2),]$pred_Inland <- 1
  }

  # Generate confusion matrix.
  preds <- as.factor(df$pred_Inland)
  names(preds) <- rownames(df)
  ans <- get_confusion(preds, df[, "Inland", drop=FALSE])

  mat <- as.matrix(ans[[1]])
  fscore <- round(as.numeric(ans[[2]]), 4)
  acc <- round(sum(diag(mat))/floor(sum(mat)), 4)
  type2 <- round((0.4 * acc + 0.6 * ans[[2]]), 4)
  FN <- as.numeric(mat[2,1])
  FP <- as.numeric(mat[1,2])
  datout[as.character(cur_seed), 1:5] <- c(fscore, acc, type2, FN, FP)
}
return(datout)
}

```

In [98]: # Get km1p2_b's scores on the testset data. For each of
 # the 250 seeds, I sample 4K (no replacement) from the 10K
 # set of testdat records.

```

set.seed(1821)
seed_vector <- sample(1:9999, 250, replace=FALSE)

start <- Sys.time()
paste("Start time: ", start, sep="")
dat_result <- get_testdatScores_km1p2_b(seed_vector, kmtest)

```

```
stop <- Sys.time()
round(stop - start, 2)
# Time difference of 8.23 mins
```

'Start time: 2021-04-30 10:04:23'

Time difference of 8.23 mins

```
In [99]: dim(dat_result)
head(dat_result)
```

250 · 5

A data.frame: 6 × 5

	fscore	Acc	Type2	FN	FP
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
5934	0.8356	0.8965	0.8600	225	189
1953	0.8456	0.9010	0.8678	216	180
7591	0.8327	0.8932	0.8569	262	165
1038	0.8362	0.8985	0.8611	233	173
49	0.8337	0.8932	0.8575	231	196
3203	0.8361	0.8970	0.8605	228	184

```
In [100]: fscore_mean <- round(mean(dat_result$fscore), 4)
fscore_sd <- round(sd(dat_result$fscore), 4)
paste0("fscore mean: ", as.character(fscore_mean))
# 0.8355
paste0("fscore StdDev: ", as.character(fscore_sd))
# 0.006
""
summary(dat_result$fscore)
```

'fscore mean: 0.8355'

'fscore StdDev: 0.006'

"

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	0.821	0.831	0.836	0.836	0.839	0.852

```
In [101]: Acc_mean <- round(mean(dat_result$Acc), 4)
Acc_sd <- round(sd(dat_result$Acc), 4)
paste0("accuracy mean: ", as.character(Acc_mean))
# 0.8963
paste0("accuracy StdDev: ", as.character(Acc_sd))
# 0.0037
""
summary(dat_result$Acc)
```

'accuracy mean: 0.8963'

'accuracy StdDev: 0.0037'

"

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	0.887	0.894	0.896	0.896	0.899	0.907

```
In [102]: Type2_mean <- round(mean(dat_result$Type2), 4)
Type2_sd <- round(sd(dat_result$Type2), 4)
paste0("Type2 mean: ", as.character(Type2_mean))
# 0.8599
paste0("Type2 StdDev: ", as.character(Type2_sd))
```

```
# 0.005
""
summary(dat_result$Type2)

'Type2 mean: 0.8599'

'Type2 StdDev: 0.005'

"

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
0.848   0.857   0.860   0.860   0.863   0.874
```

```
In [103]: FN_mean <- round(mean(dat_result$FN), 2)
FN_sd <- round(sd(dat_result$FN), 2)
paste0("FN mean: ", as.character(FN_mean))
# 234.2
paste0("FN StdDev: ", as.character(FN_sd))
# 11.4
""
summary(dat_result$FN)

'FN mean: 234.18'

'FN StdDev: 11.43'

"

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
195      227      234      234      242      262
```

```
In [104]: FP_mean <- round(mean(dat_result$FP), 2)
FP_sd <- round(sd(dat_result$FP), 2)
paste0("FP mean: ", as.character(FP_mean))
# 180.5
paste0("FP StdDev: ", as.character(FP_sd))
# 10.2
""
summary(dat_result$FP)

'FP mean: 180.53'

'FP StdDev: 10.18'

"

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
156      173      181      181      189      206
```

Section 7 Comments

The scores in Section 7 are essentially equivalent to those in Section 6. So there is no difference between the two approaches.

Table Summarizing Part 3 Results

```
In [7]: results <- read.csv("/home/greg/Documents/stat/github_repos/cows/model_results_Part3_CAhous
      header=TRUE, row.names=1)
dim(results)
```

```
9 10
```

```
In [8]: results
```

A data.frame: 9 × 10

	fscore	fscore_sd	Type2	Type2_sd	accuracy	acc_sd	FN	FN_sd	FP	FP_sd
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
rfclf	0.8366	0.0060	0.8611	0.0049	0.8979	0.0037	242.3	11.6	166.0	9.8
gb	0.8204	0.0065	0.8475	0.0054	0.8882	0.0039	266.0	11.9	181.2	10.5
svm02	0.8265	0.0065	0.8536	0.0053	0.8943	0.0038	280.3	11.8	142.6	8.9
f03	0.7765	0.0076	0.8149	0.0060	0.8725	0.0042	401.6	14.3	108.6	8.4
kmp1_pca4	0.8348	0.0061	0.8592	0.0050	0.8957	0.0037	233.1	11.3	184.2	10.6
kmp1_pca4_wghts	0.8349	0.0060	0.8593	0.0050	0.8959	0.0037	234.7	11.3	181.9	10.2
kmp1p2	0.8346	0.0059	0.8593	0.0049	0.8962	0.0037	240.2	11.4	174.9	9.7
kmp1p2_pca5	0.8357	0.0061	0.8600	0.0050	0.8965	0.0037	234.6	11.5	179.5	10.3
kmp1p2_pca4	0.8355	0.0060	0.8599	0.0050	0.8963	0.0037	234.2	11.4	180.5	10.2

Part 3 Final Comments

The k-means approach that produced better supervised learning models in Parts 1 and 2 is not yielding the same positive result for the California housing dataset. For this dataset, our best k-means model still does not quite measure up to our best SL model (in this case, rfclf_best). On average, rfclf_best classifies 5.6 more districts correctly out of 4K districts (0.14%) than does kmp1p2_pca5.

When comparing the results of Part 3 with those from Parts 1 and 2, it is important to keep in mind that the predictors used in Part 3 were not collected with the aim of distinguishing between INLAND and OCEAN districts; whereas with the cow data, the predictors were specifically chosen to help distinguish between survivors and non-survivors. One of the implications of this is that our first k-means models in Parts 1 and 2, absent any probability columns, were very competitive with the very best SL models. I was then able to add the probabilities from one or more of the best SL models, employ weights, and get an even stronger model. Here in Part 3, though, I start with a very weak k-means "base" model. It is not competitive until I add the prob01 column. Applying weights is not able to boost the score in any significant way. Adding a second probability column also did not make an important positive difference.

The results of Part 3 should thus not be seen as a mark against the k-means hybrid approach to model construction. We shouldn't expect the hybrid approach to always provide us with the strongest model.

The above work with weights is important. The more predictors we have, the more difficult it is to search the parameter space for the optimal weights. This difficulty increases as the number of records in our dataset increases. We saw above, however, that this search may not always be necessary. But more work is required in order to find out the conditions which, if satisfied, make a search for weights superfluous.

In []: